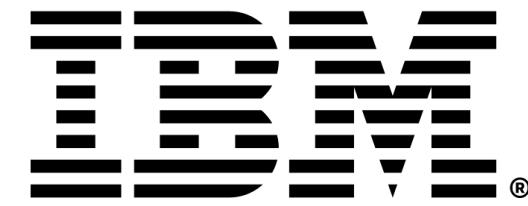


Red Hat



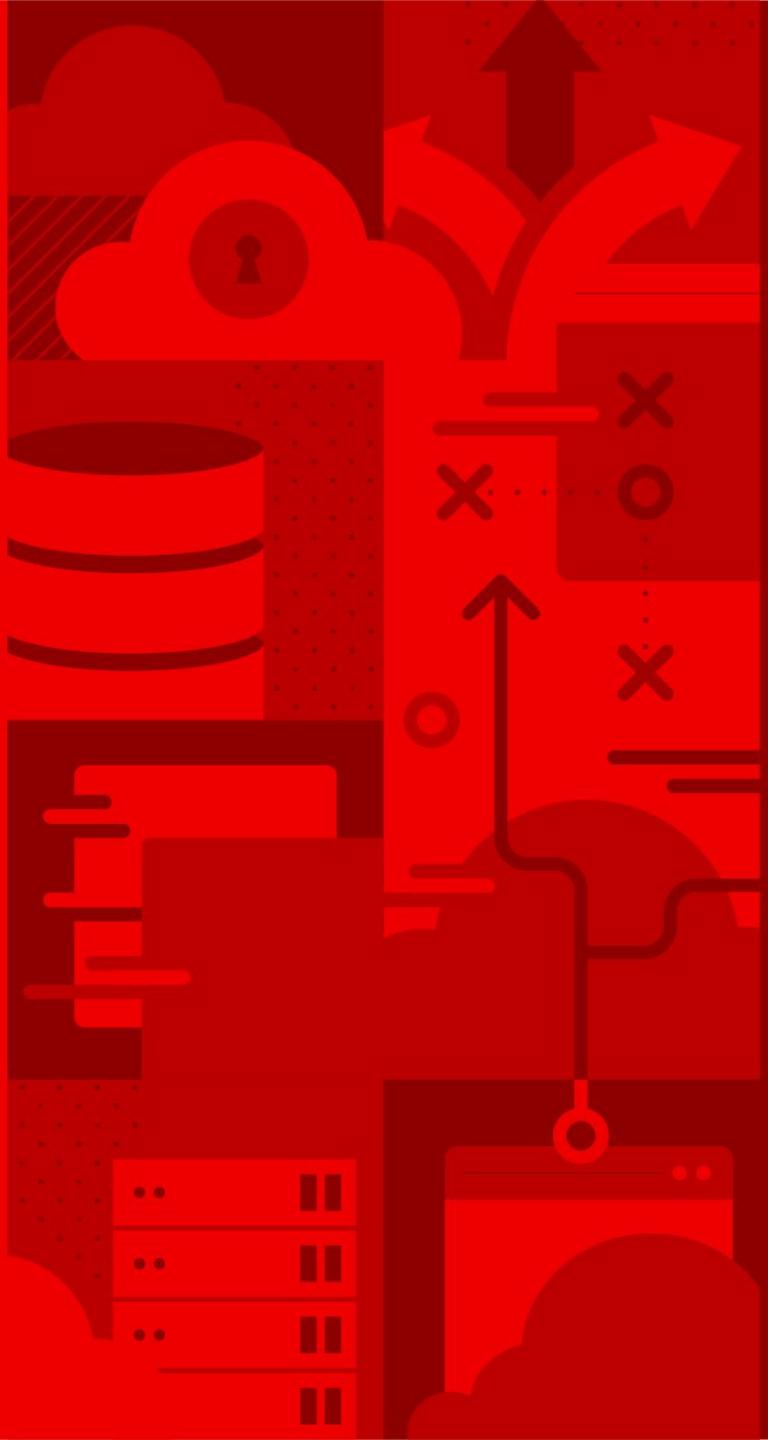


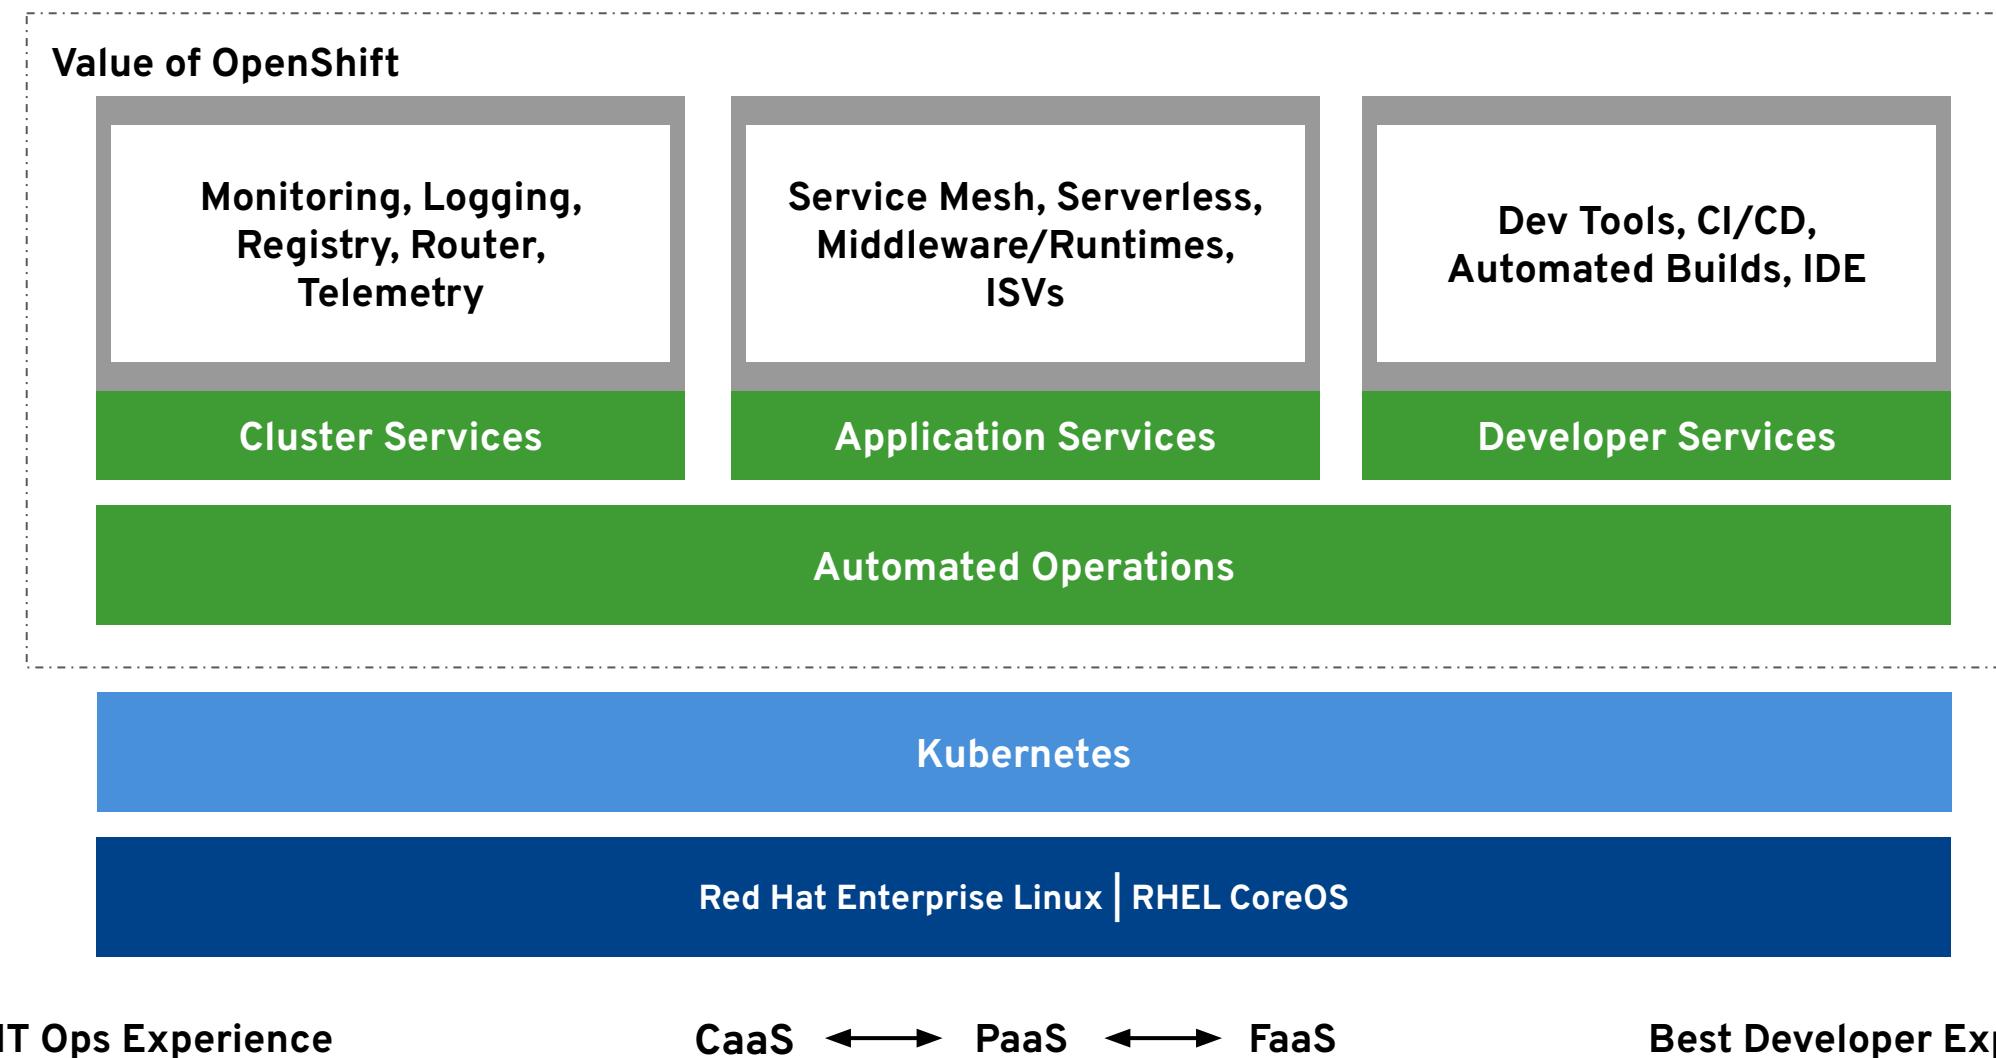
OpenShift Architecture

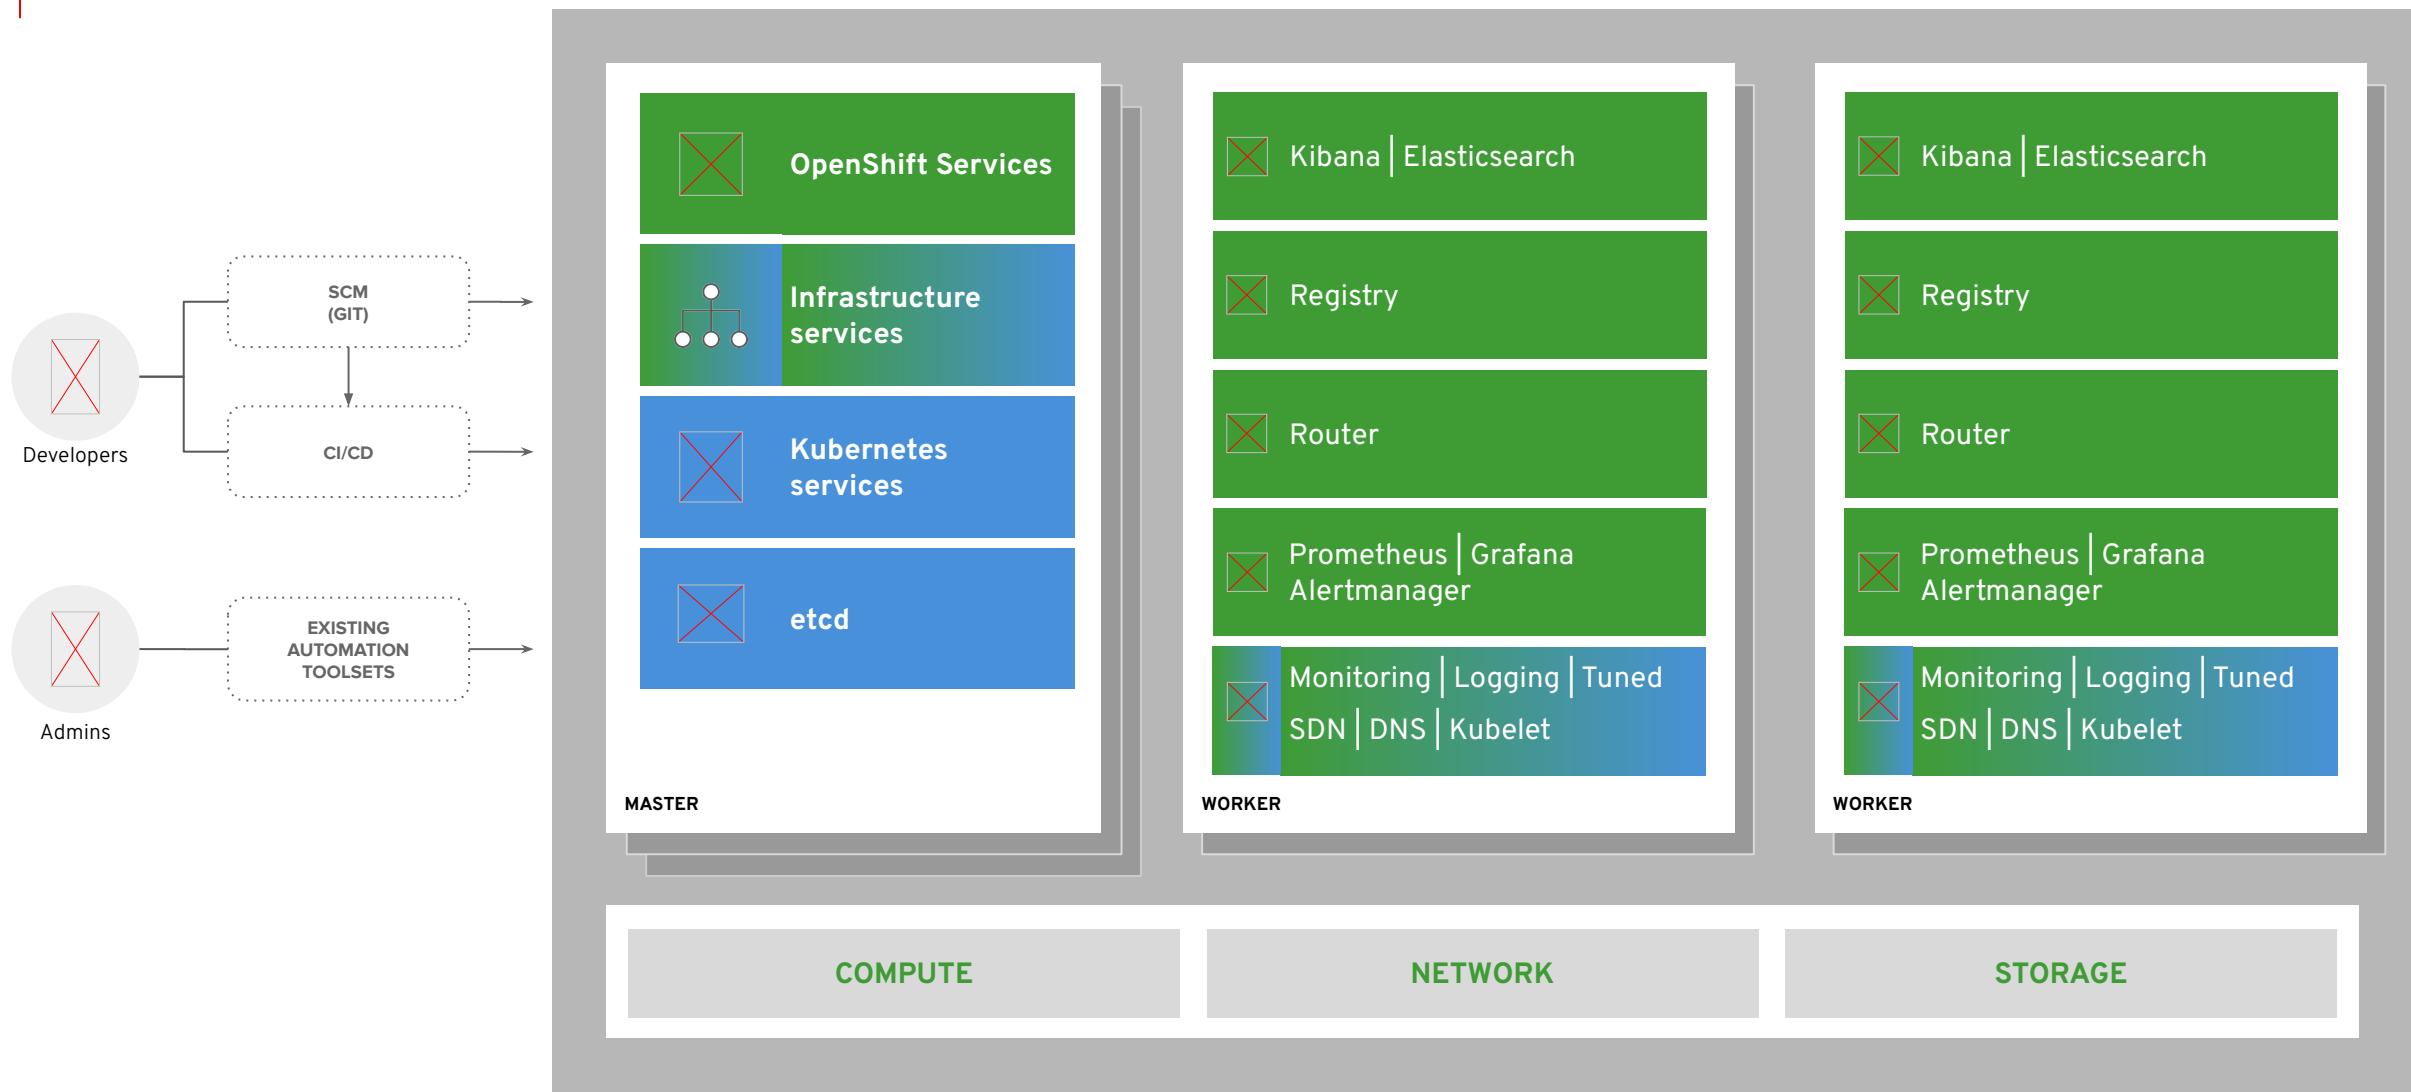
OpenShift Architecture Workshop for IBM Nordics

Alfred Bach
PESA RedHat EMEA
abach@redhat.com

Functional overview









OpenShift and Kubernetes core concepts

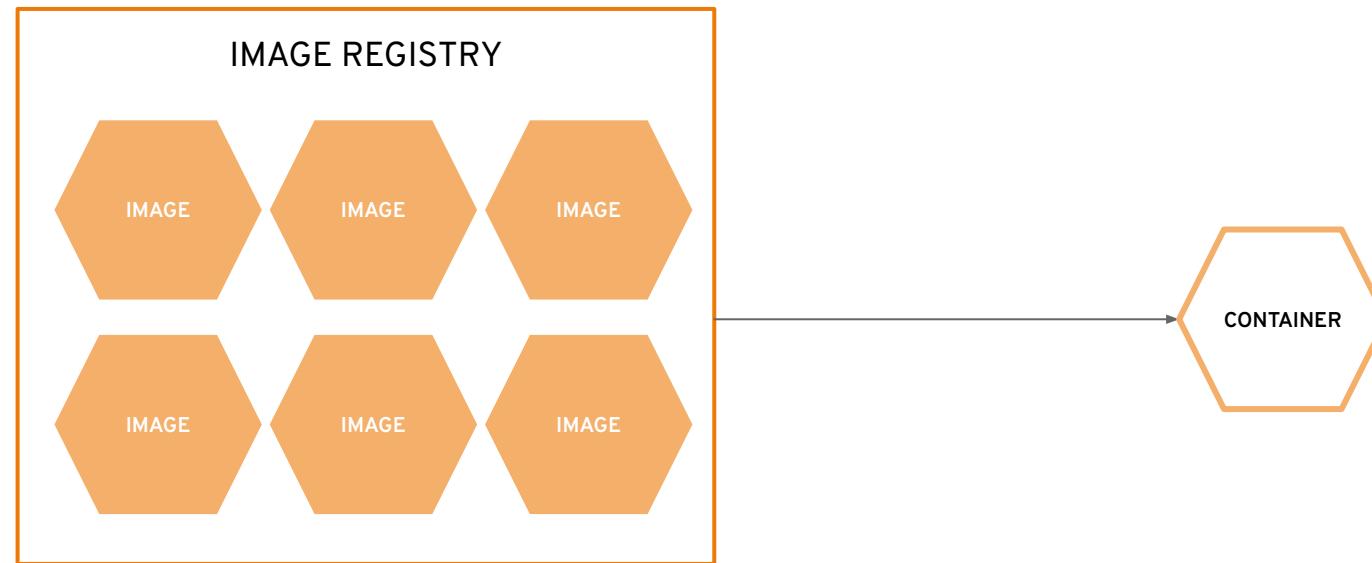
a container is the smallest compute unit



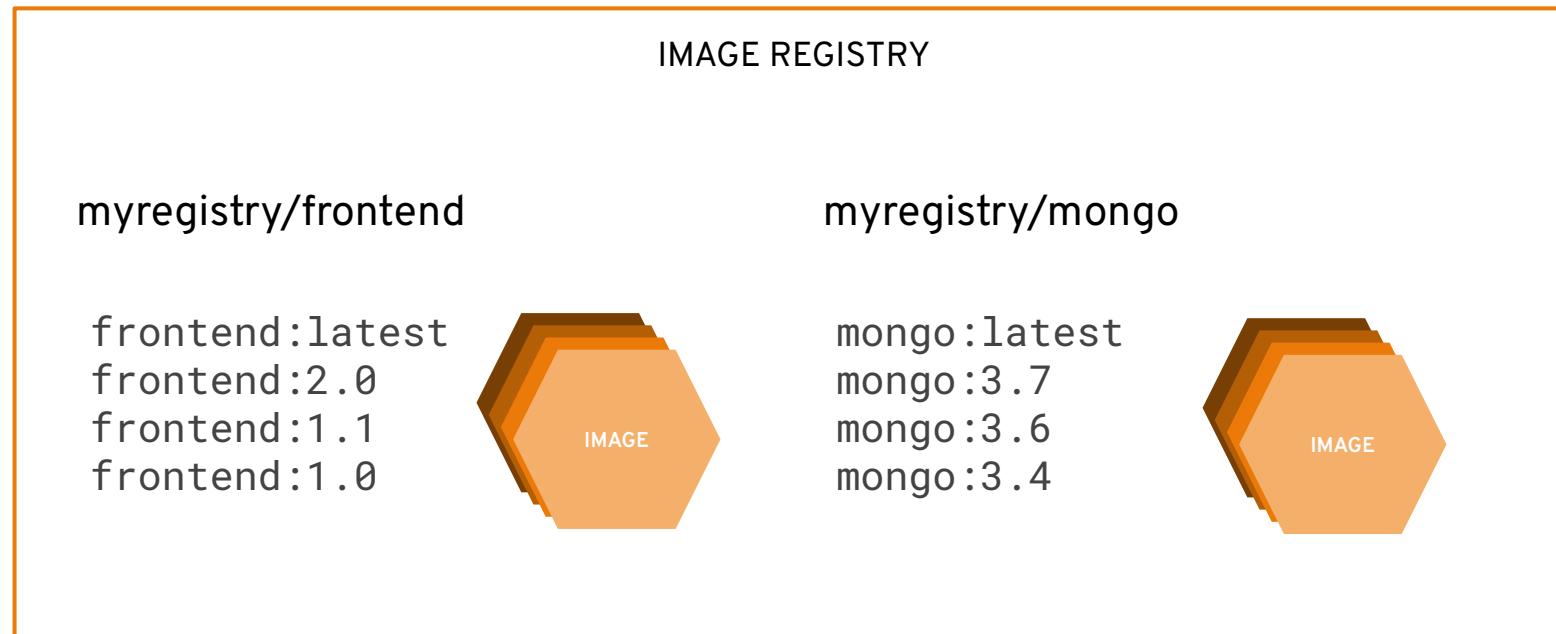
containers are created from container images



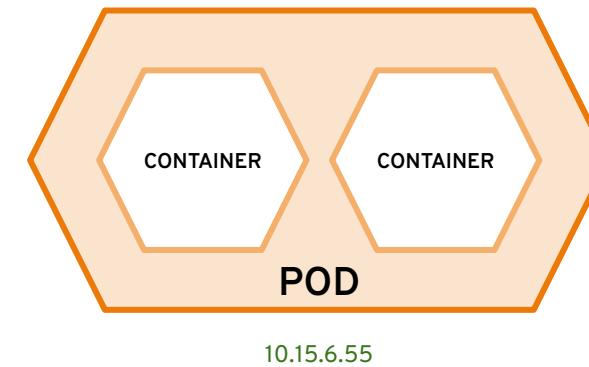
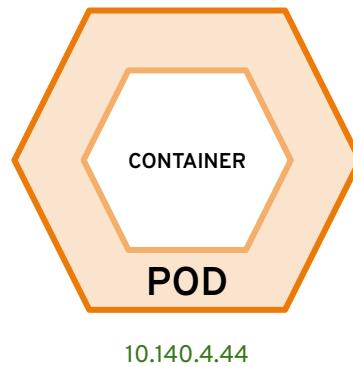
container images are stored in an image registry



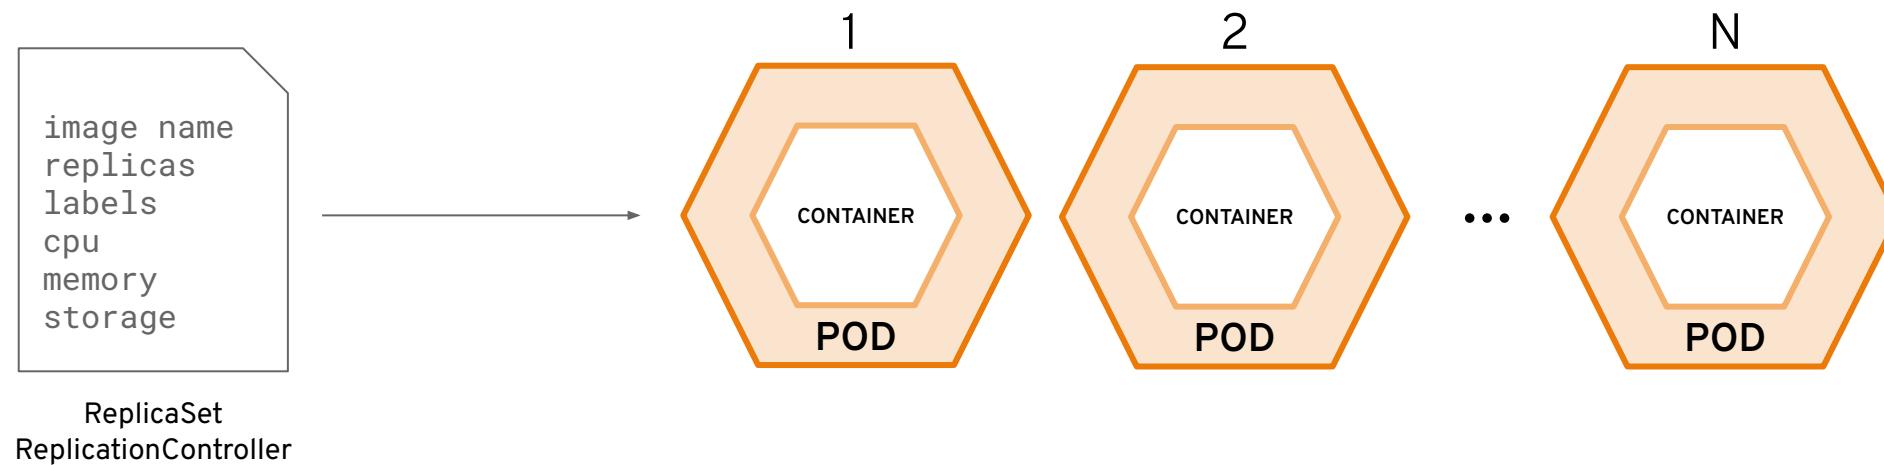
an image repository contains all versions of an image in the image registry



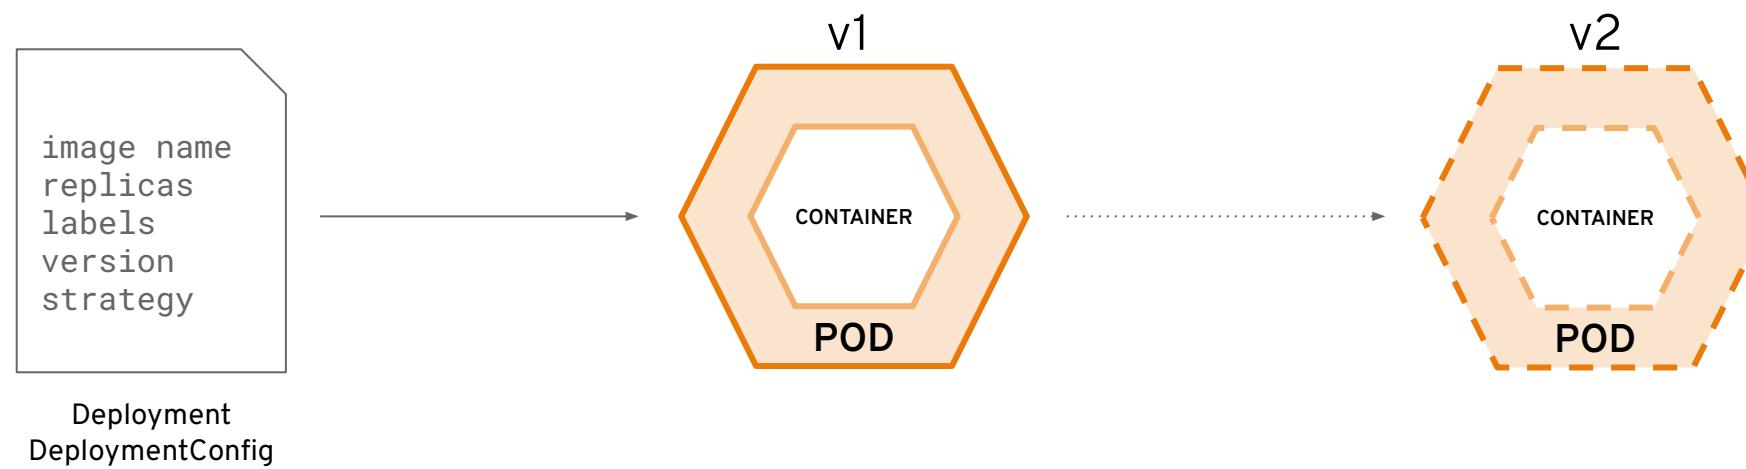
containers are wrapped in pods which are units of deployment and management



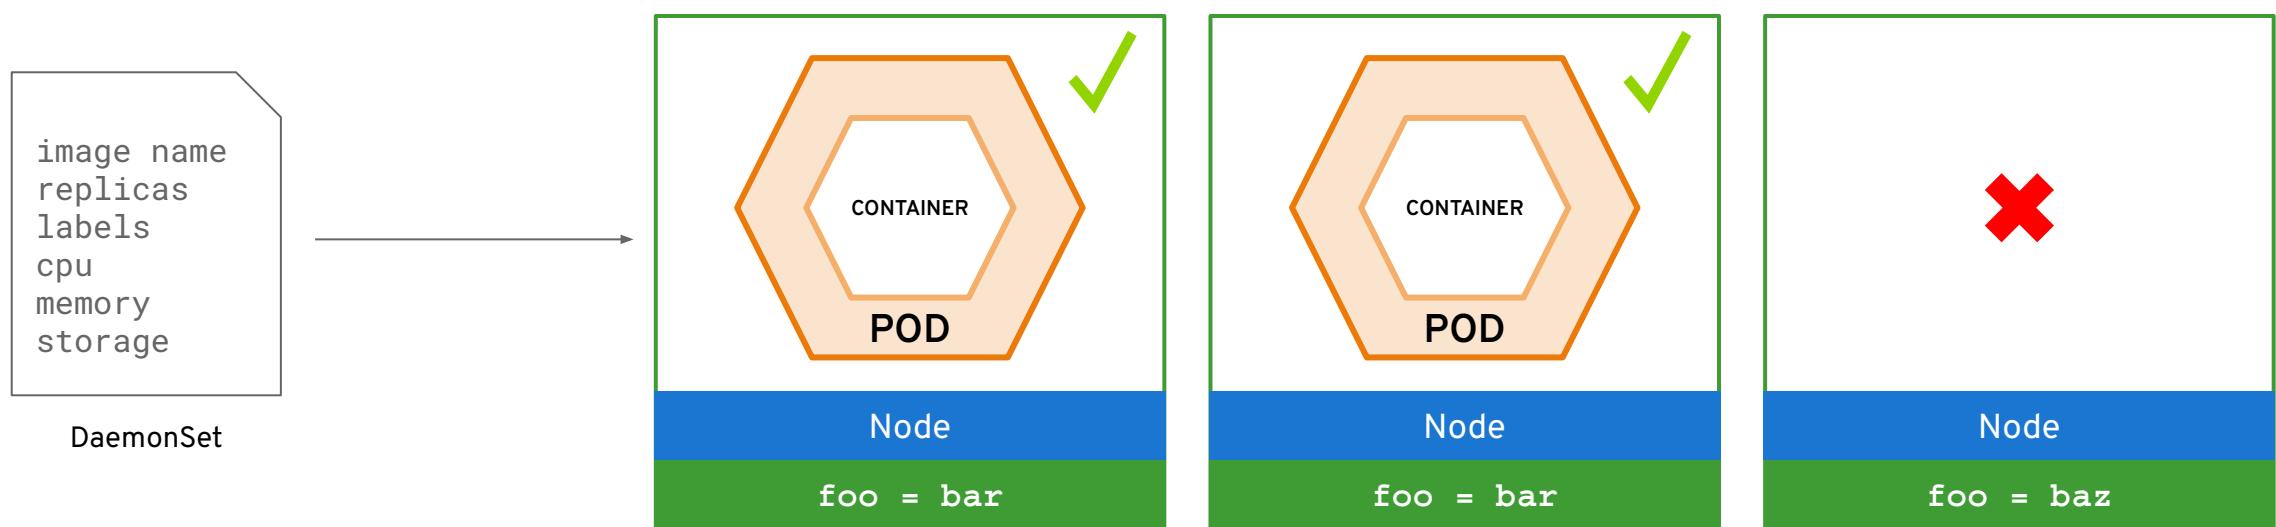
ReplicationControllers & ReplicaSets ensure a specified number of pods are running at any given time



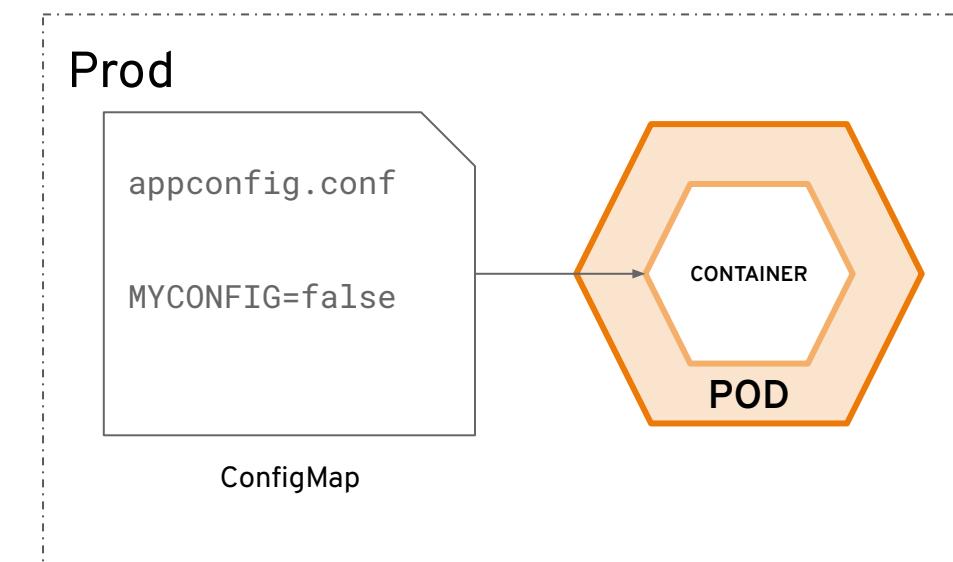
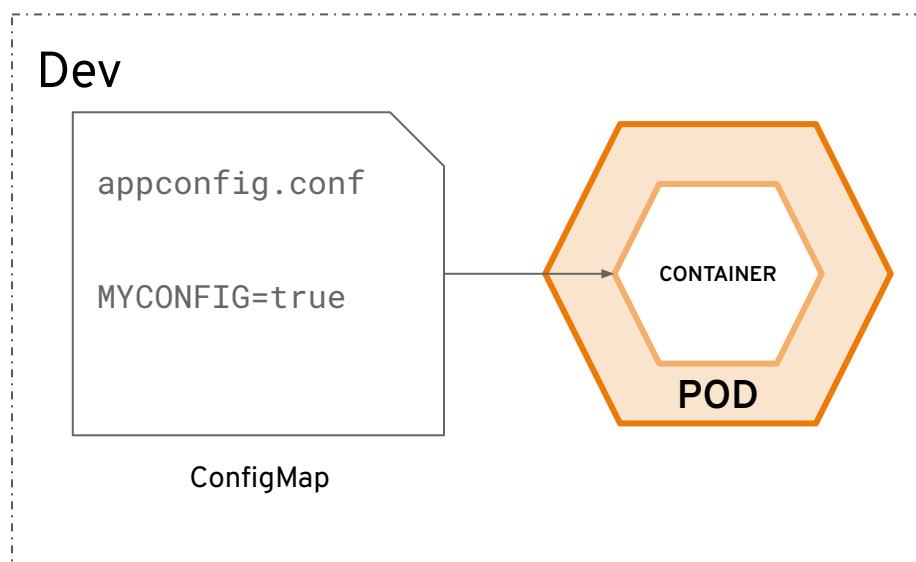
Deployments and DeploymentConfigurations define how to roll out new versions of Pods



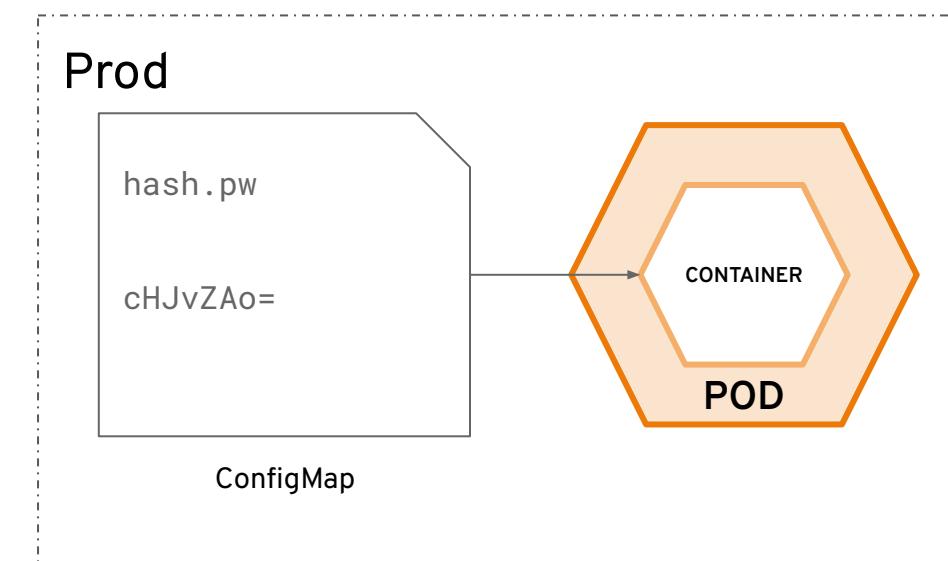
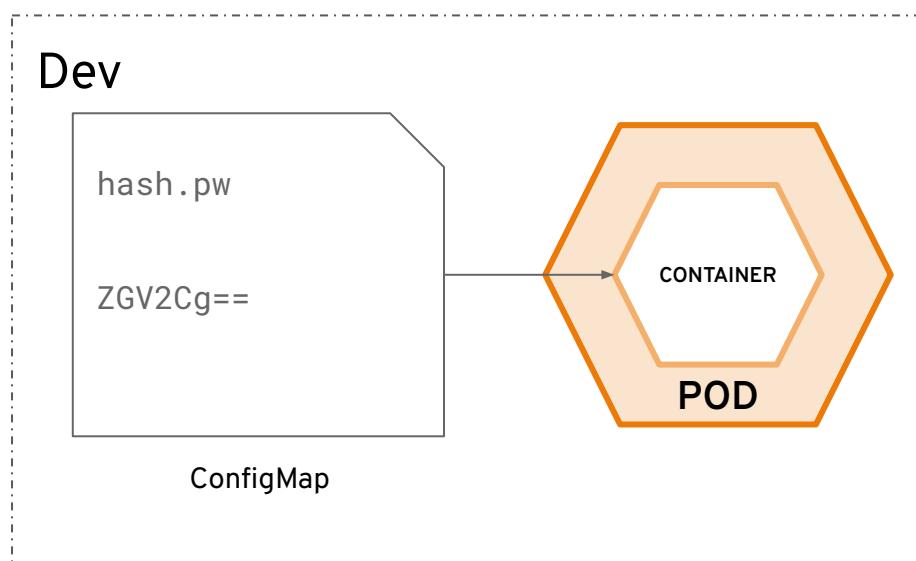
a daemonset ensures that all (or some) nodes run a copy of a pod



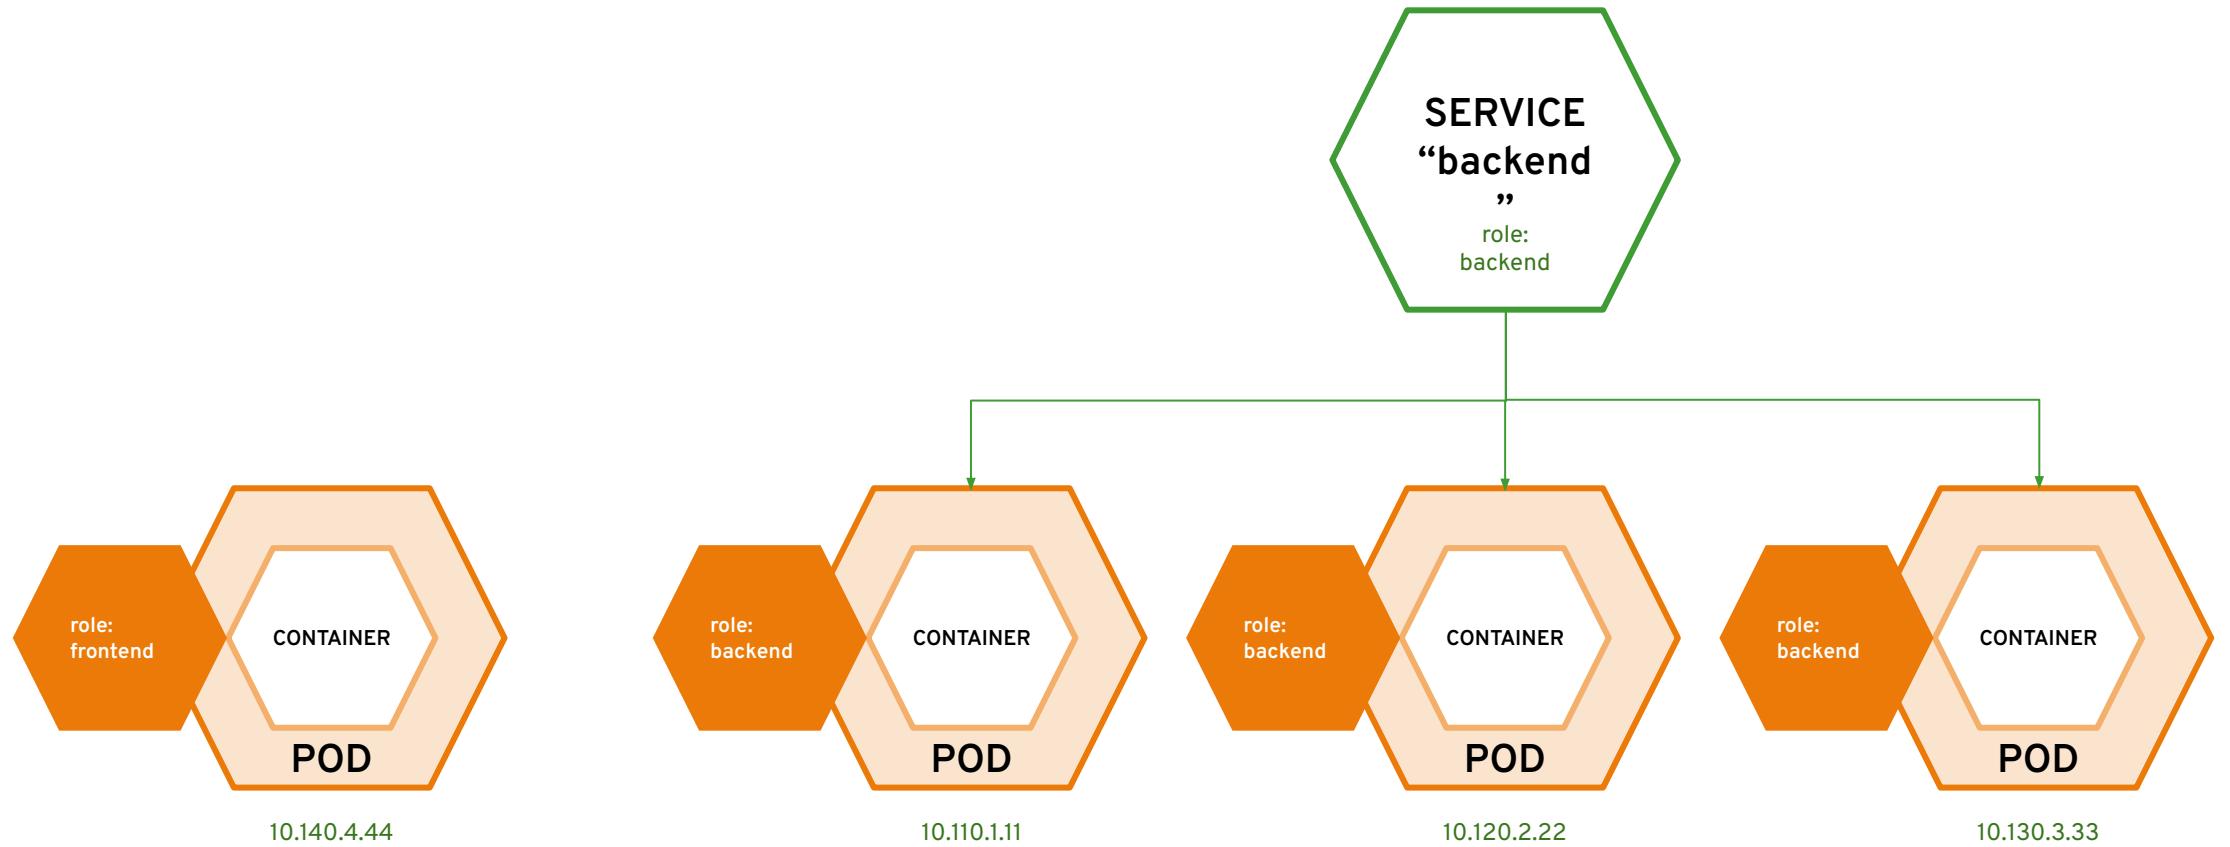
configmaps allow you to decouple configuration artifacts from image content



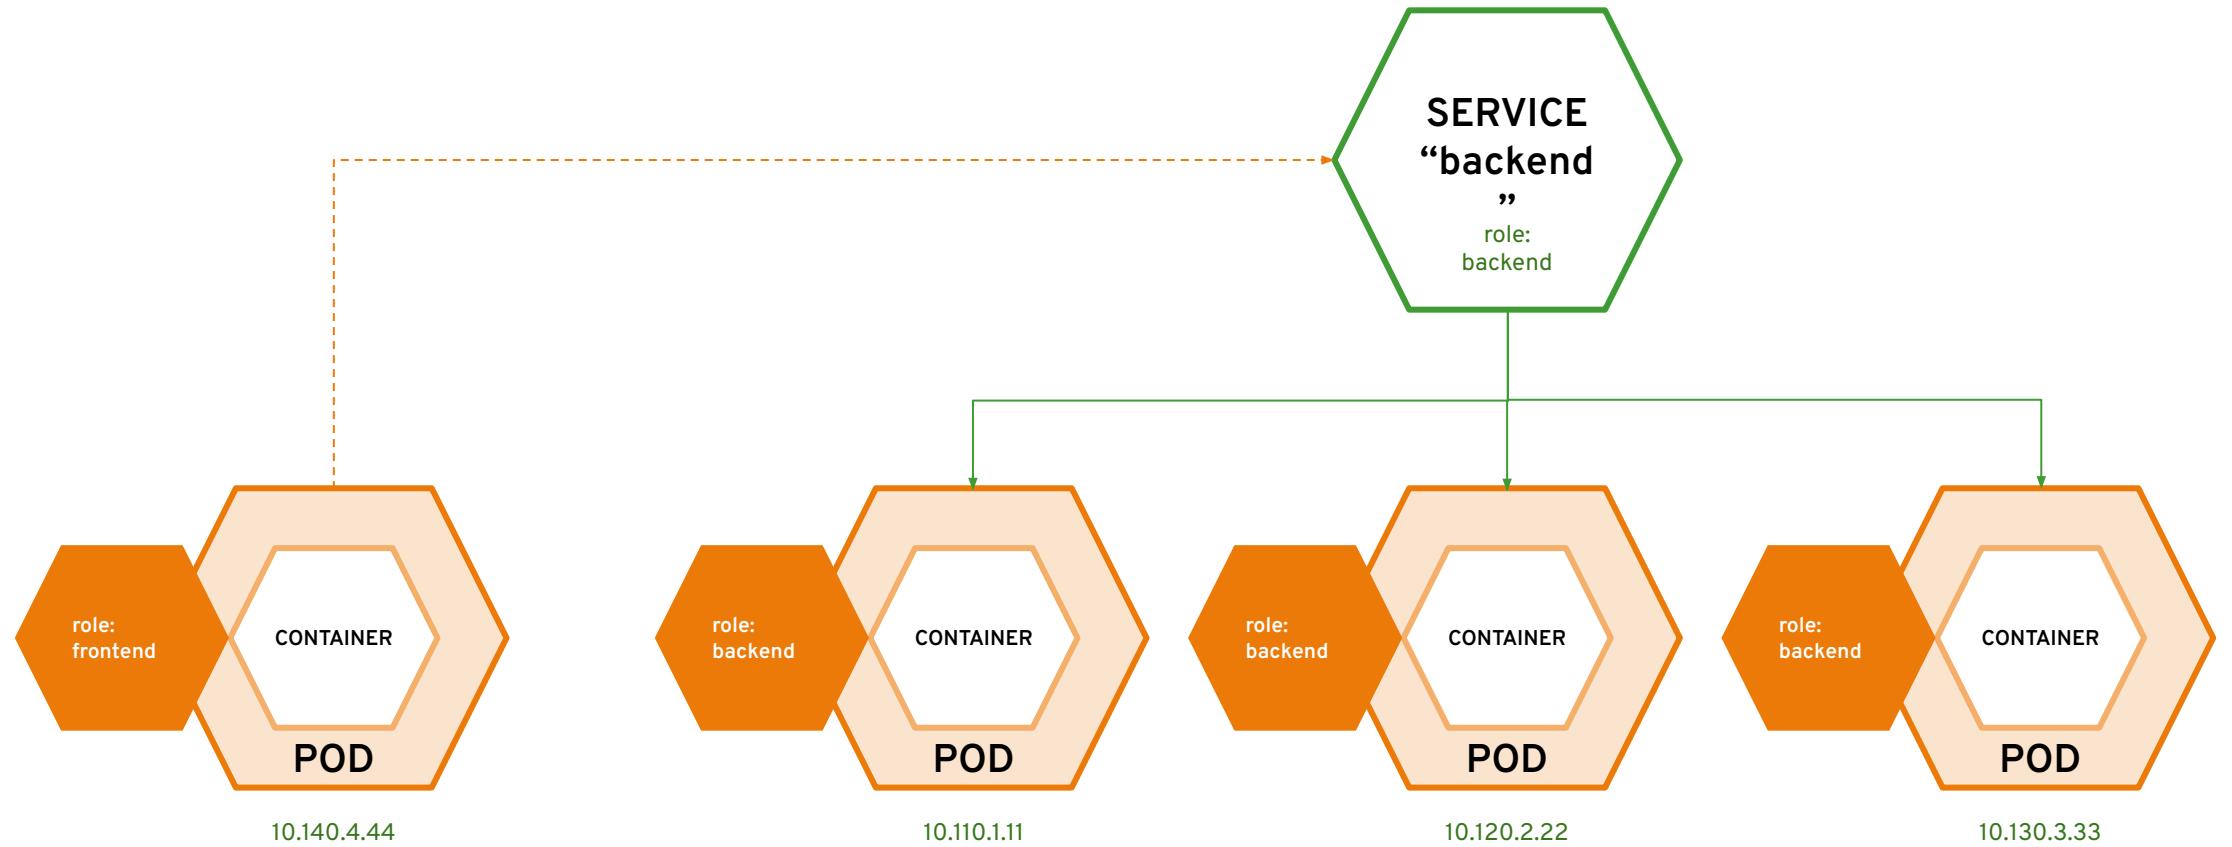
secrets provide a mechanism to hold sensitive information such as passwords



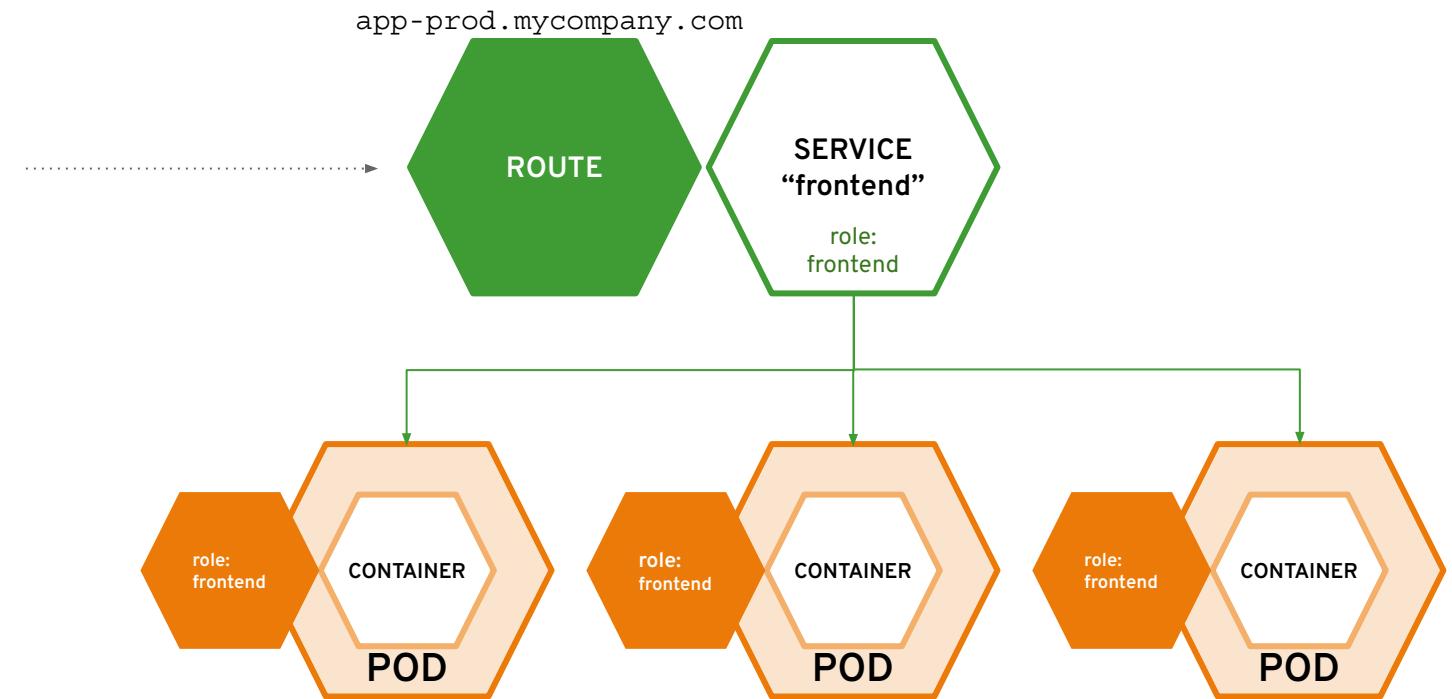
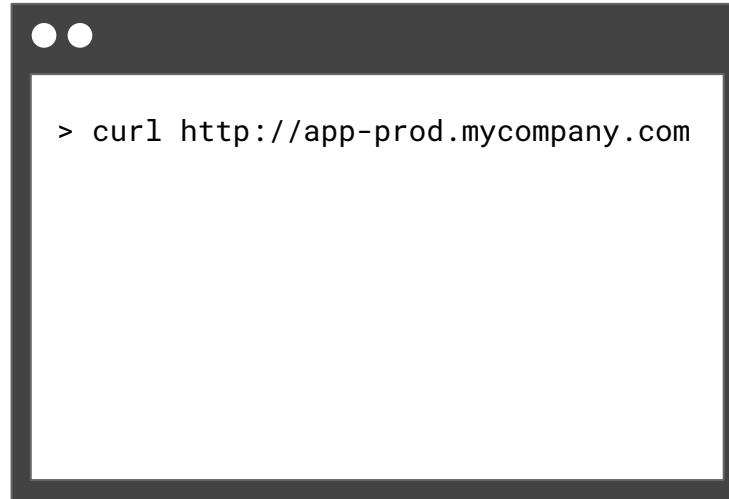
services provide internal load-balancing and service discovery across pods



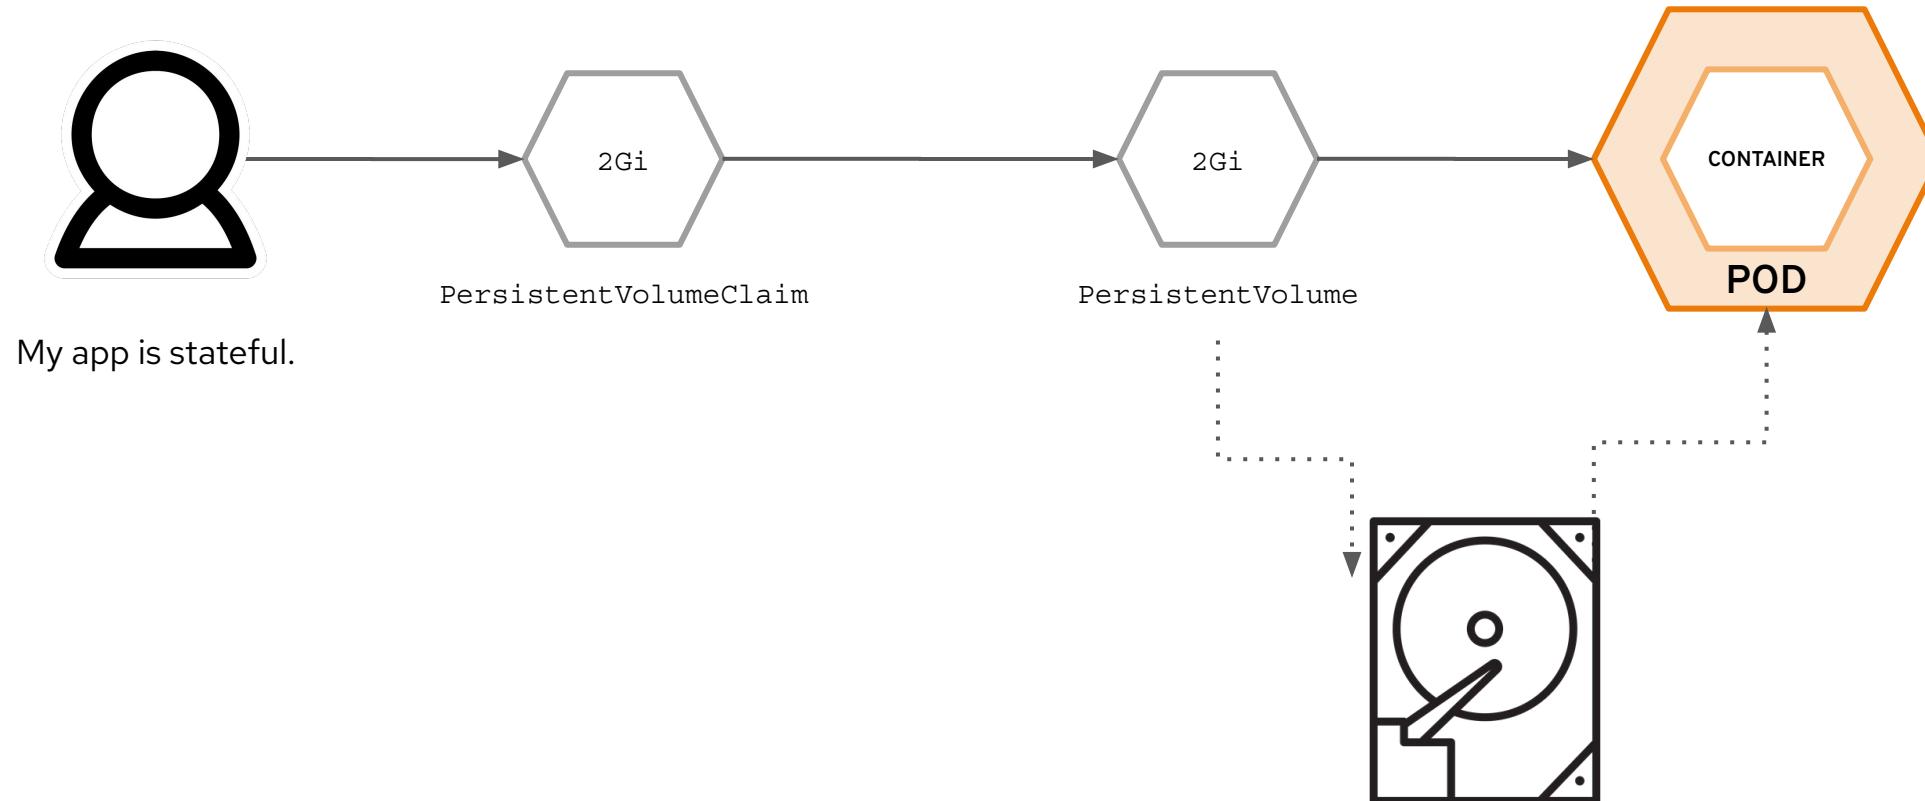
apps can talk to each other via services



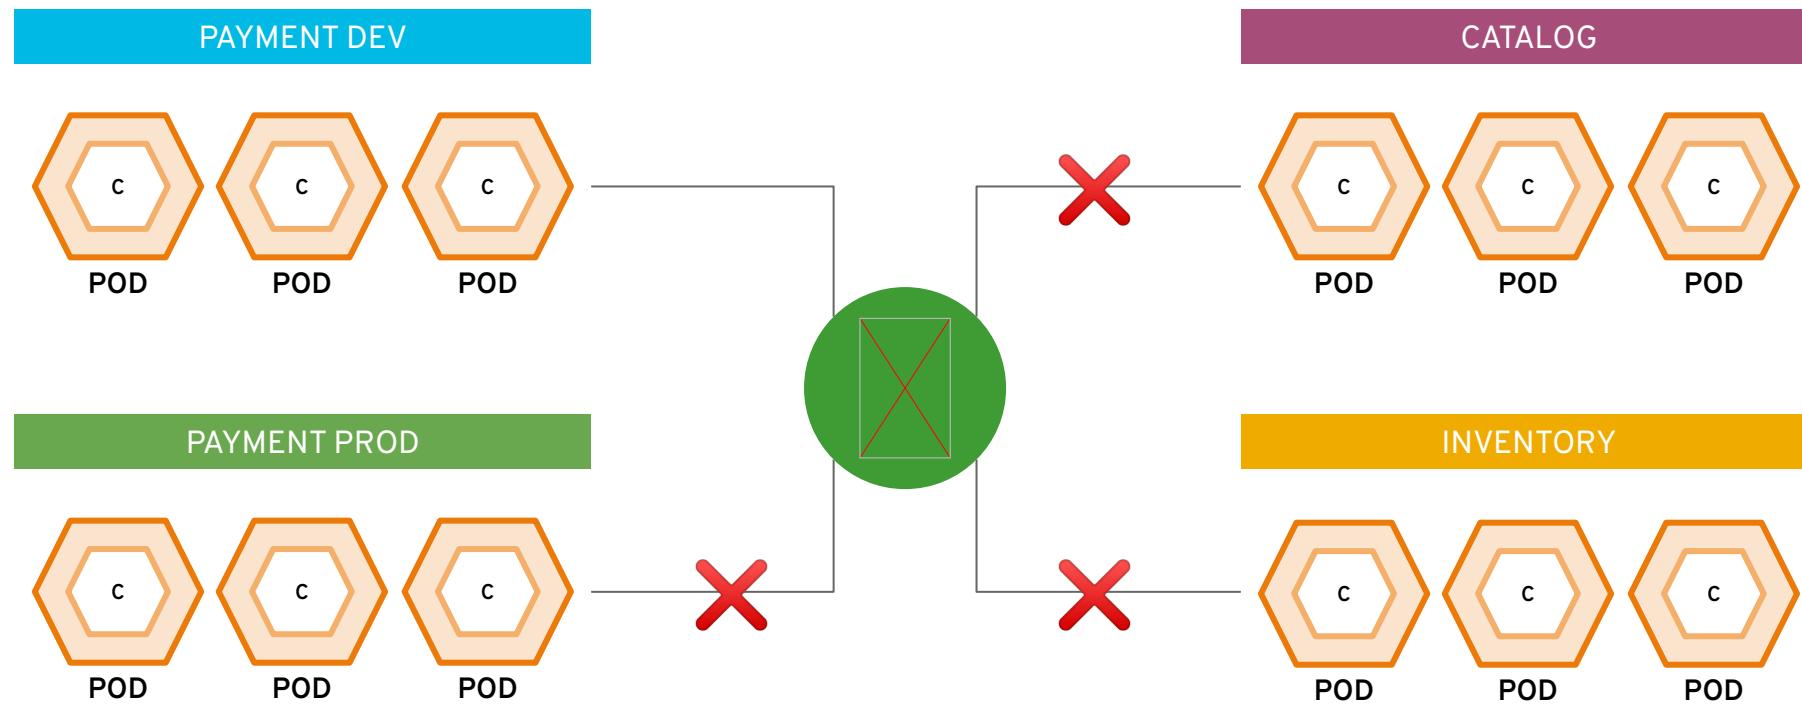
routes make services accessible to clients outside the environment via real-world urls



Persistent Volume and Claims



projects isolate apps across environments,
teams, groups and departments





OpenShift 4 Architecture

your choice of infrastructure

COMPUTE

NETWORK

STORAGE



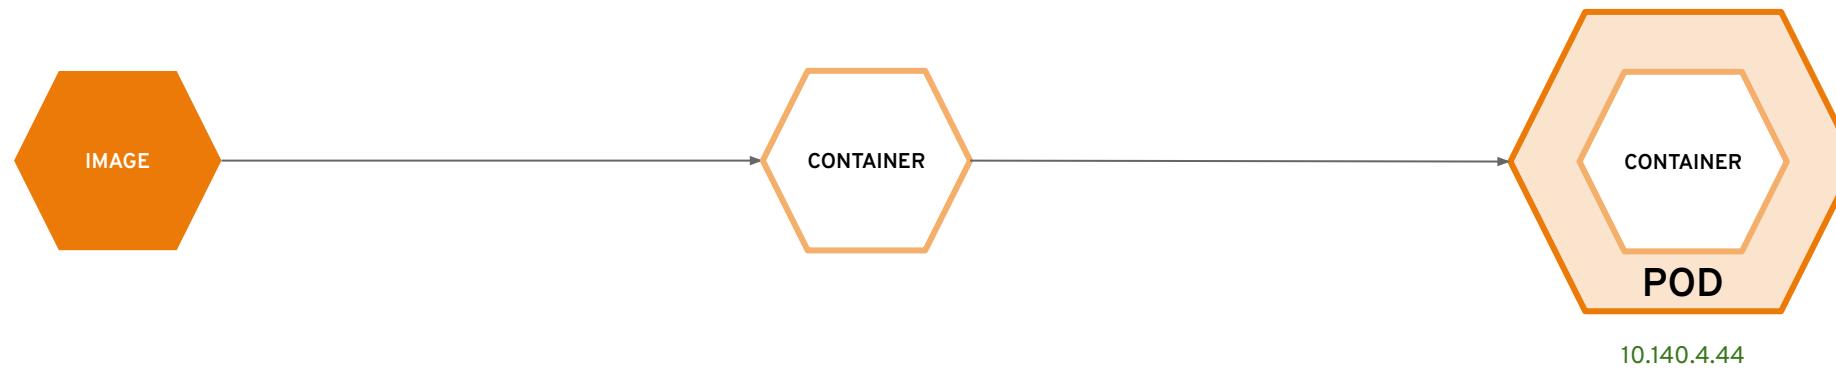
workers run workloads



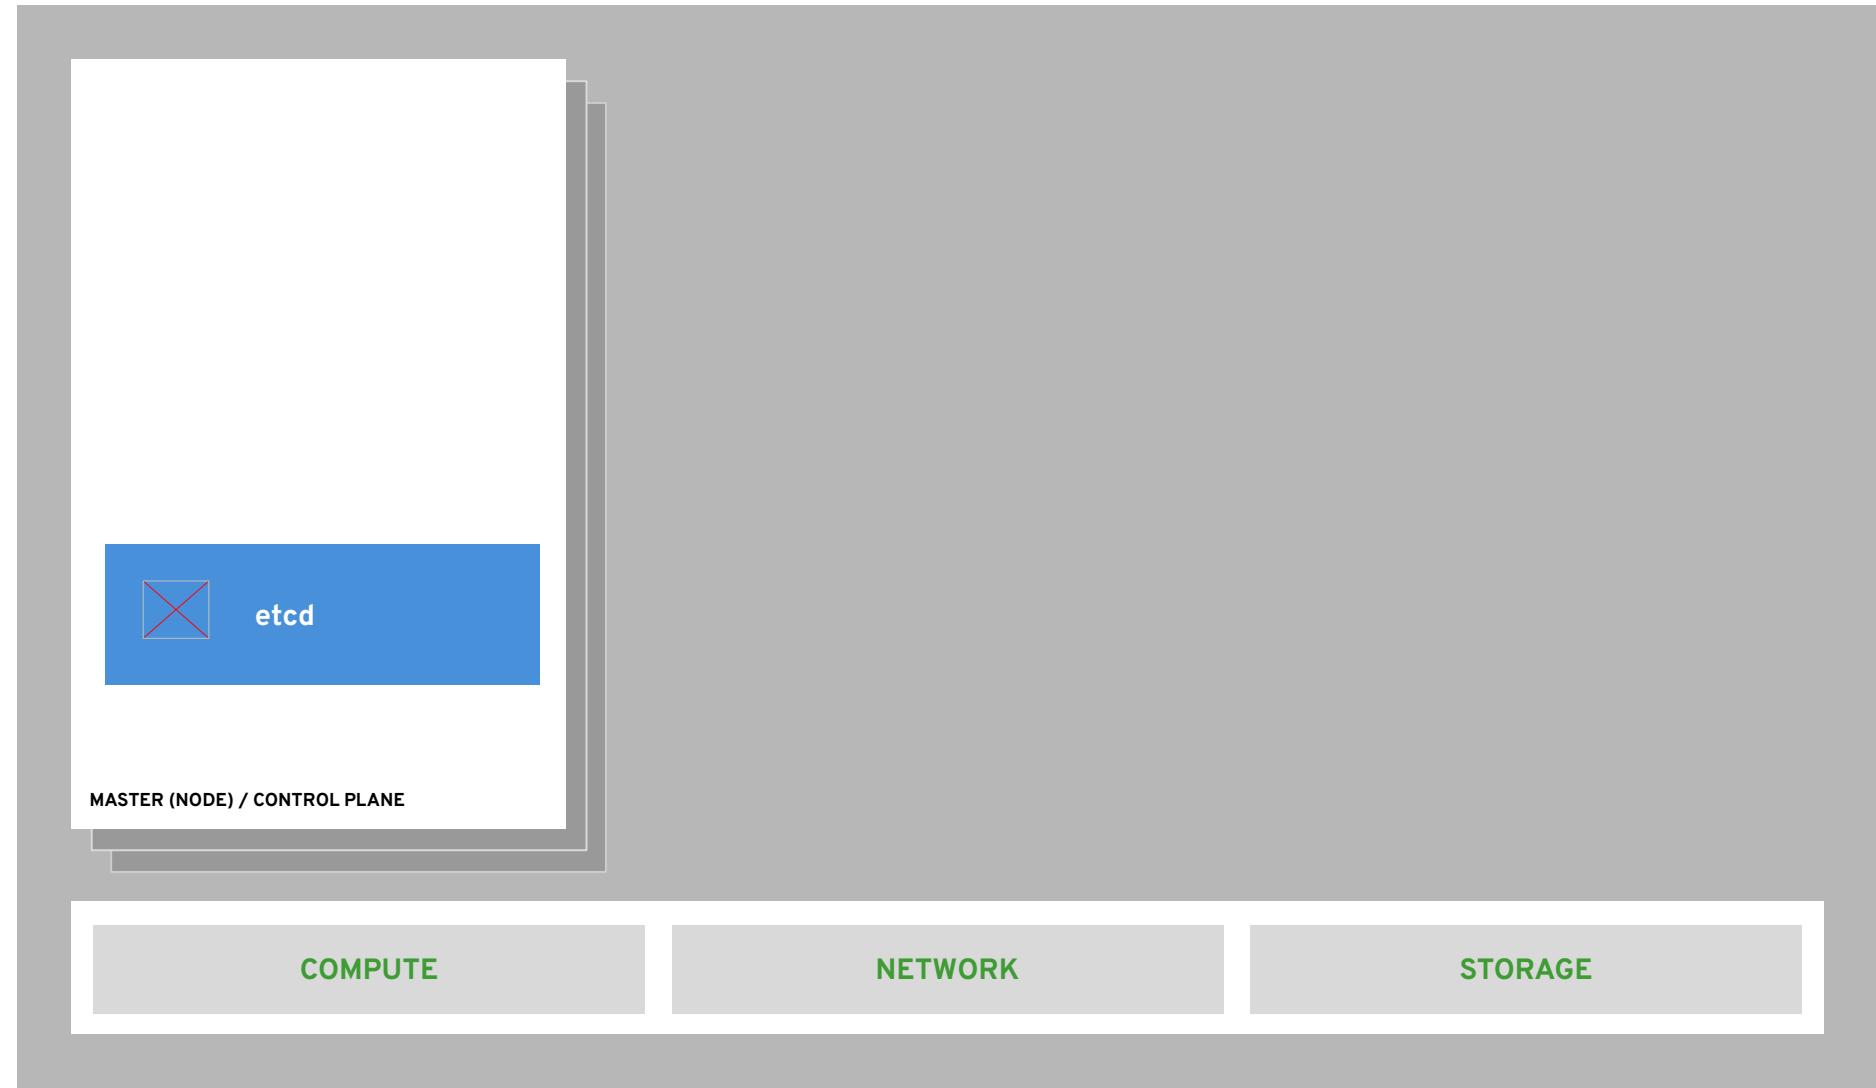
masters are the control plane



everything runs in pods



state of everything



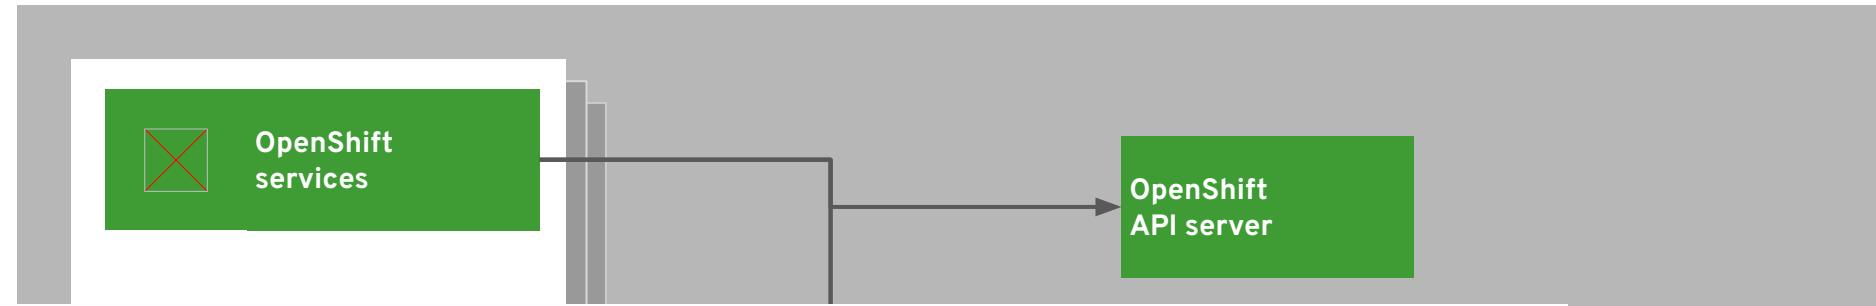
Introducing the Declarative Architecture of Kubernetes

The architecture of OpenShift is based on the declarative nature of Kubernetes. Most system administrators are used to imperative architectures, where you perform actions that indirectly change the state of the system, such as starting and stopping containers on a given server. In a declarative architecture, you change the state of the system and the system updates itself to comply with the new state. For example, with Kubernetes, you define a pod resource that specifies that a certain container should run under specific conditions. Then Kubernetes finds a server (a node) that can run that container under these specific conditions.

Declarative architectures allow for self-optimizing and self-healing systems that are easier to manage than imperative architectures.

Kubernetes defines the state of its cluster, including the set of deployed applications, as a set of resources stored in the etcd database. Kubernetes also runs controllers that monitor these resources and compares them to the current state of the cluster. These controllers take any action necessary to reconcile the state of the cluster with the state of the resources, for example by finding a node with sufficient CPU capacity to start a new container from a new pod resource.

Kubernetes provides a REST API to manage these resources. All actions that an OpenShift user takes, either using the command-line interface or the web console, are performed by invoking this REST API.



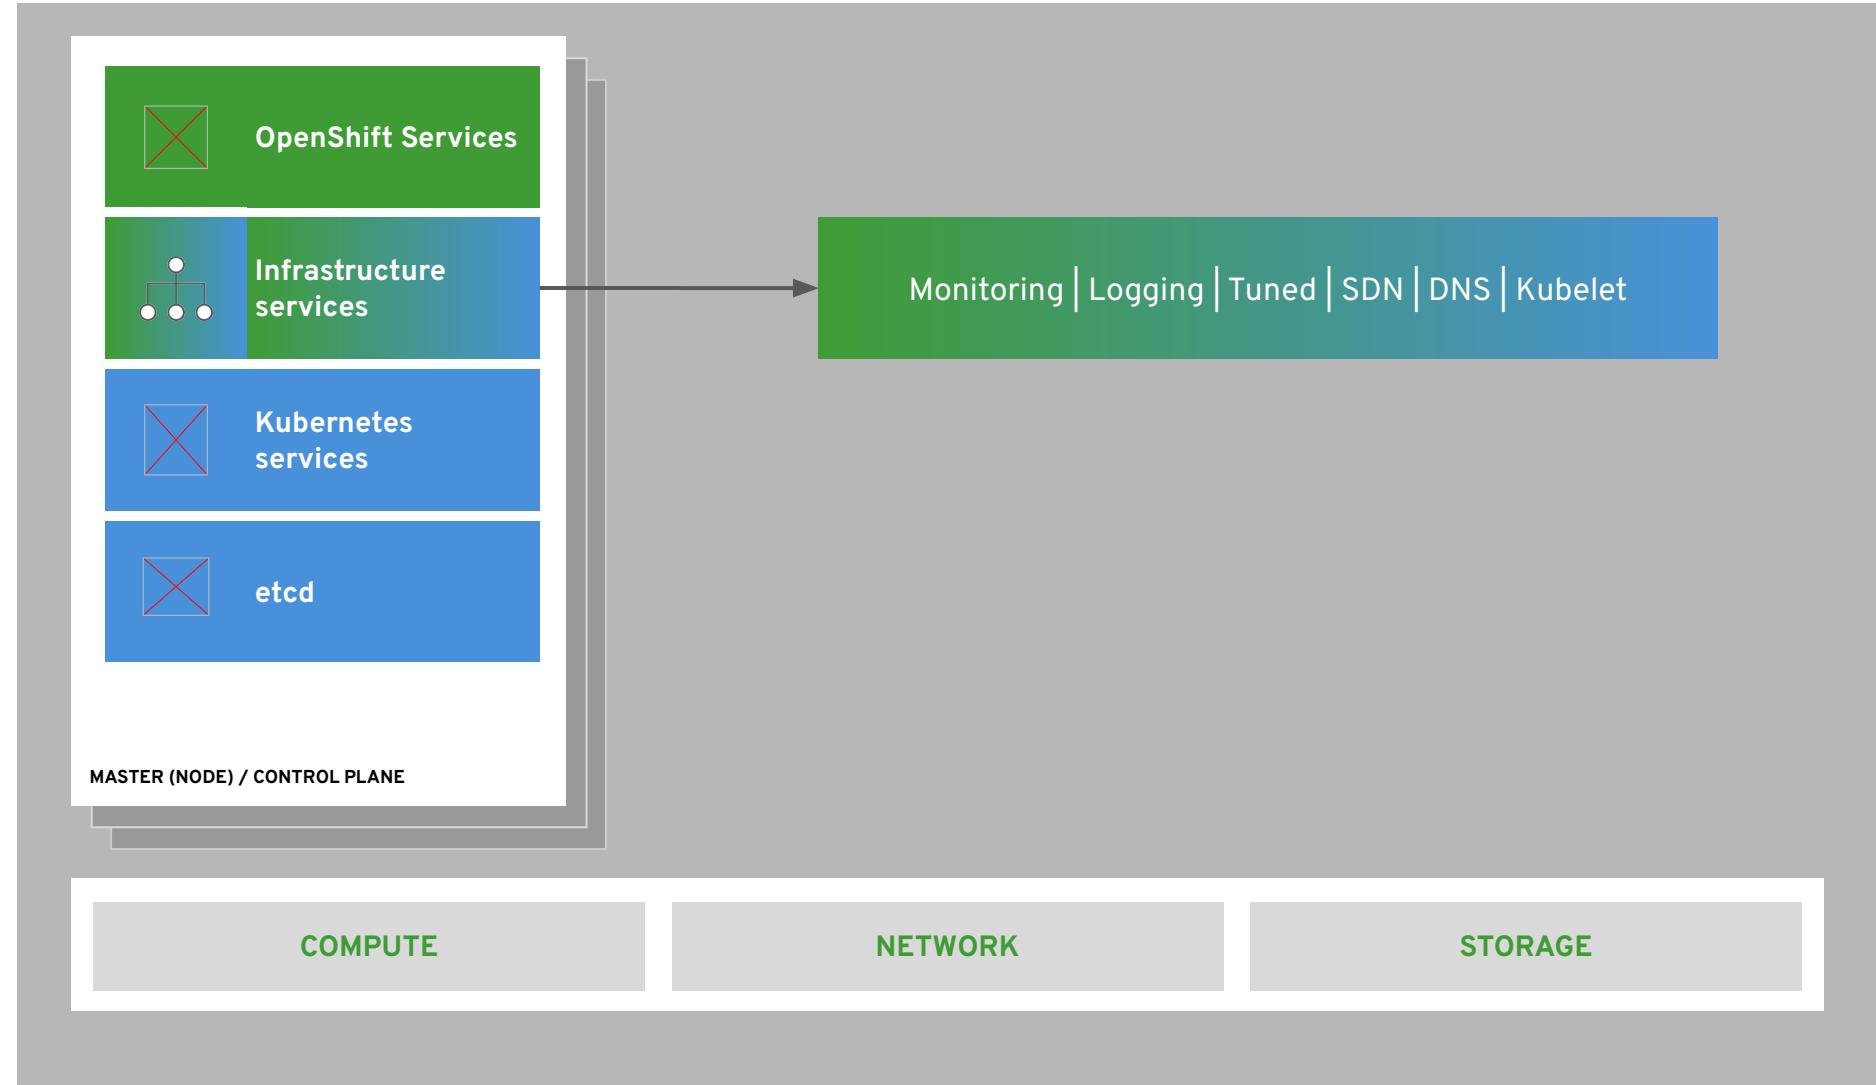
Describing OpenShift Extensions

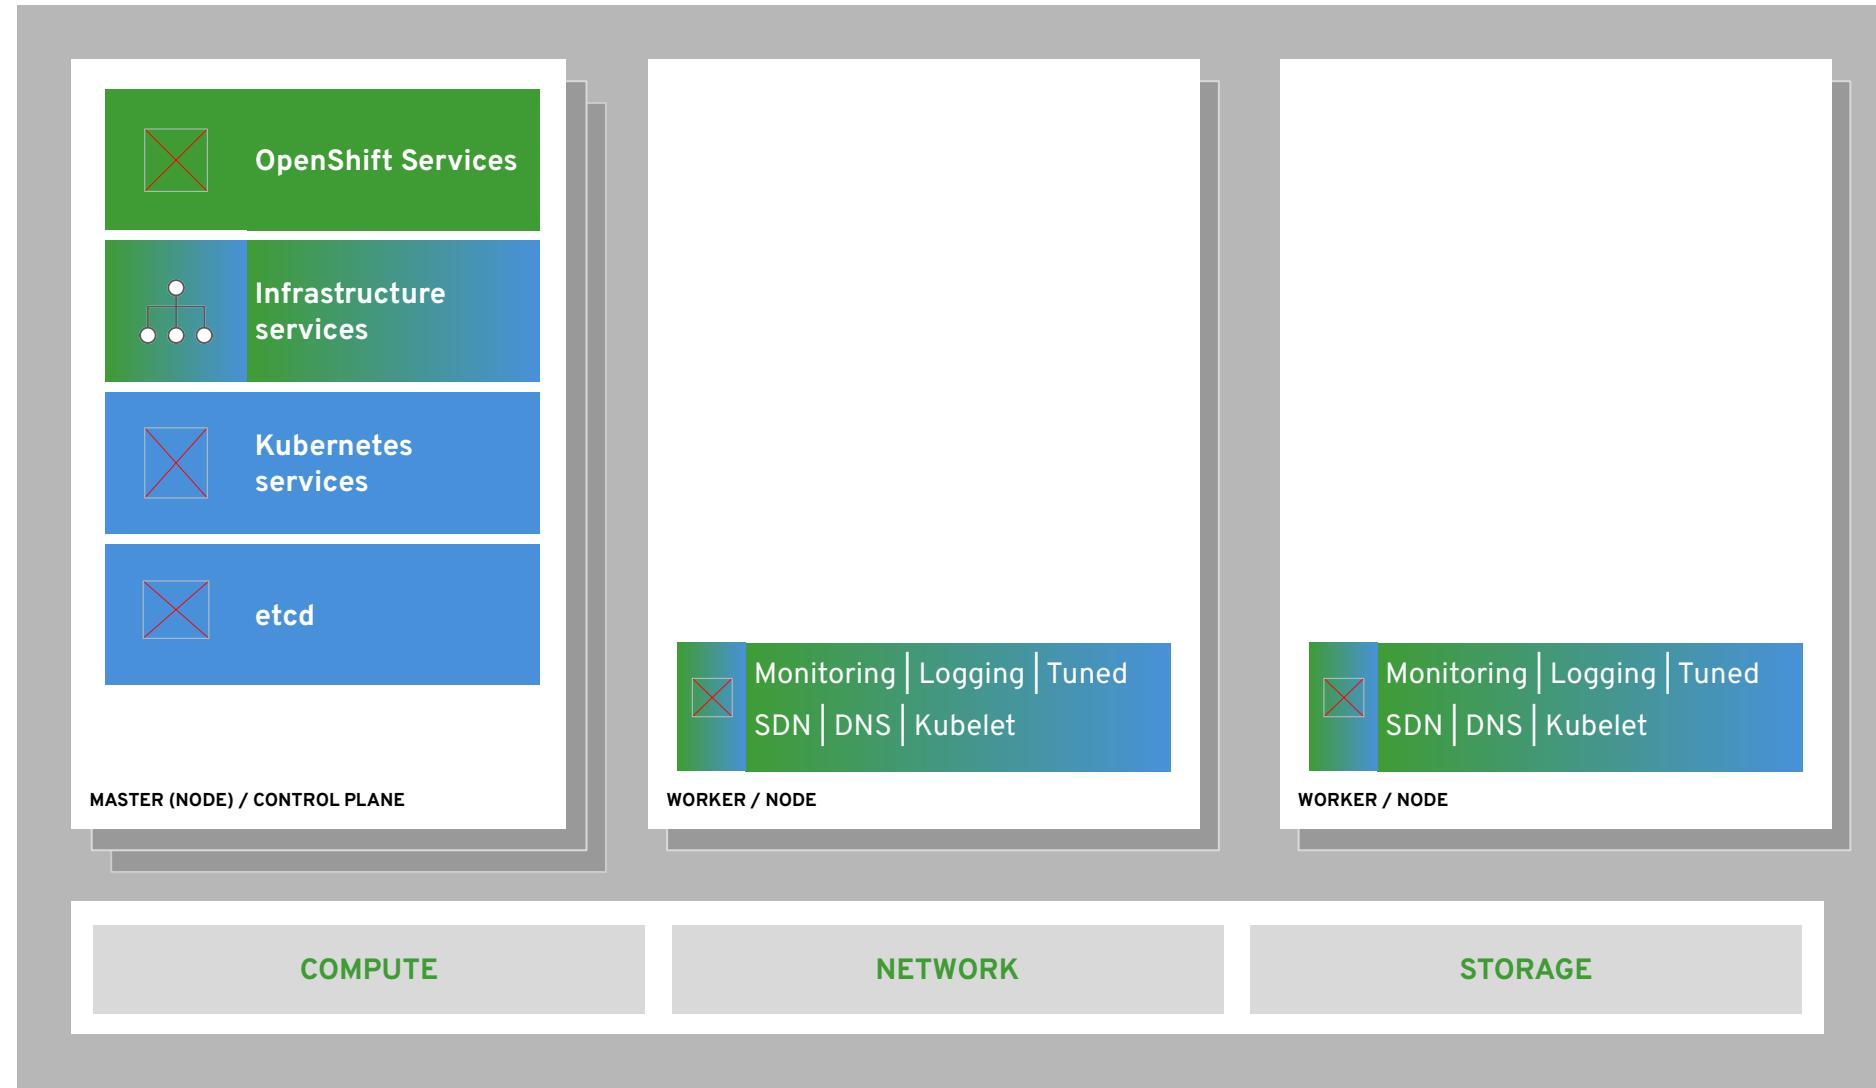
A lot of functionality from Kubernetes depends on external components, such as ingress controllers, storage plug-ins, network plug-ins, and authentication plug-ins. Similar to Linux distributions, there are many ways to build a Kubernetes distribution by picking and choosing different components.

A lot of functionality from Kubernetes also depends on extension APIs, such as access control and network isolation.

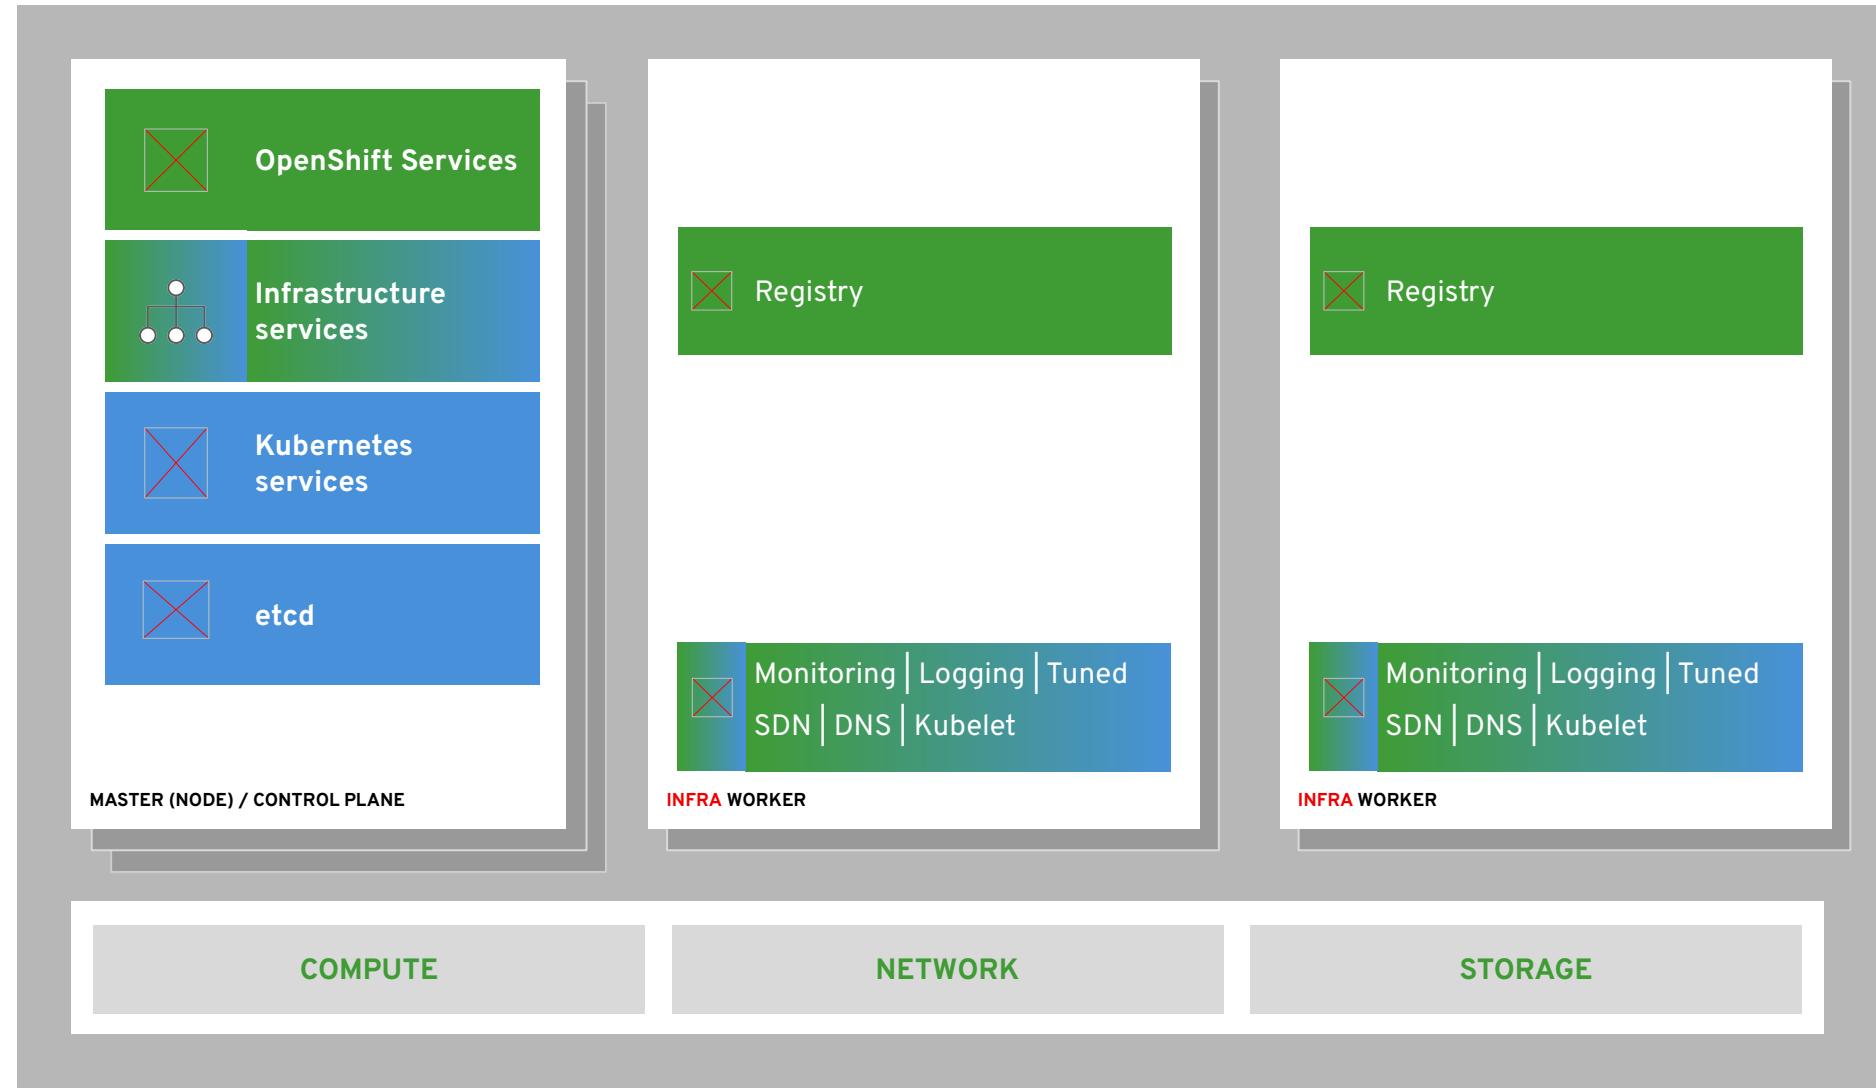
OpenShift is a Kubernetes distribution that provides many of these components already integrated and configured, and managed by operators. OpenShift also provides preinstalled applications, such as a container image registry and a web console, managed by operators.

internal and support infrastructure services

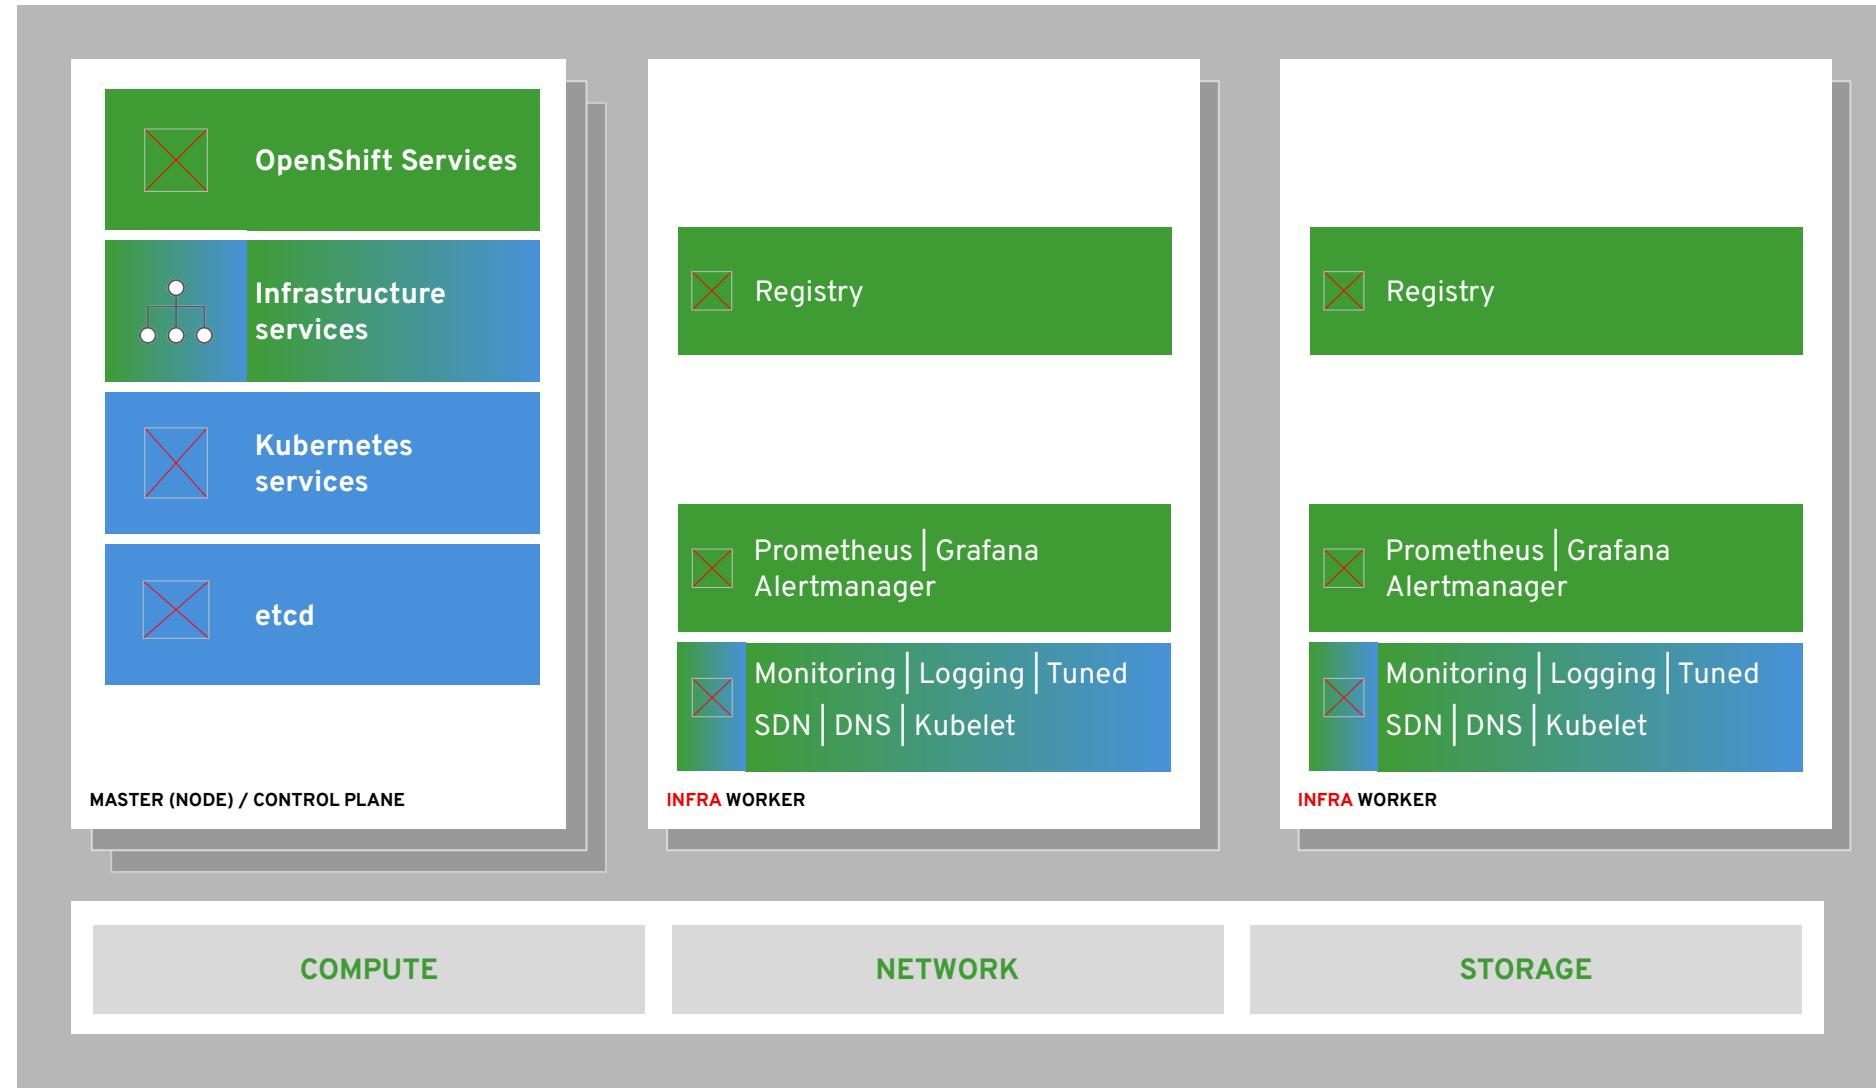




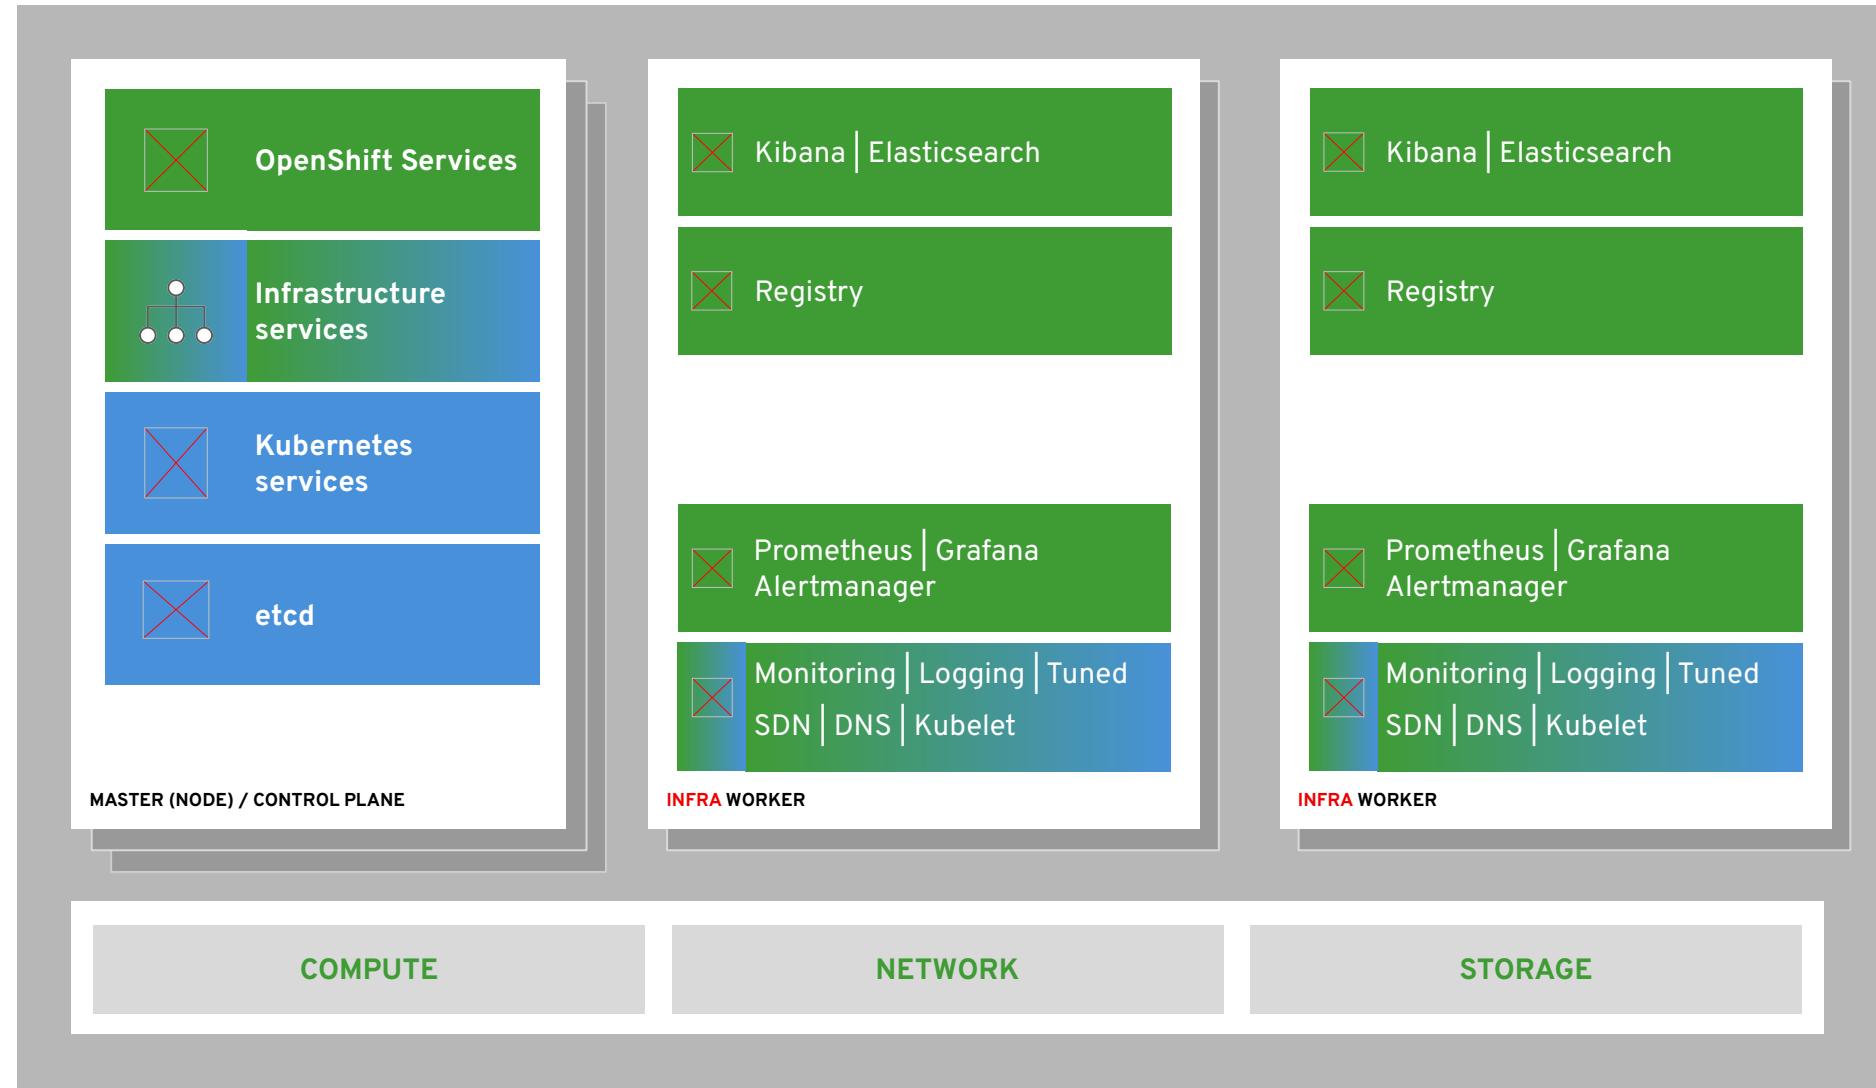
integrated image registry



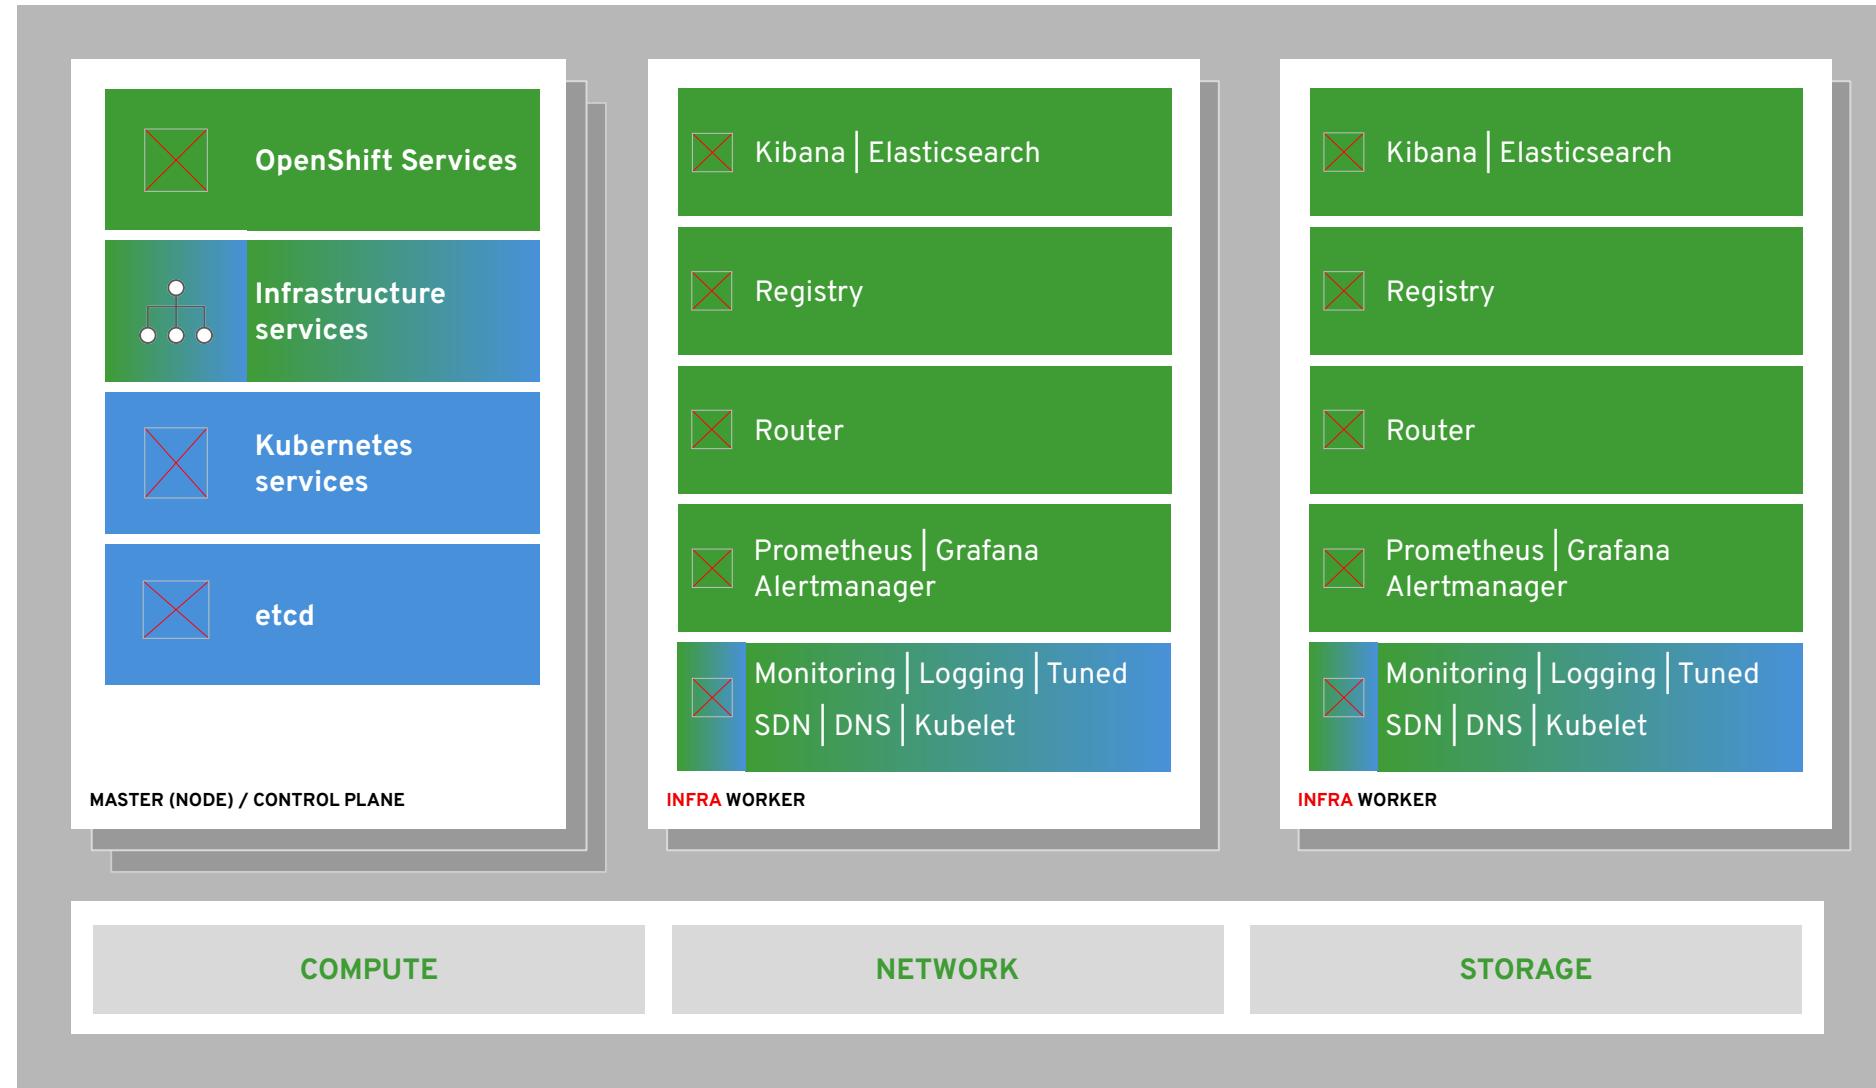
cluster monitoring

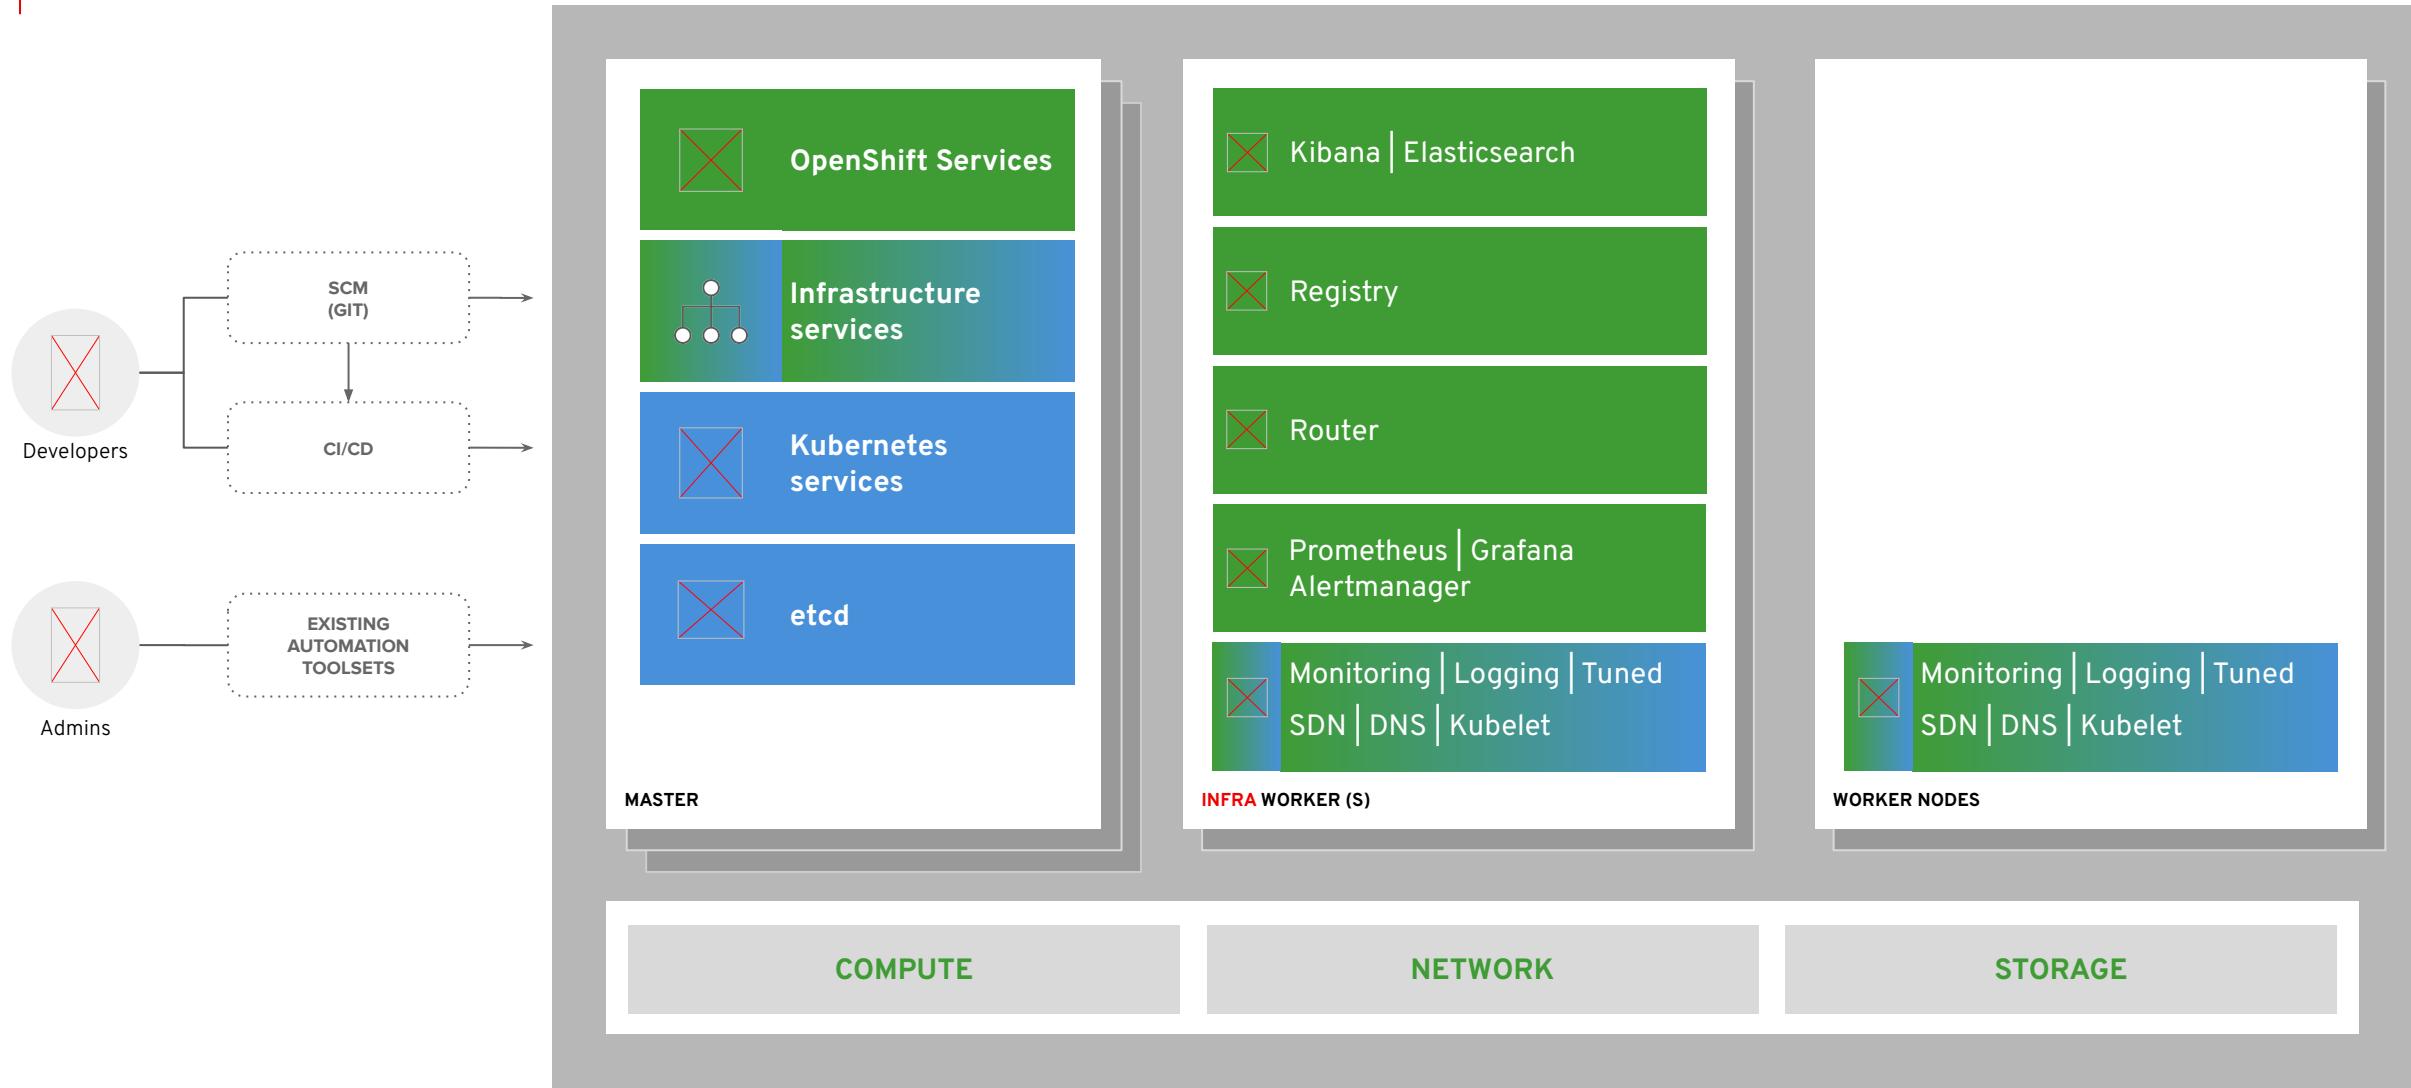


log aggregation



integrated routing





Quiz: Describing the Architecture of OpenShift

<https://rol.redhat.com/rol/app/courses/do280-4.10/pages/ch01>

1. OpenShift is based on which of the following container orchestration technologies?

- A Docker Swarm
- B Rancher
- C Kubernetes
- D Mesosphere Marathon
- E CoreOS Fleet

CHECK **RESET** **SHOW SOLUTION**

2. Which two of the following statements are true of OpenShift Container Platform? (Choose two.)

- A OpenShift provides an OAuth server that authenticates calls to its REST API.
- B OpenShift requires the CRI-O container engine.
- C Kubernetes follows a declarative architecture, but OpenShift follows a more traditional imperative architecture.
- D OpenShift extension APIs run as system services.

CHECK **RESET** **SHOW SOLUTION**

3. Which of the following servers runs Kubernetes API components?

- A Compute nodes
- B Nodes
- C Control plane nodes

CHECK **RESET** **SHOW SOLUTION**

4. Which of the following components does OpenShift add to upstream Kubernetes?

- A The etcd database
- B A container engine
- C A registry server
- D A scheduler
- E The Kubelet

CHECK **RESET** **SHOW SOLUTION**

5. Which of the following sentences is true regarding support for storage with OpenShift?

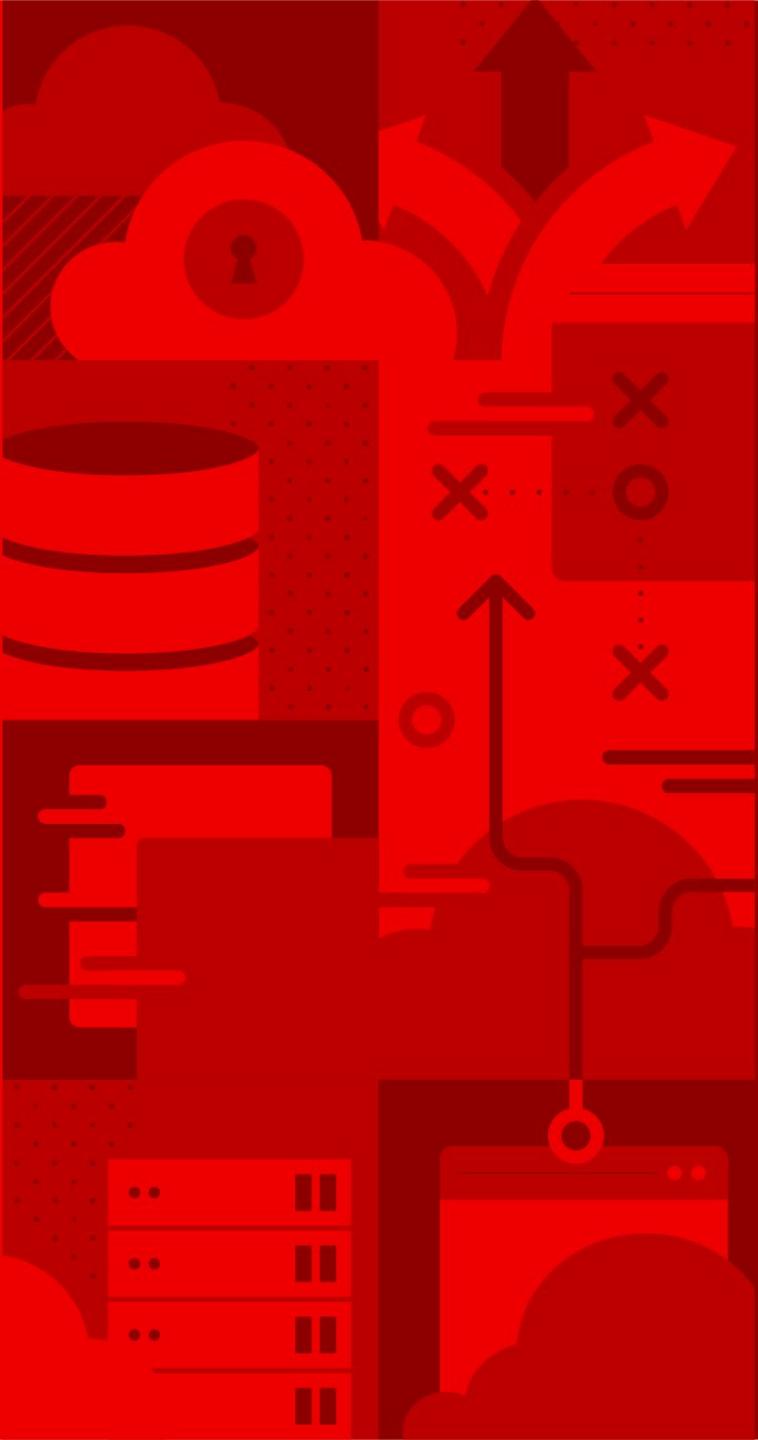
- A Users can only store persistent data in the etcd database.
- B Users can only deploy on OpenShift cloud-native applications that conform to the Twelve-Factor App methodology.
- C Administrators must configure storage plug-ins appropriate for their cloud providers.
- D Administrators must define persistent volumes before any user can deploy applications that require persistent storage.
- E Users can deploy applications that require persistent storage by relying on the default storage class.

CHECK **RESET** **SHOW SOLUTION**



Networking

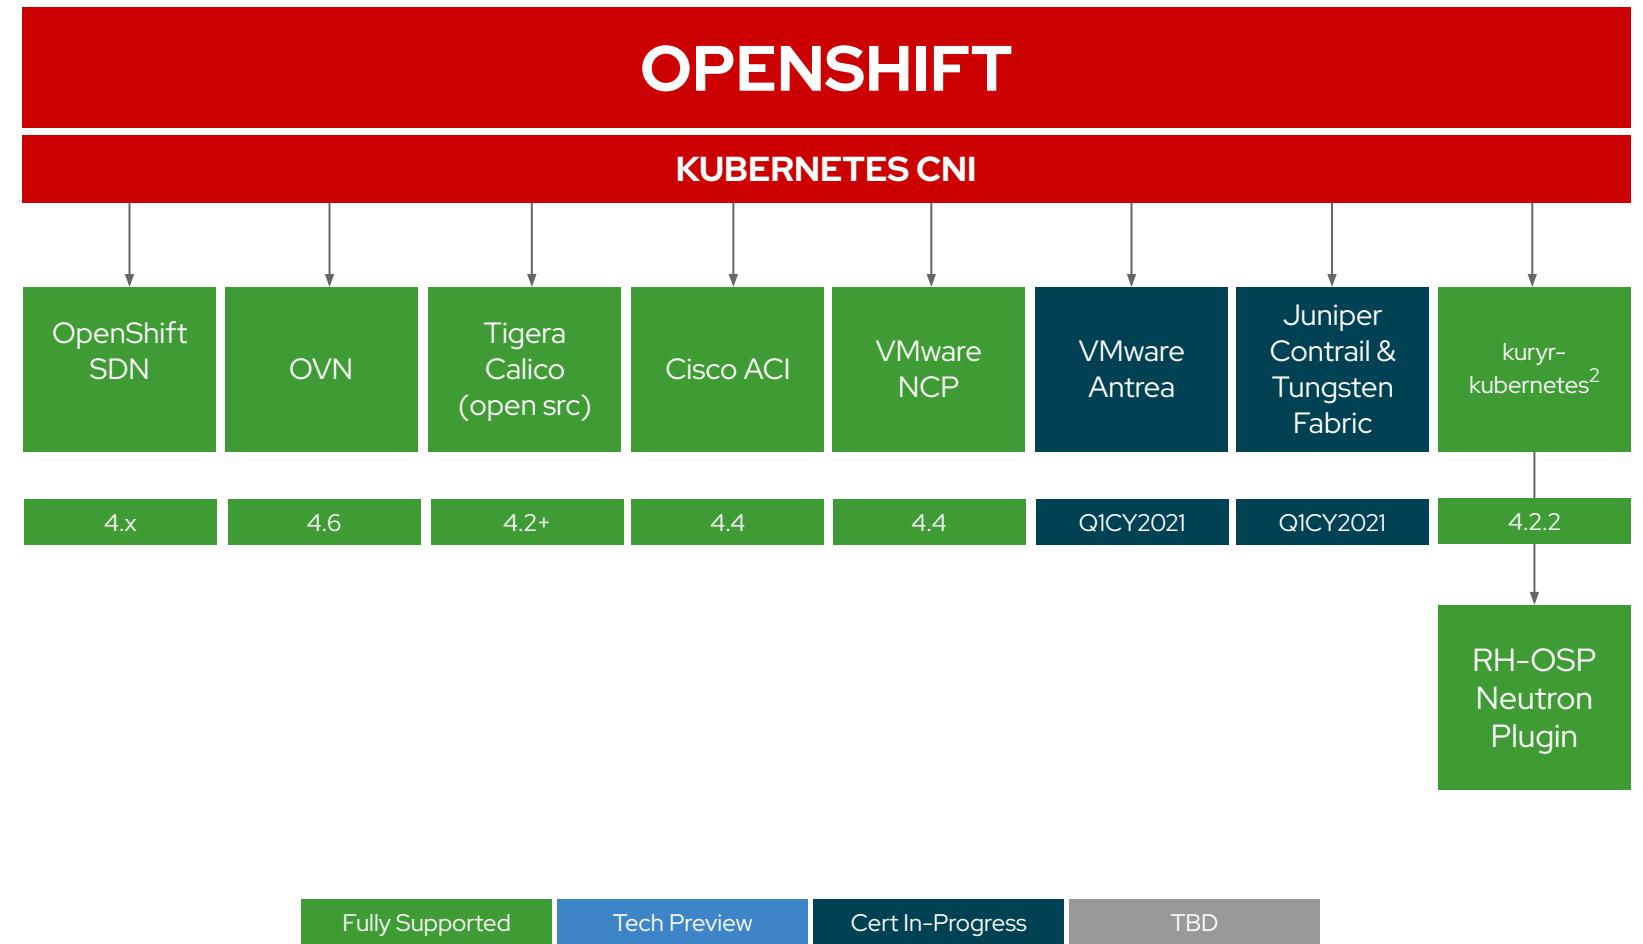
A pluggable model for network interface controls in kubernetes



OpenShift Networking Plug-ins

3rd-party Kubernetes CNI plug-in certification primarily consists of:

1. Formalizing the partnership
2. Certifying the container(s)
3. Certifying the Operator
4. Successfully passing the same Kubernetes networking conformance tests that OpenShift uses to validate its own SDN



OpenShift SDN

An Open Virtual Network OVN Software Defined Network for kubernetes

OpenShift implements a software-defined network (SDN) to manage the network infrastructure of the cluster and user applications. Software-defined networking is a networking model that allows you to manage network services through the abstraction of several networking layers. It decouples the software that handles the traffic, called the *control plane*, and the underlying mechanisms that route the traffic, called the *data plane*. Among the many features of SDN, open standards enable vendors to propose their solutions, centralized management, dynamic routing, and tenant isolation.

OpenShift SDN high-level architecture

OpenShift Architecture

CONFIDENTIAL designator

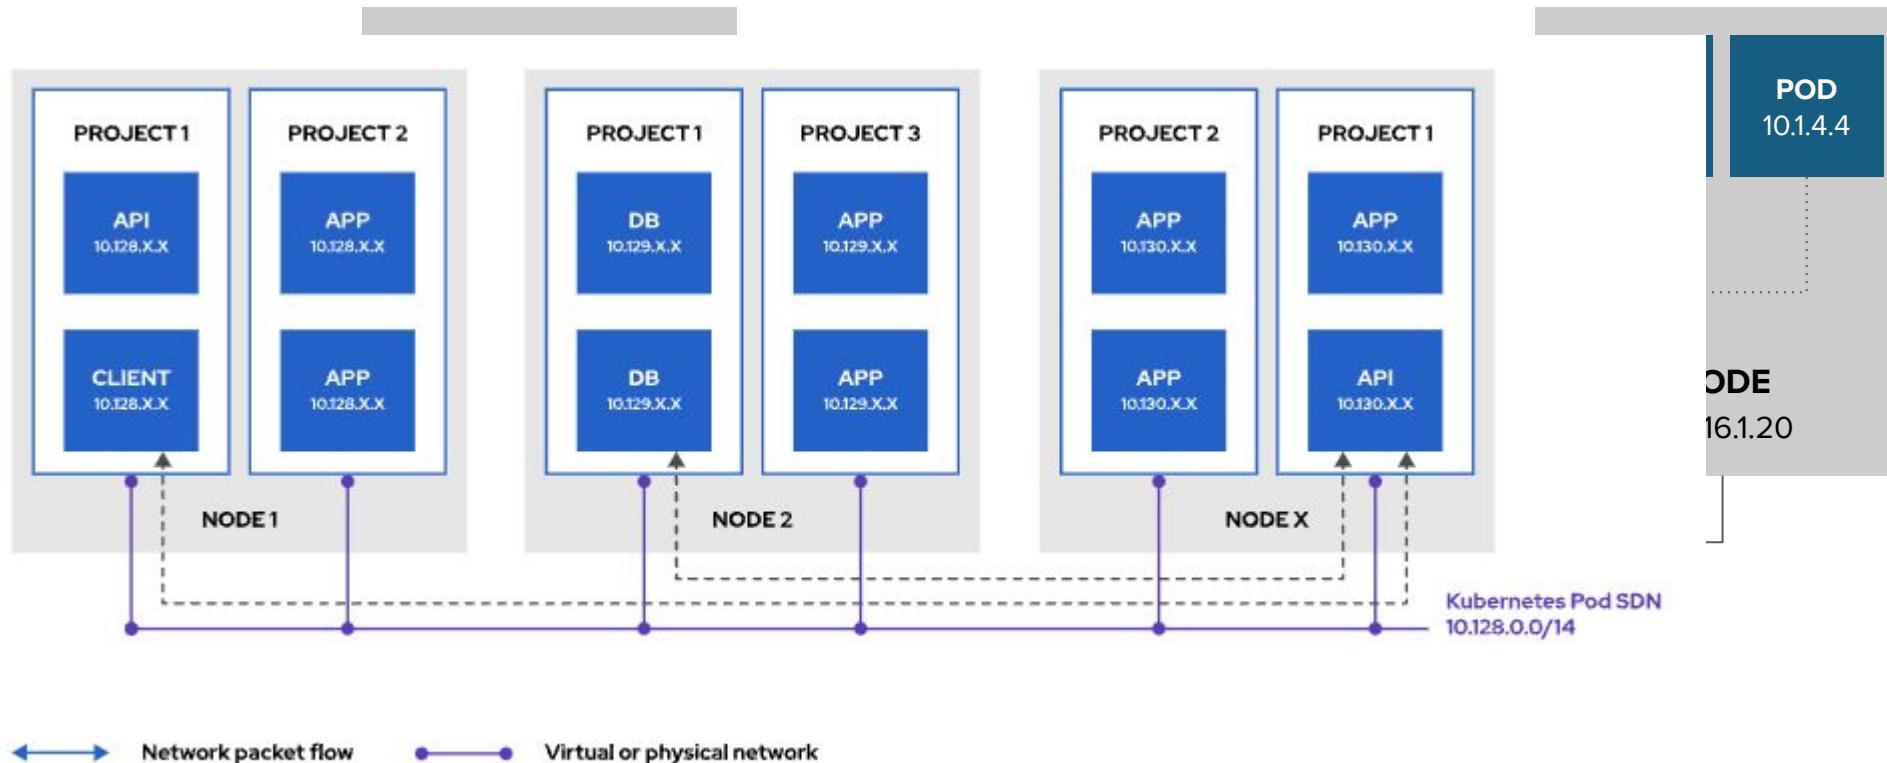
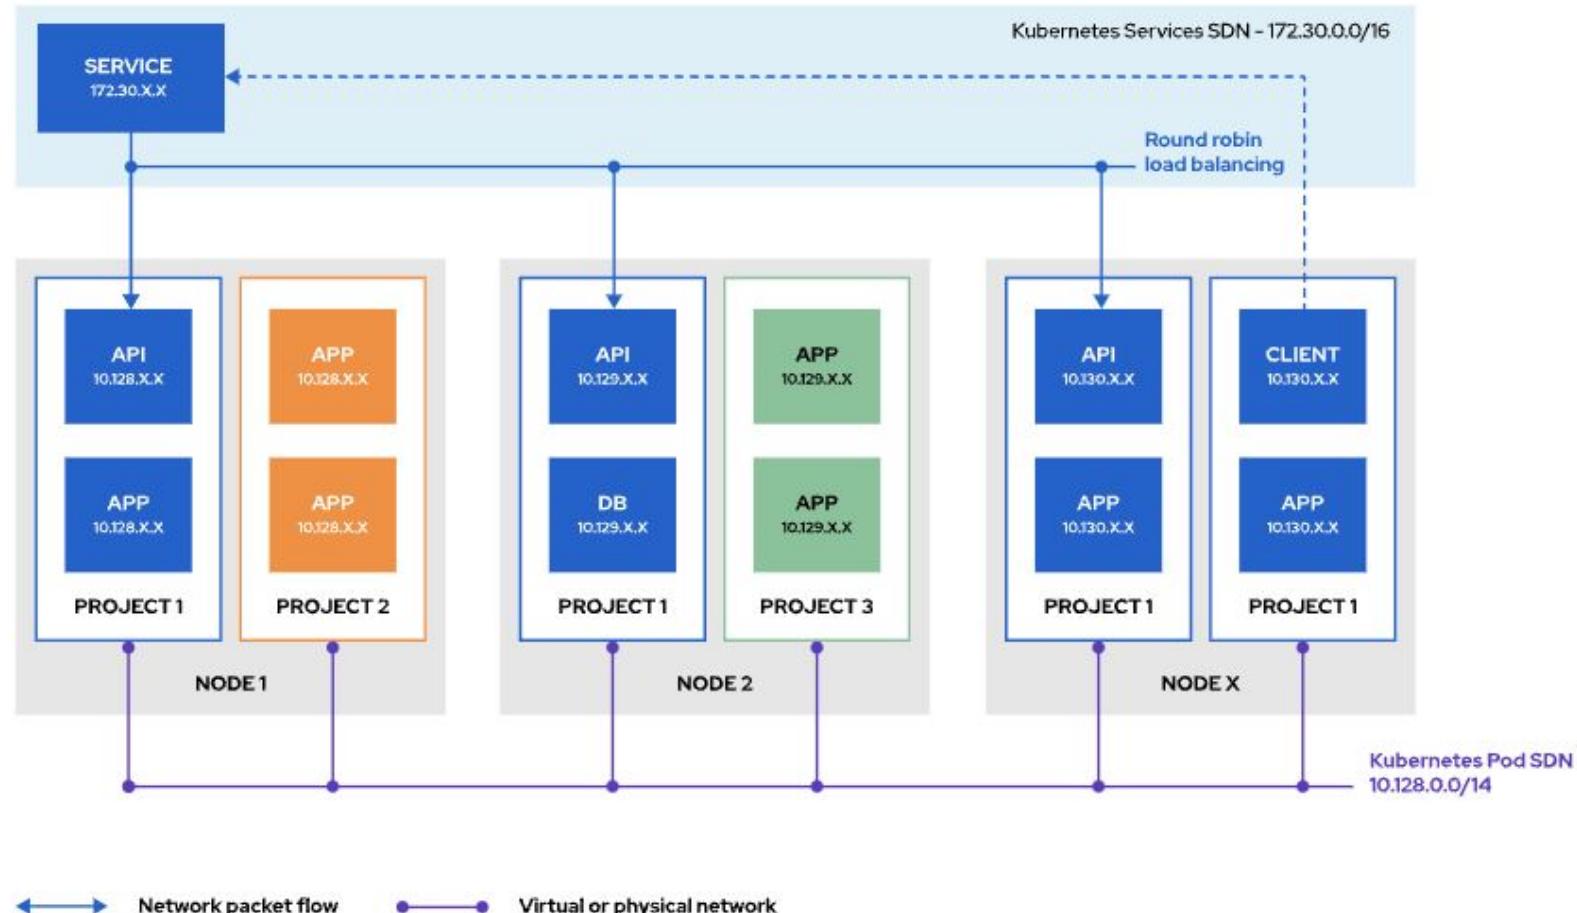


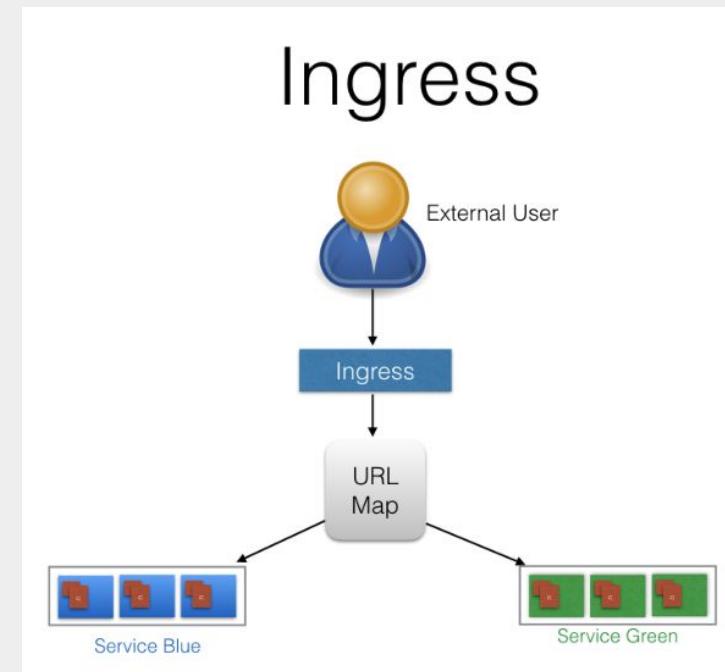
Figure 5.1: Kubernetes basic networking

Using Services for Accessing Pods



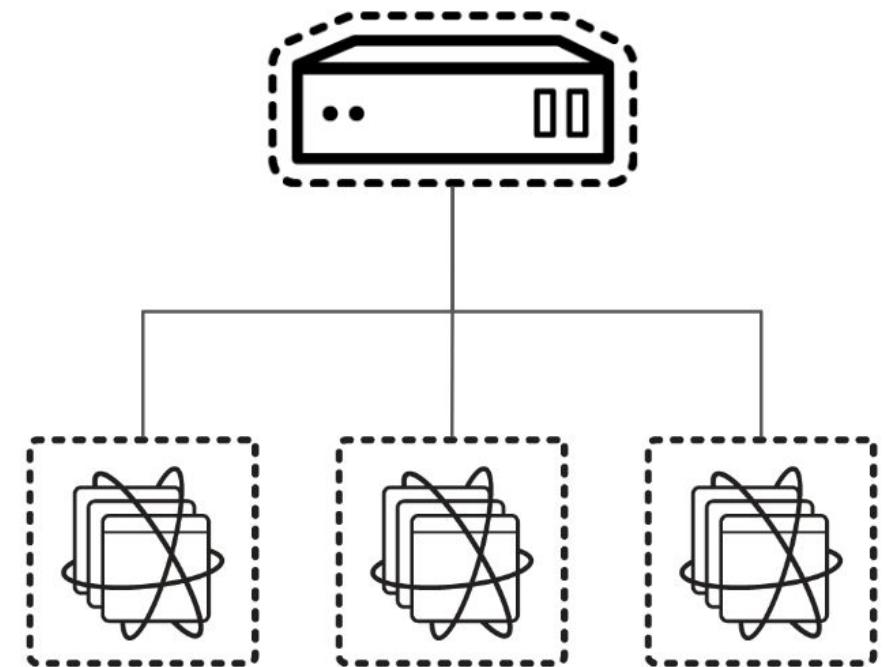
routes and ingress

How traffic enters the cluster



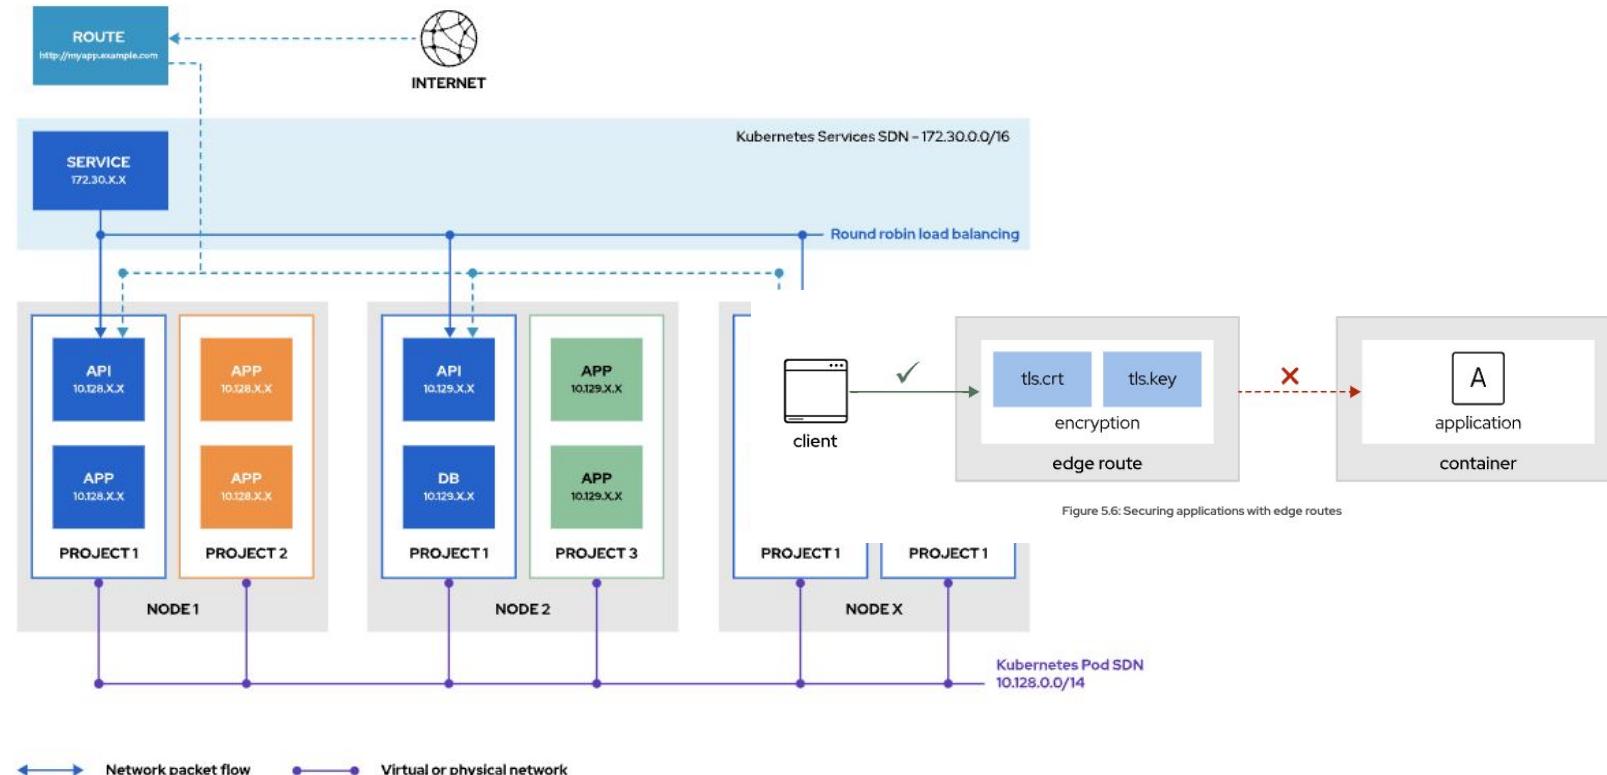
Routing and Load Balancing

- ▶ Pluggable routing architecture
 - HAProxy Router
 - F5 Router
- ▶ Multiple-routers with traffic sharding
- ▶ Router supported protocols
 - HTTP/HTTPS
 - WebSockets
 - TLS with SNI
- ▶ Non-standard ports via cloud load-balancers, external IP, and NodePort



Exposing Applications for External Access

Guided Exercise: Exposing Applications for External Access



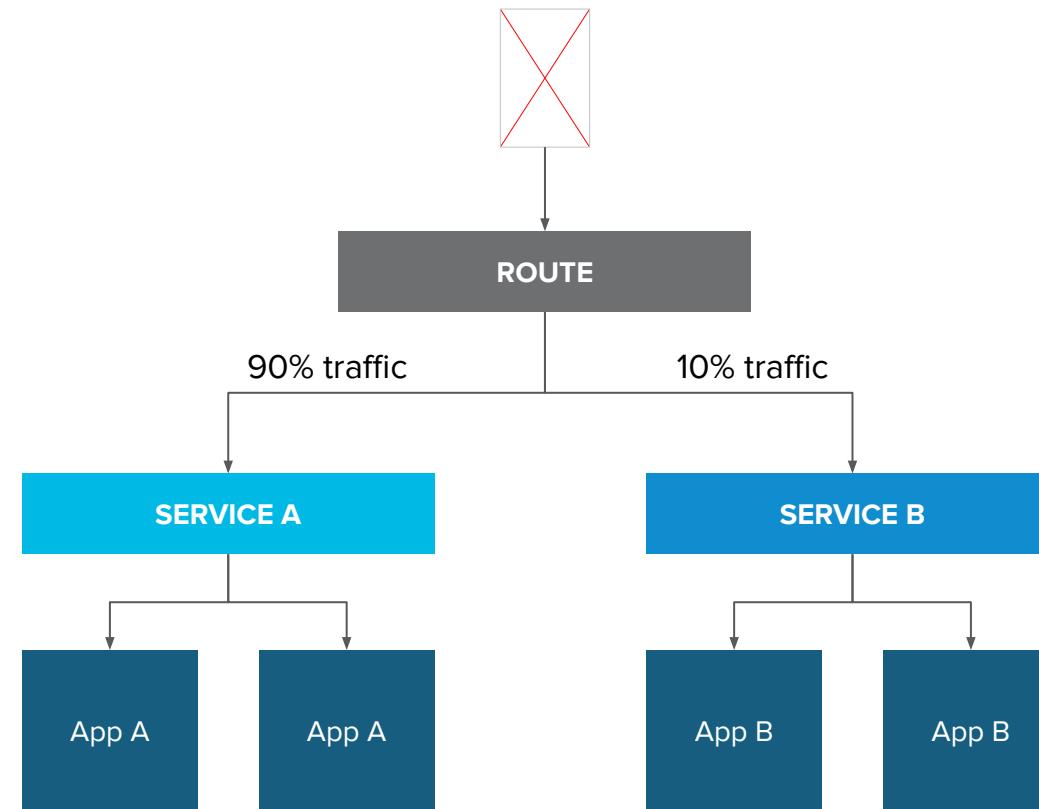
Routes vs Ingress

Feature	Ingress	Route
Standard Kubernetes object	X	
External access to services	X	X
Persistent (sticky) sessions	X	X
Load-balancing strategies (e.g. round robin)	X	X
Rate-limit and throttling	X	X
IP whitelisting	X	X
TLS edge termination	X	X
TLS re-encryption	X	X
TLS passthrough	X	X
Multiple weighted backends (split traffic)		X
Generated pattern-based hostnames		X
Wildcard domains		X



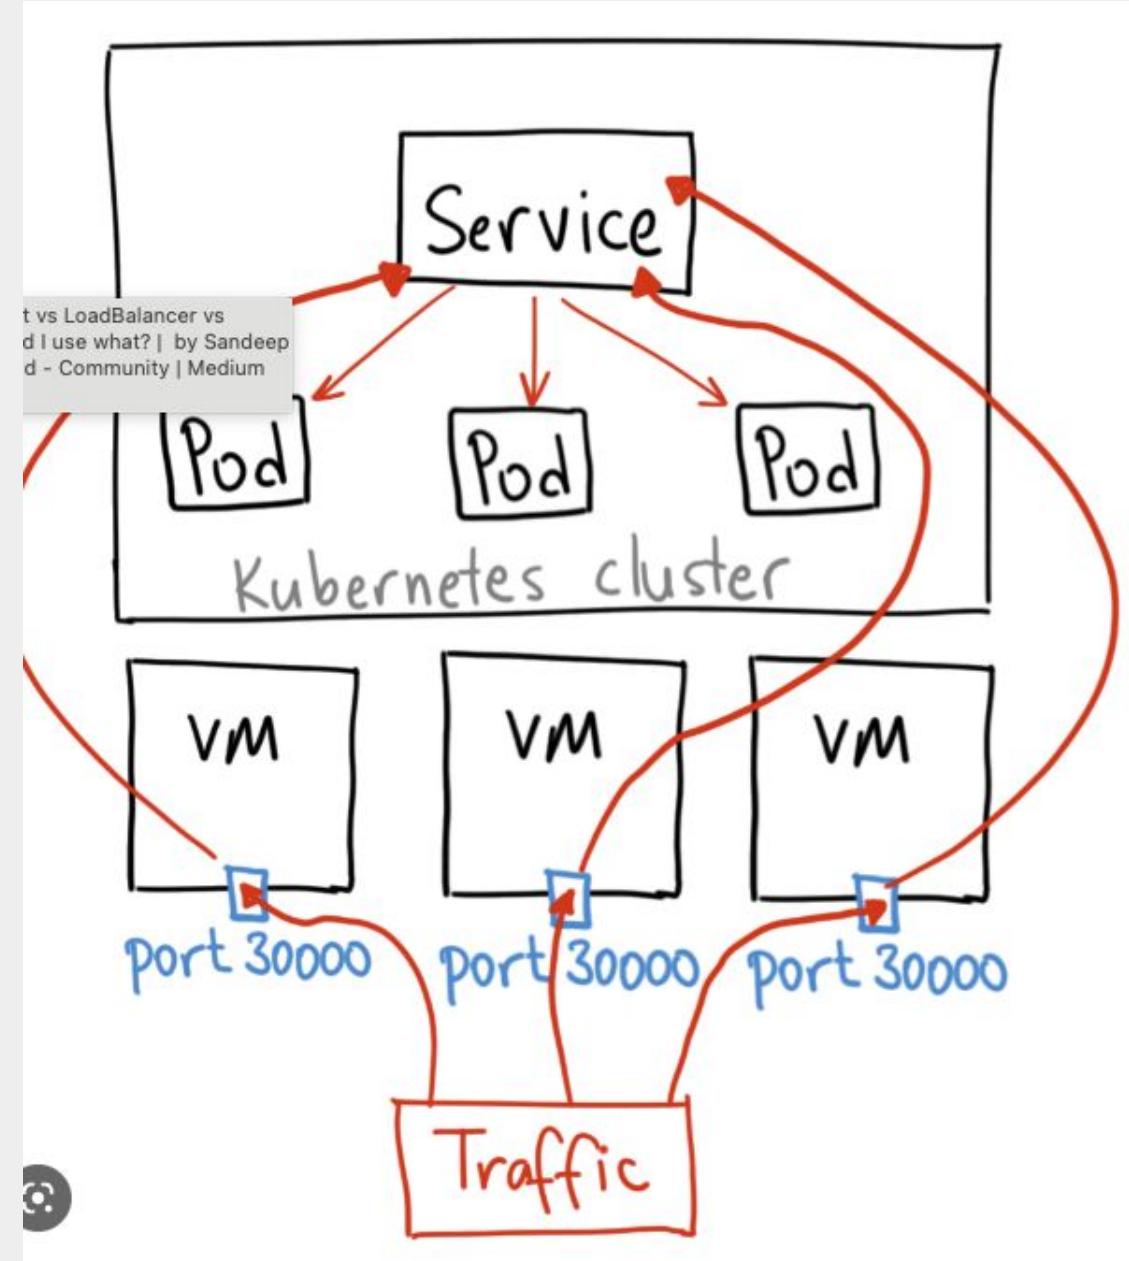
Router-based deployment methodologies

Split Traffic Between
Multiple Services For A/B
Testing, Blue/Green and
Canary Deployments



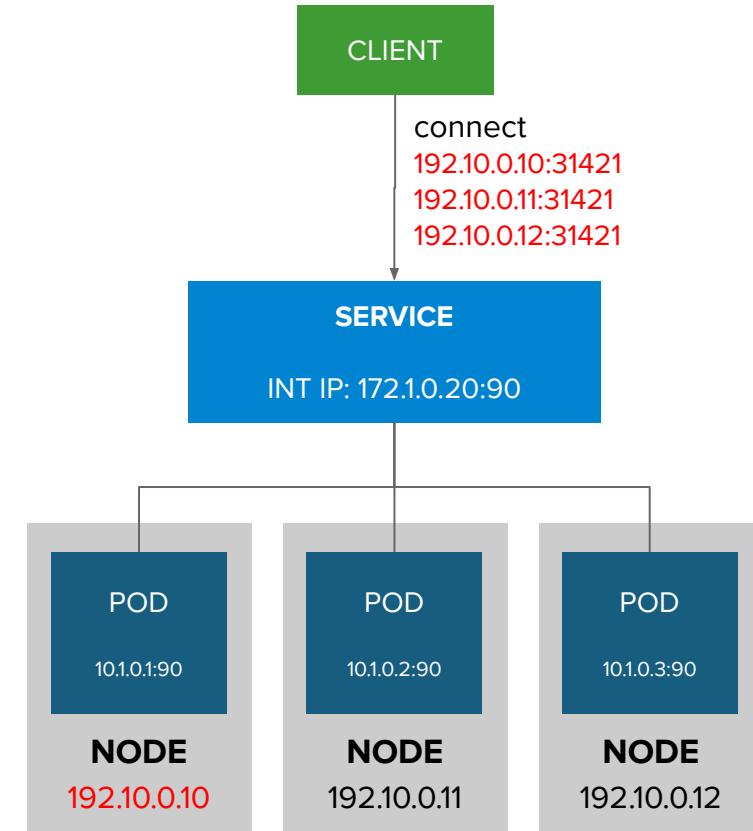
Alternative methods for ingress

Different ways that traffic can enter the cluster without the router



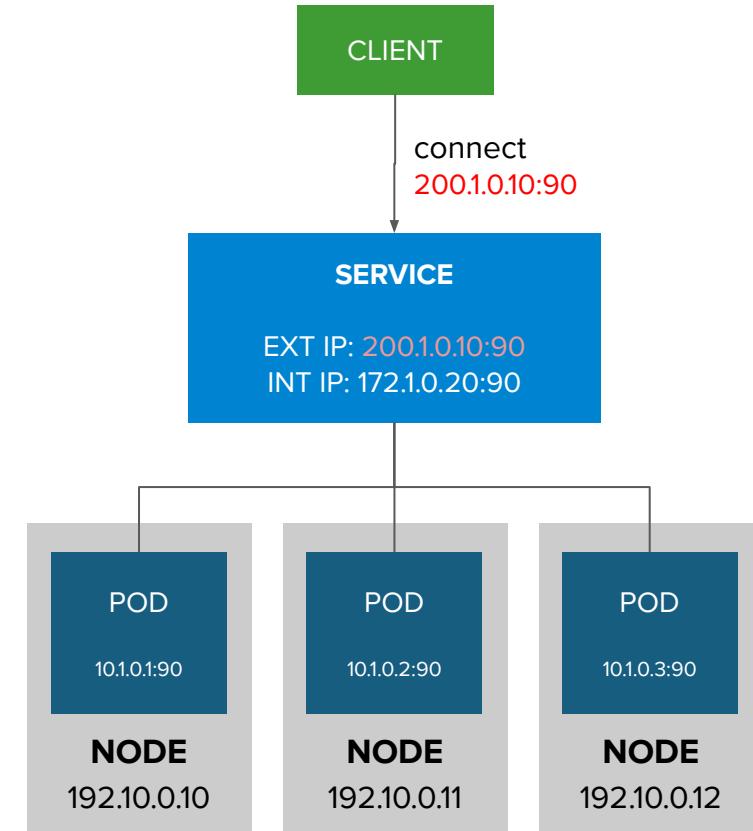
Entering the cluster on a random port with service nodeports

- ▶ NodePort binds a service to a unique port on all the nodes
- ▶ Traffic received on any node redirects to a node with the running service
- ▶ Ports in 30K-60K range which usually differs from the service
- ▶ Firewall rules must allow traffic to all nodes on the specific port



External traffic to a service on any port with external IP

- ▶ Access a service with an external IP on any TCP/UDP port, such as
 - Databases
 - Message Brokers
- ▶ Automatic IP allocation from a predefined pool using Ingress IP Self-Service
- ▶ IP failover pods provide high availability for the IP pool (fully supported in 4.8)

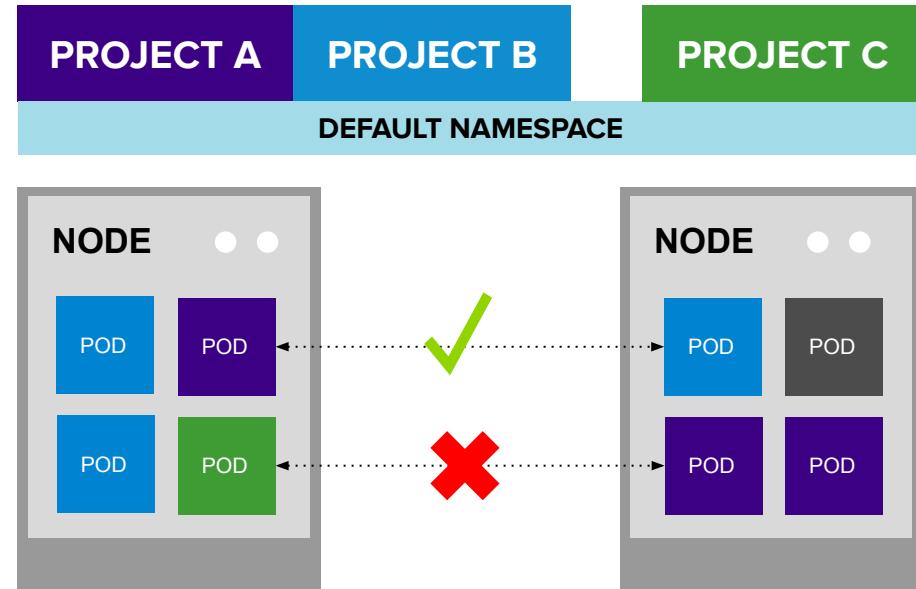


OPEN NETWORK (Default)

- ▶ All pods can communicate with each other across projects

MULTI-TENANT NETWORK

- ▶ Project-level network isolation
- ▶ Multicast support
- ▶ Egress network policies



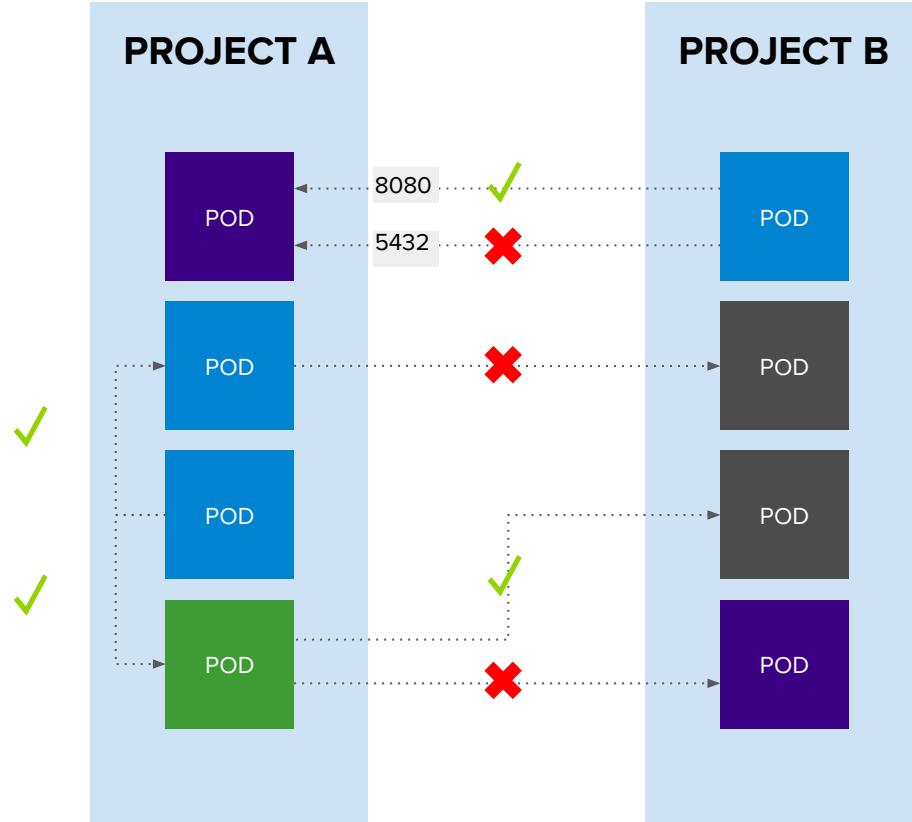
Multi-Tenant Network

Managing Network Policies in OpenShift

Network policies allow you to configure isolation policies for individual pods. Network policies do not require administrative privileges, giving developers more control over the applications in their projects. You can use network policies to create logical zones in the SDN that map to your organization network zones. The benefit of this approach is that the location of running pods becomes irrelevant because network policies allow you to segregate traffic regardless of where it originates.

The screenshot shows the Red Hat OpenShift Container Platform web interface. The left sidebar menu is visible, with 'Networking' expanded and 'Network Policies' selected. The main content area is titled 'Create Network Policy' and includes fields for 'Policy name' (set to 'my-policy') and 'Policy namespace' (set to 'my-project'). It also features a 'PodSelector' section with a radio button for 'Add pod selector' and a 'Policy type' section with a 'Deny' icon. Below these are sections for 'Ingress' and 'Egress' rules, each with a 'Remove all' button and an 'Add rule' button. A 'Create' button is at the bottom. On the right side, there is a sidebar titled 'Network Policy' with tabs for 'Schema' and 'Samples', and a detailed description of the NetworkPolicy object.

NetworkPolicy



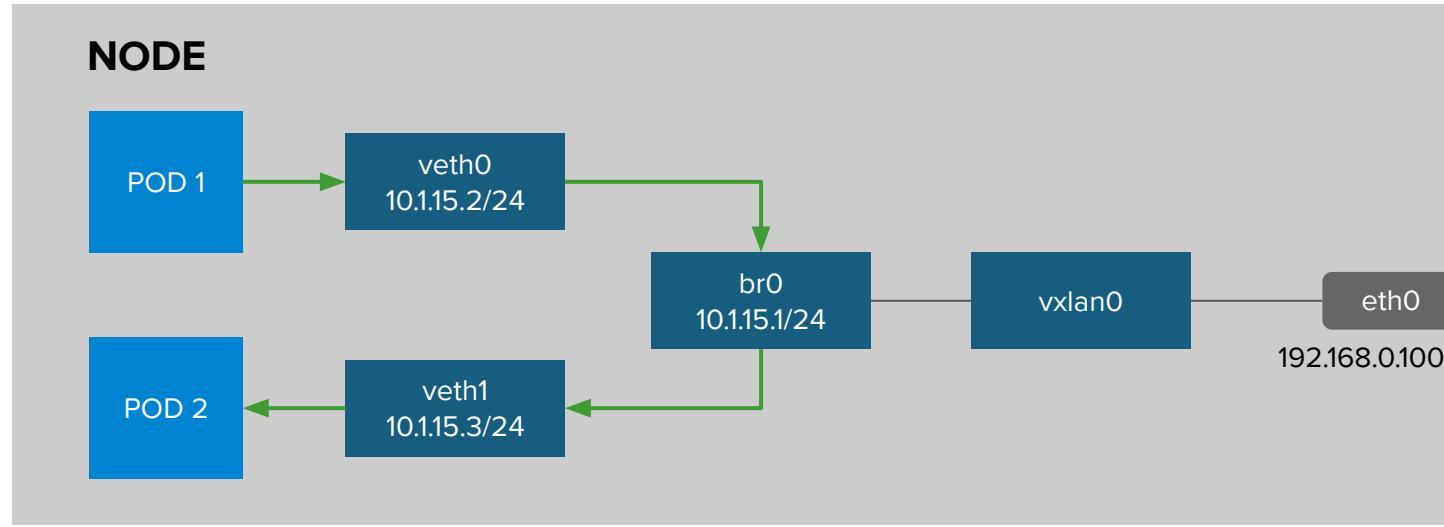
Example Policies

- Allow all traffic inside the project
- Allow traffic from green to gray
- Allow traffic to purple on 8080

```
apiVersion: extensions/v1beta1
kind: NetworkPolicy
metadata:
  name: allow-to-purple-on-8080
spec:
  podSelector:
    matchLabels:
      color: purple
  ingress:
  - ports:
    - protocol: tcp
      port: 8080
```

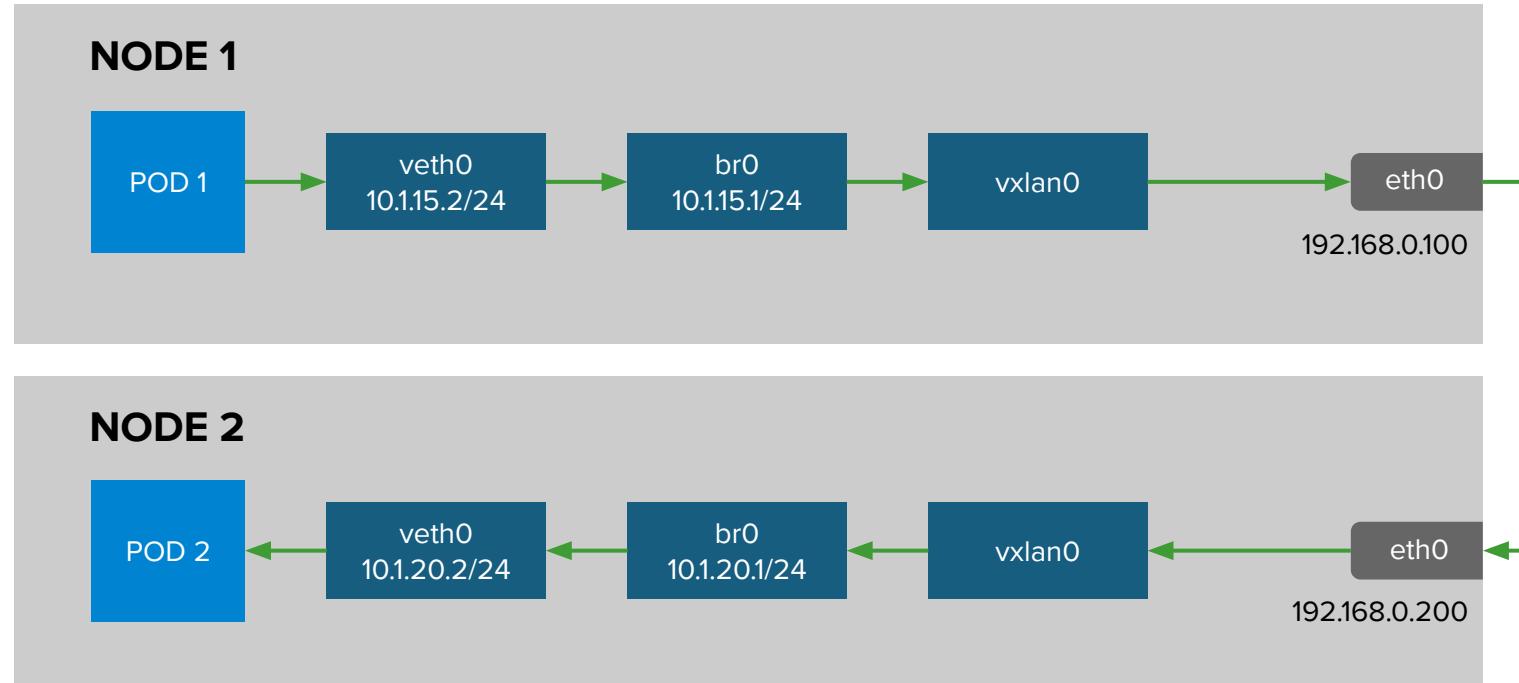
OpenShift SDN packet flows

container-container on same host



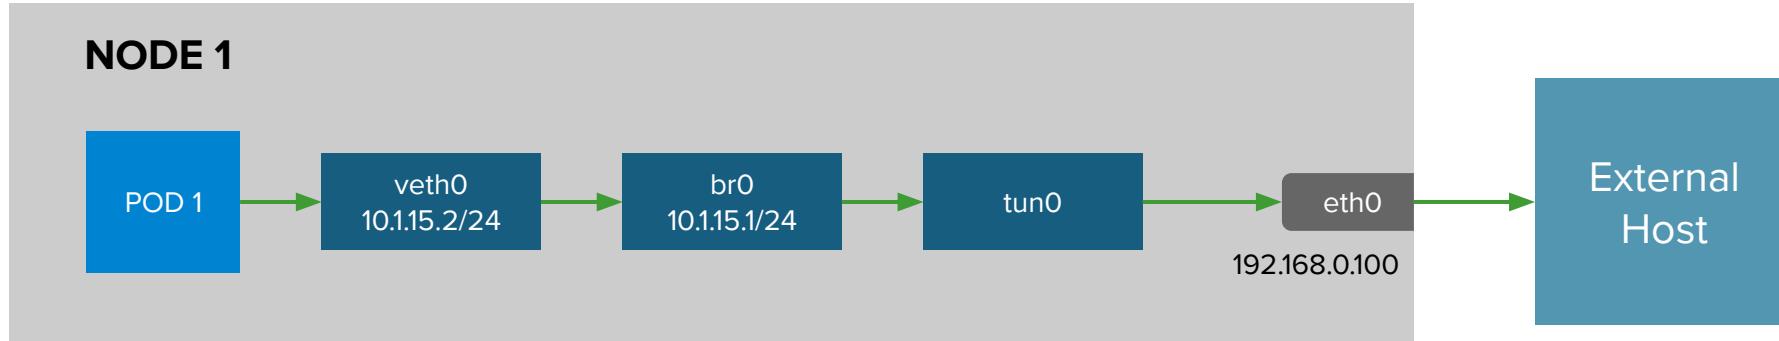
OpenShift SDN packet flows

container-container across hosts



OpenShift SDN packet flows

container leaving the host



Cluster DNS

An automated system for
providing hostname resolution
within kubernetes



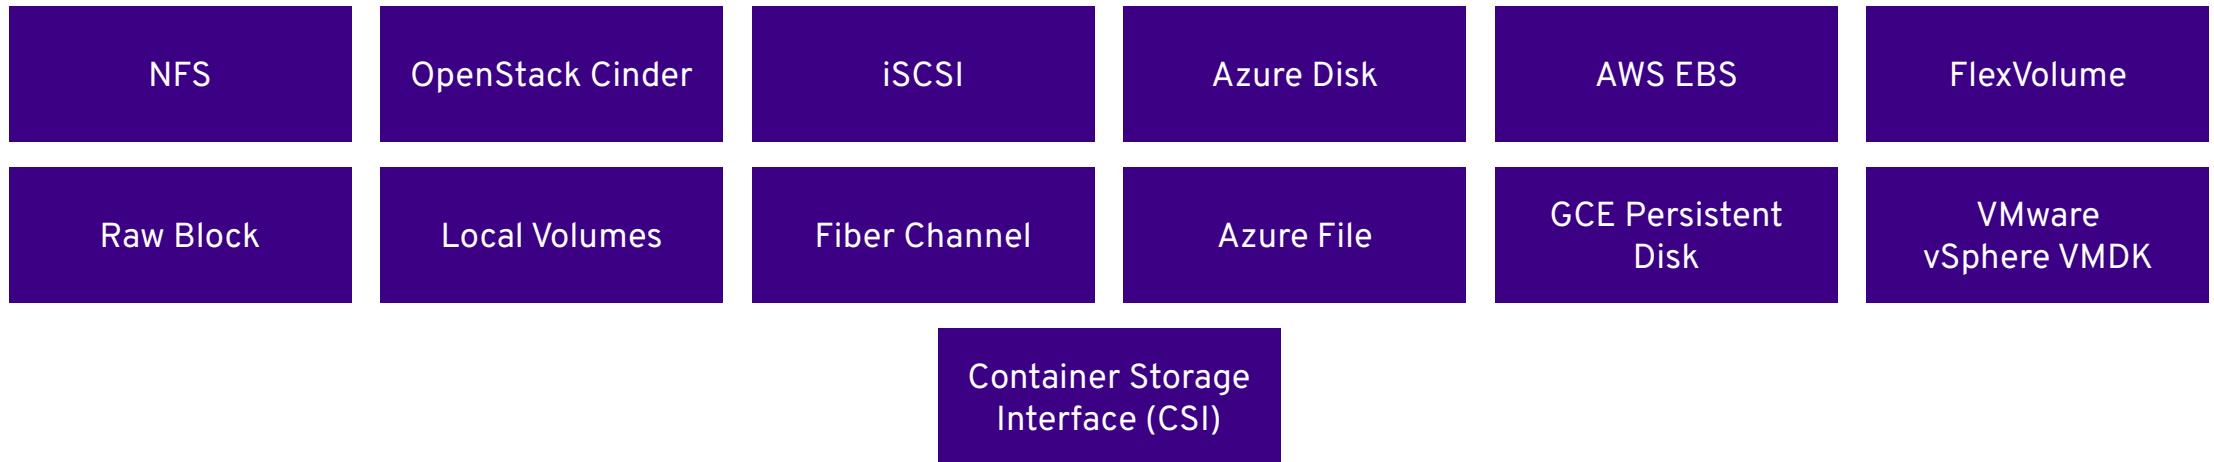
- ▶ Built-in internal DNS to reach services by a (fully qualified) hostname
- ▶ Split DNS is used with CoreDNS
 - CoreDNS answers DNS queries for internal/cluster services
 - Other defined “upstream” name servers serve the rest of the queries



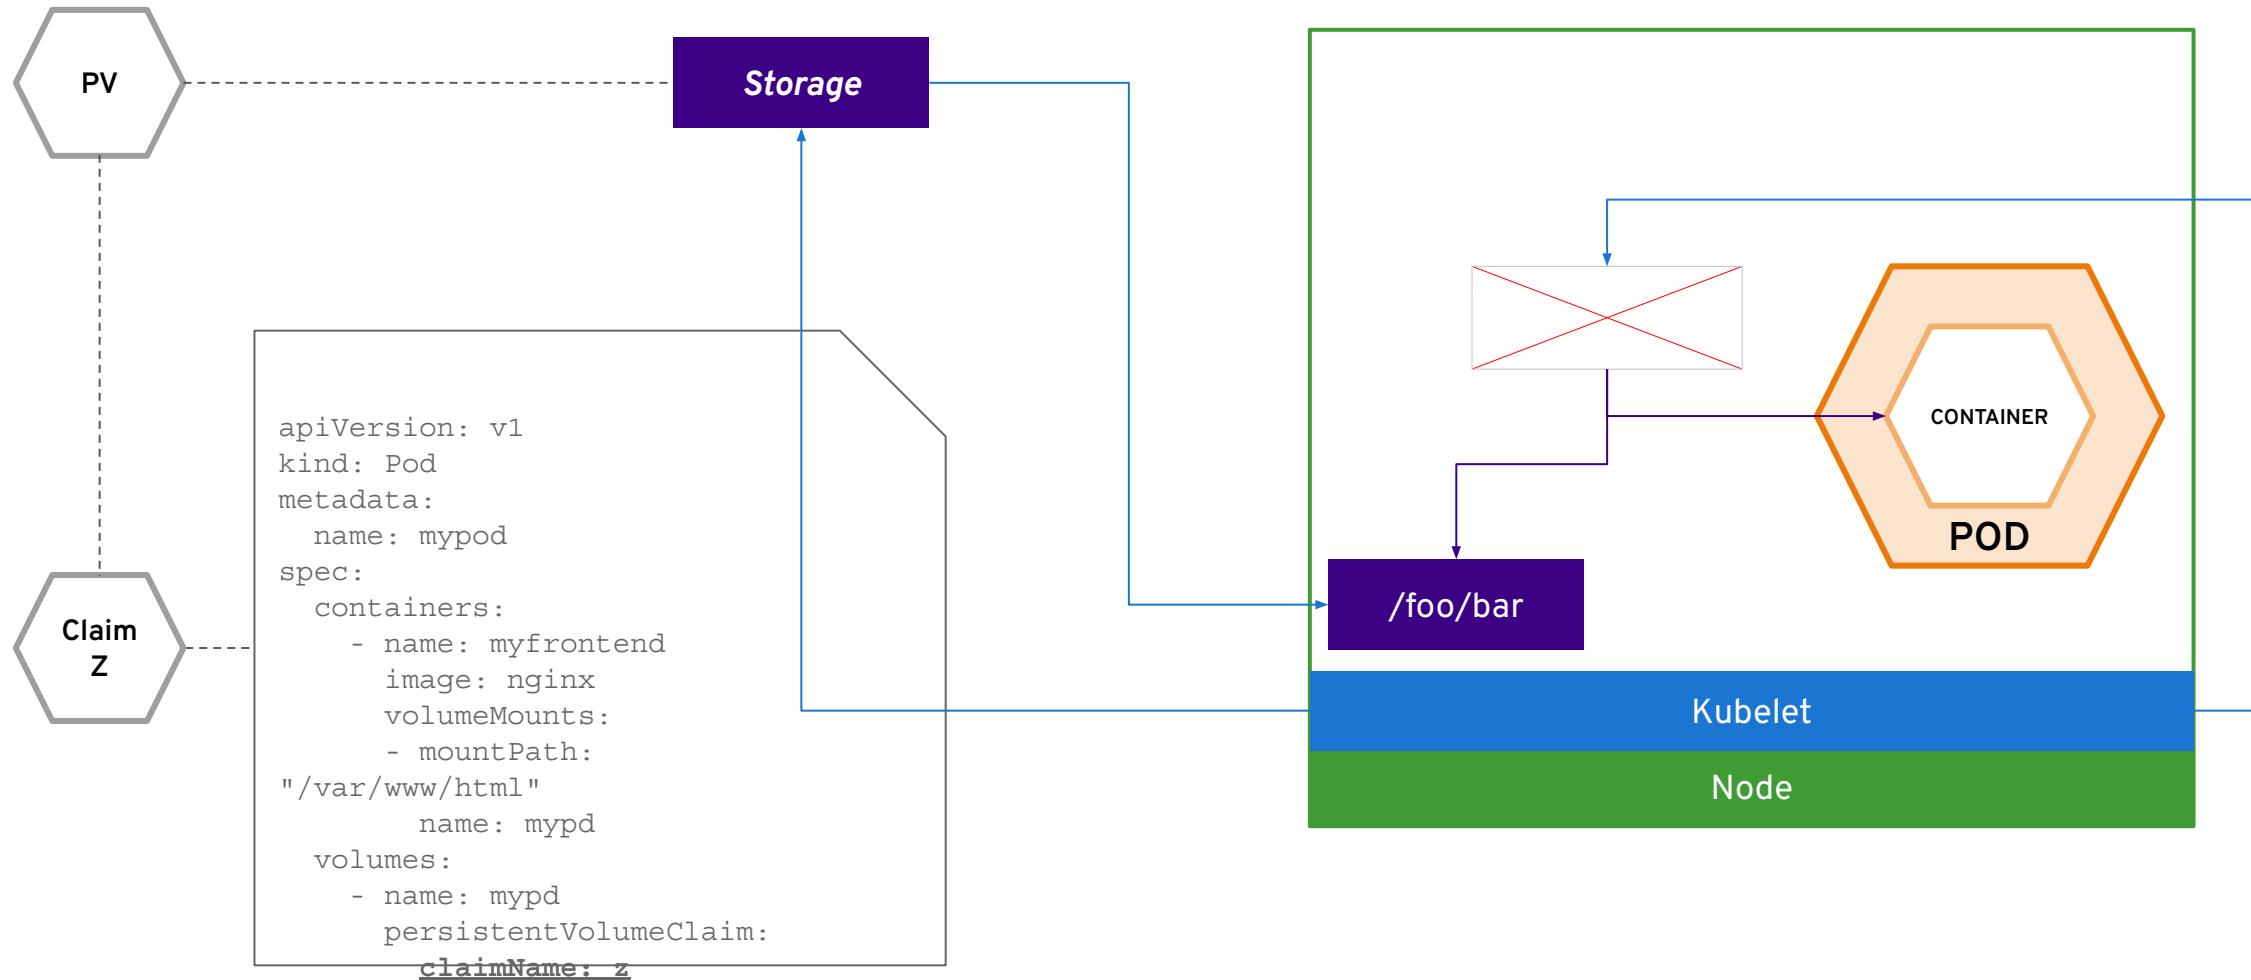


Persistent Storage

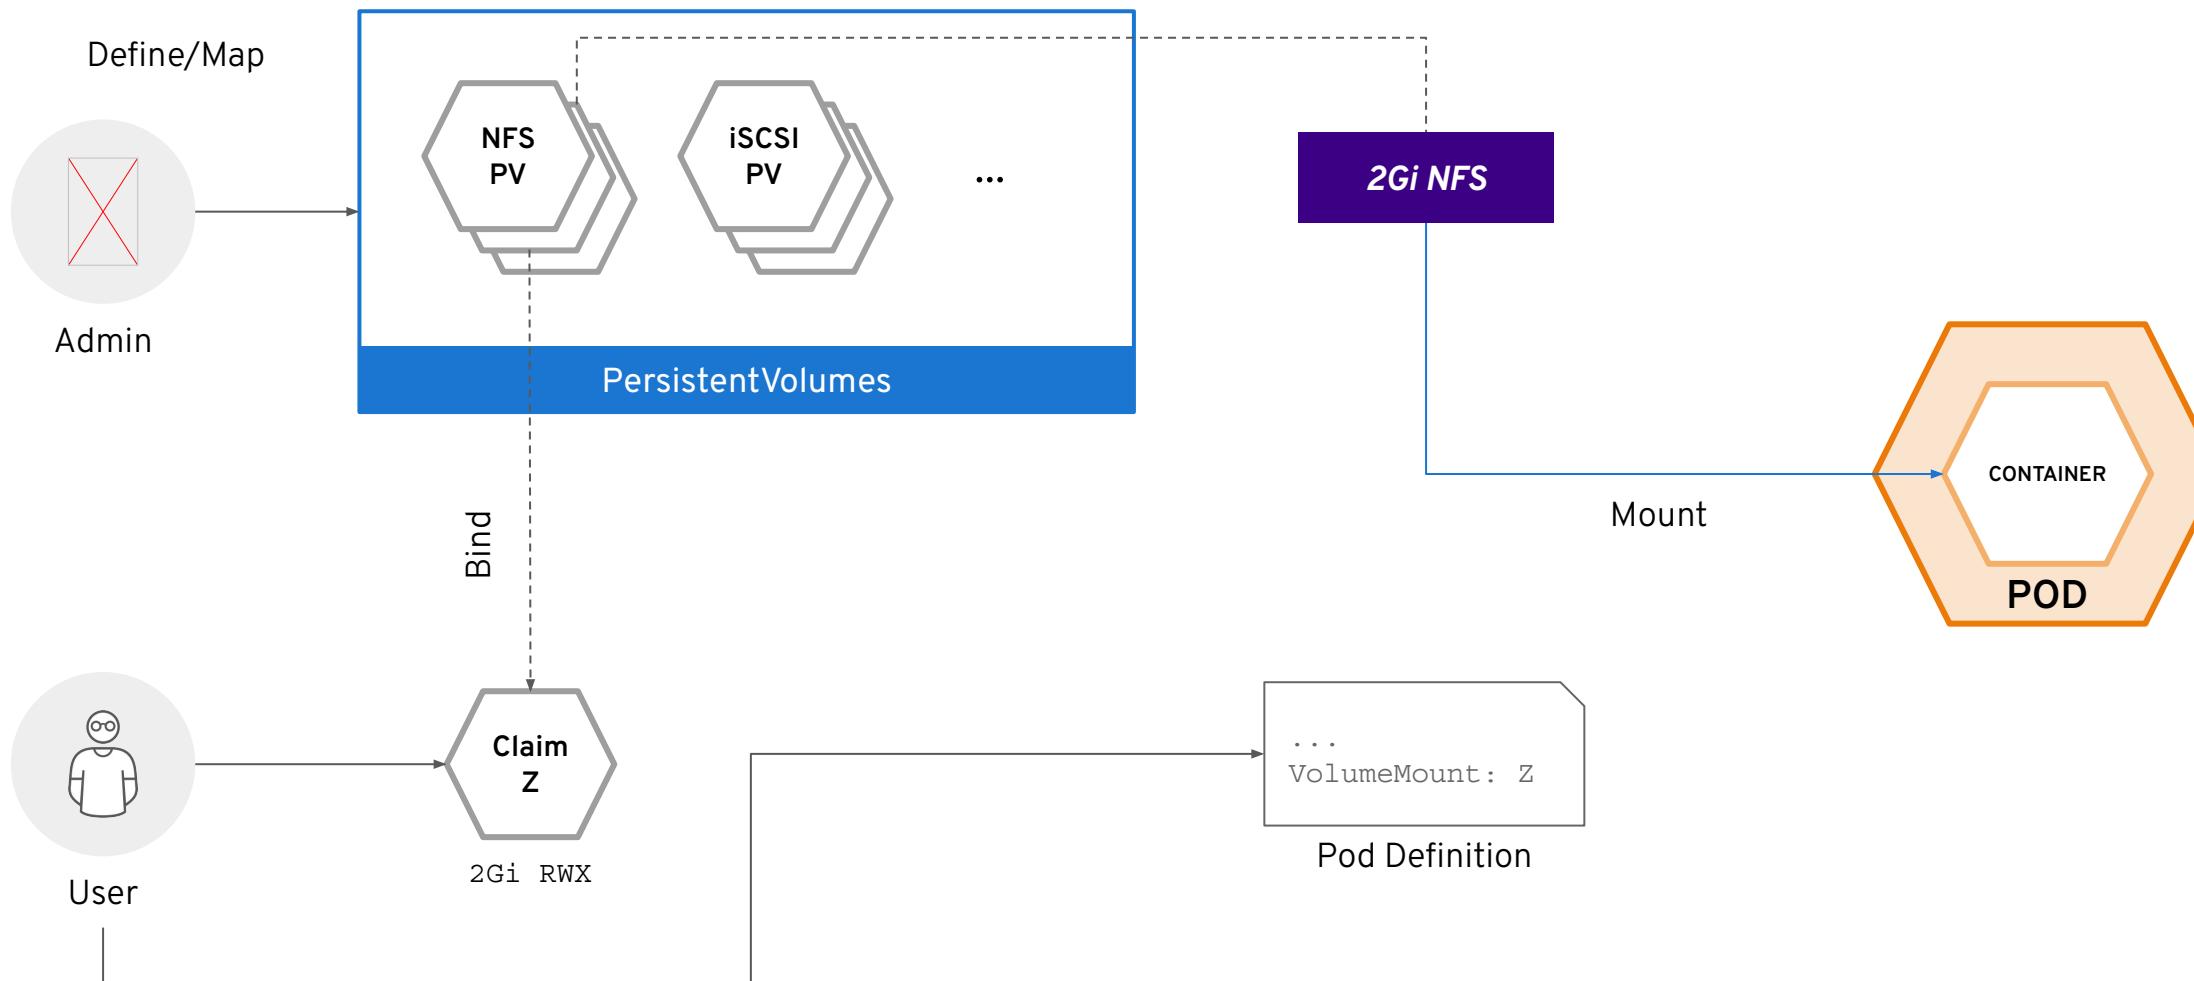
A broad spectrum of static and dynamic storage endpoints



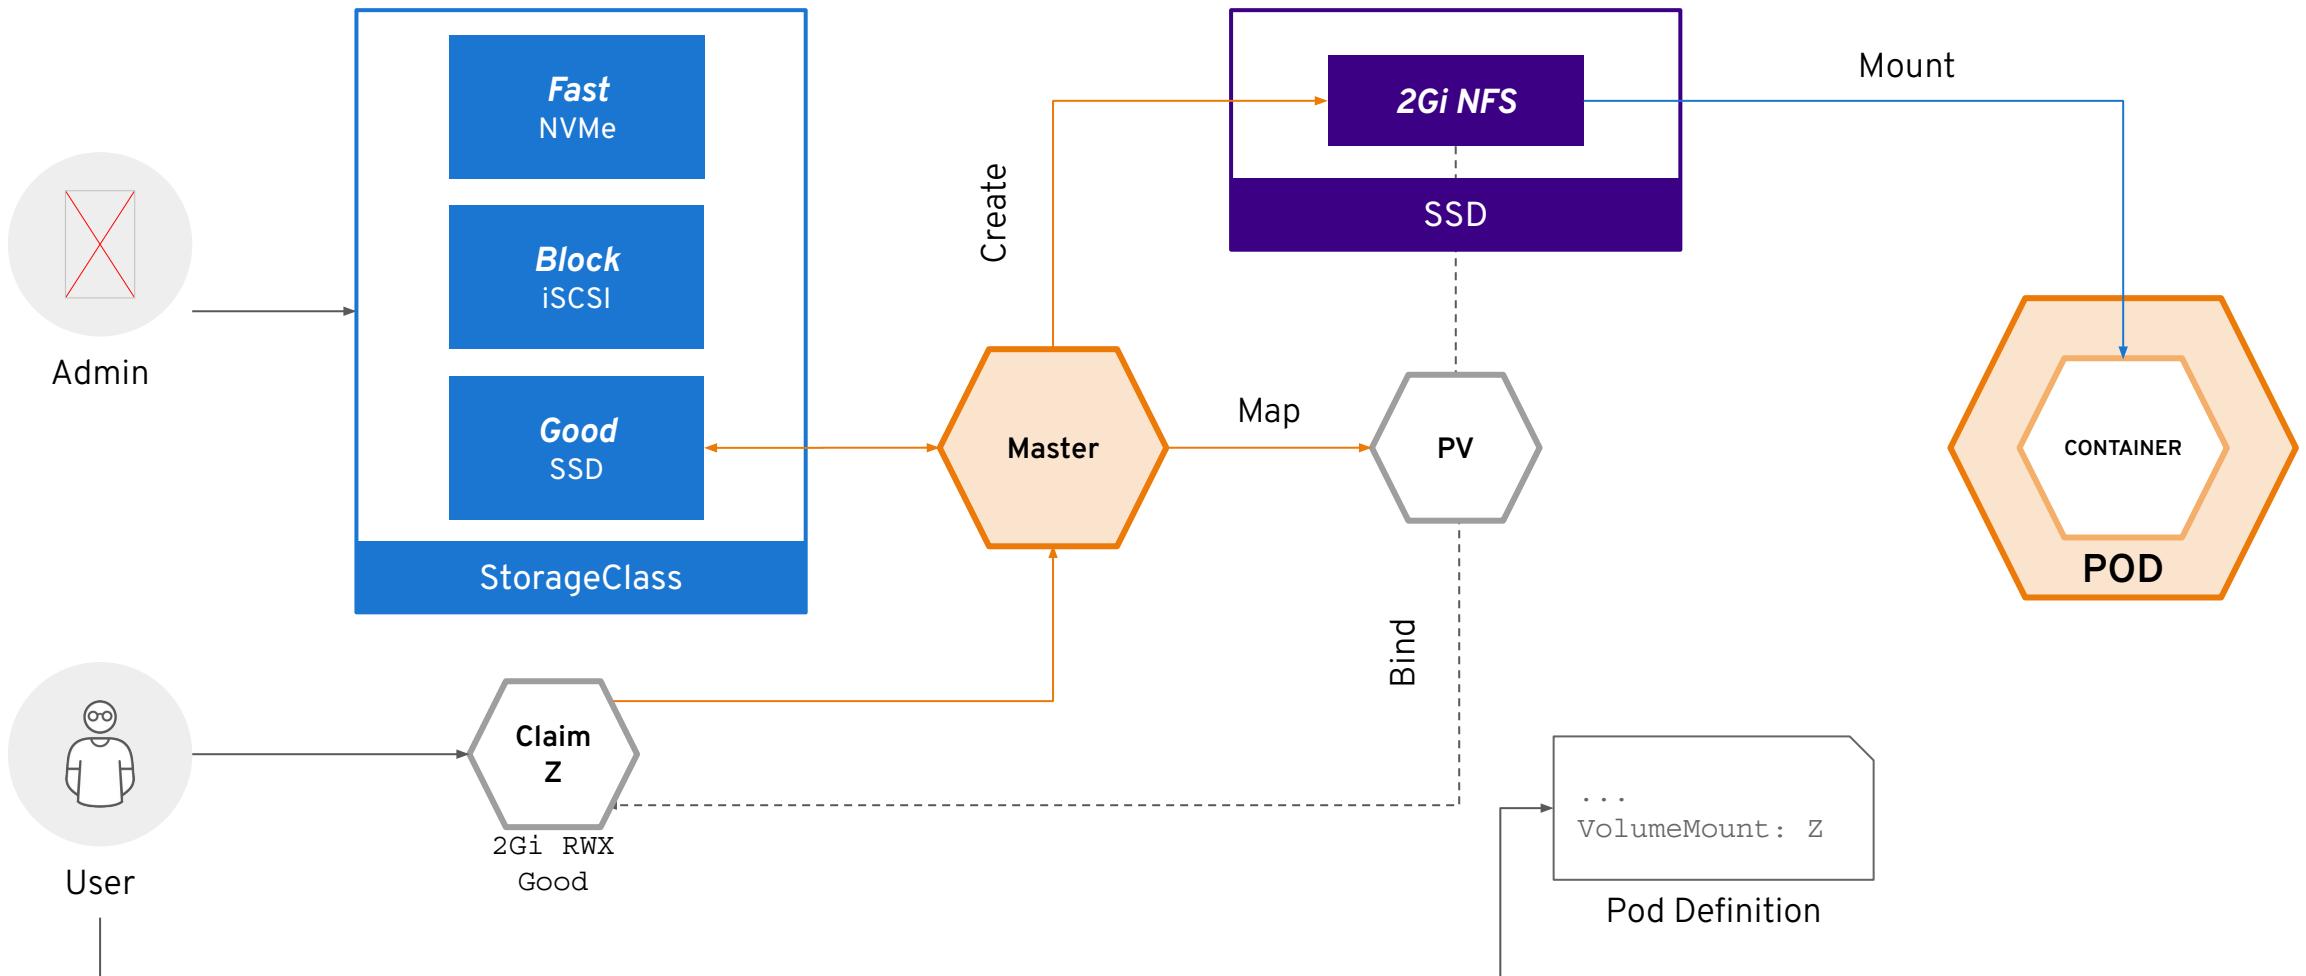
PV Consumption

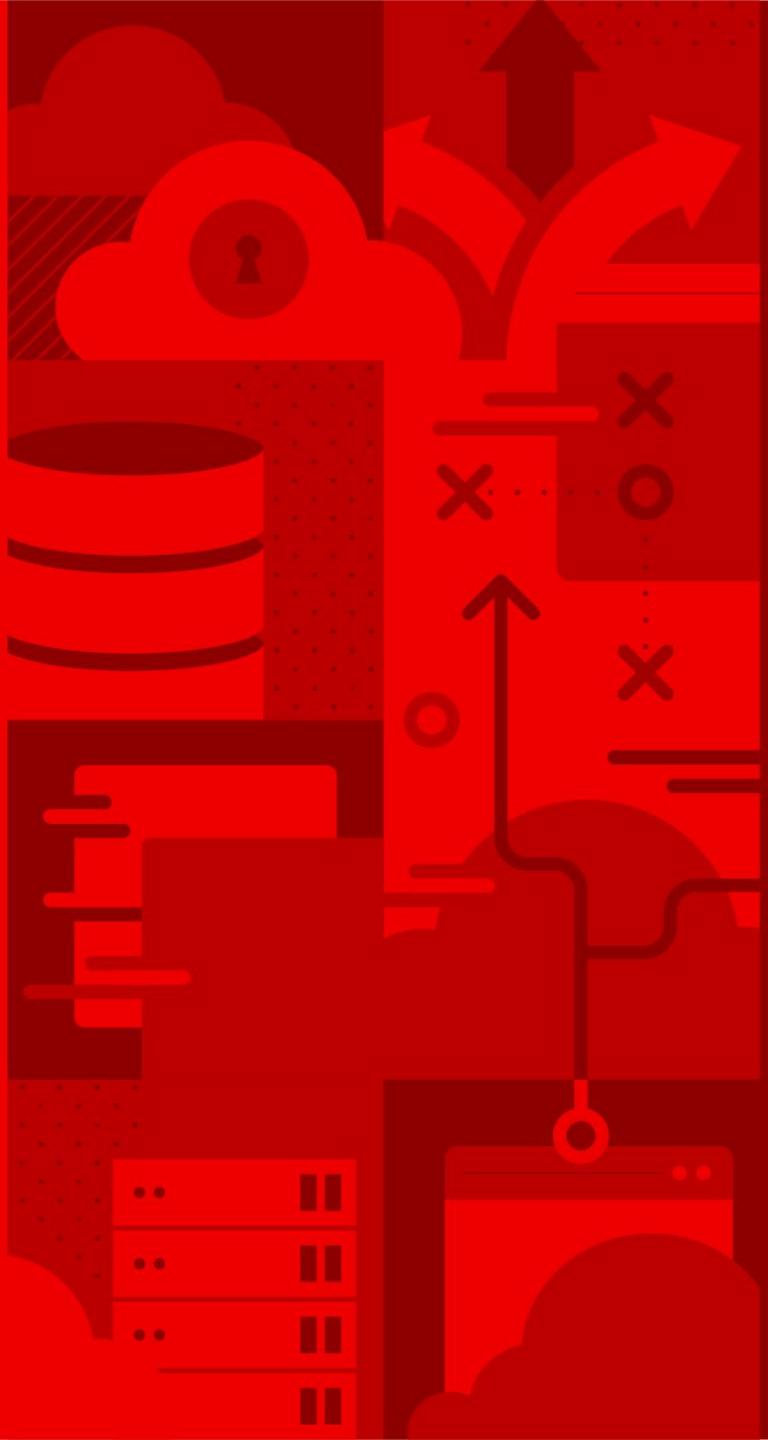


Static Storage Provisioning

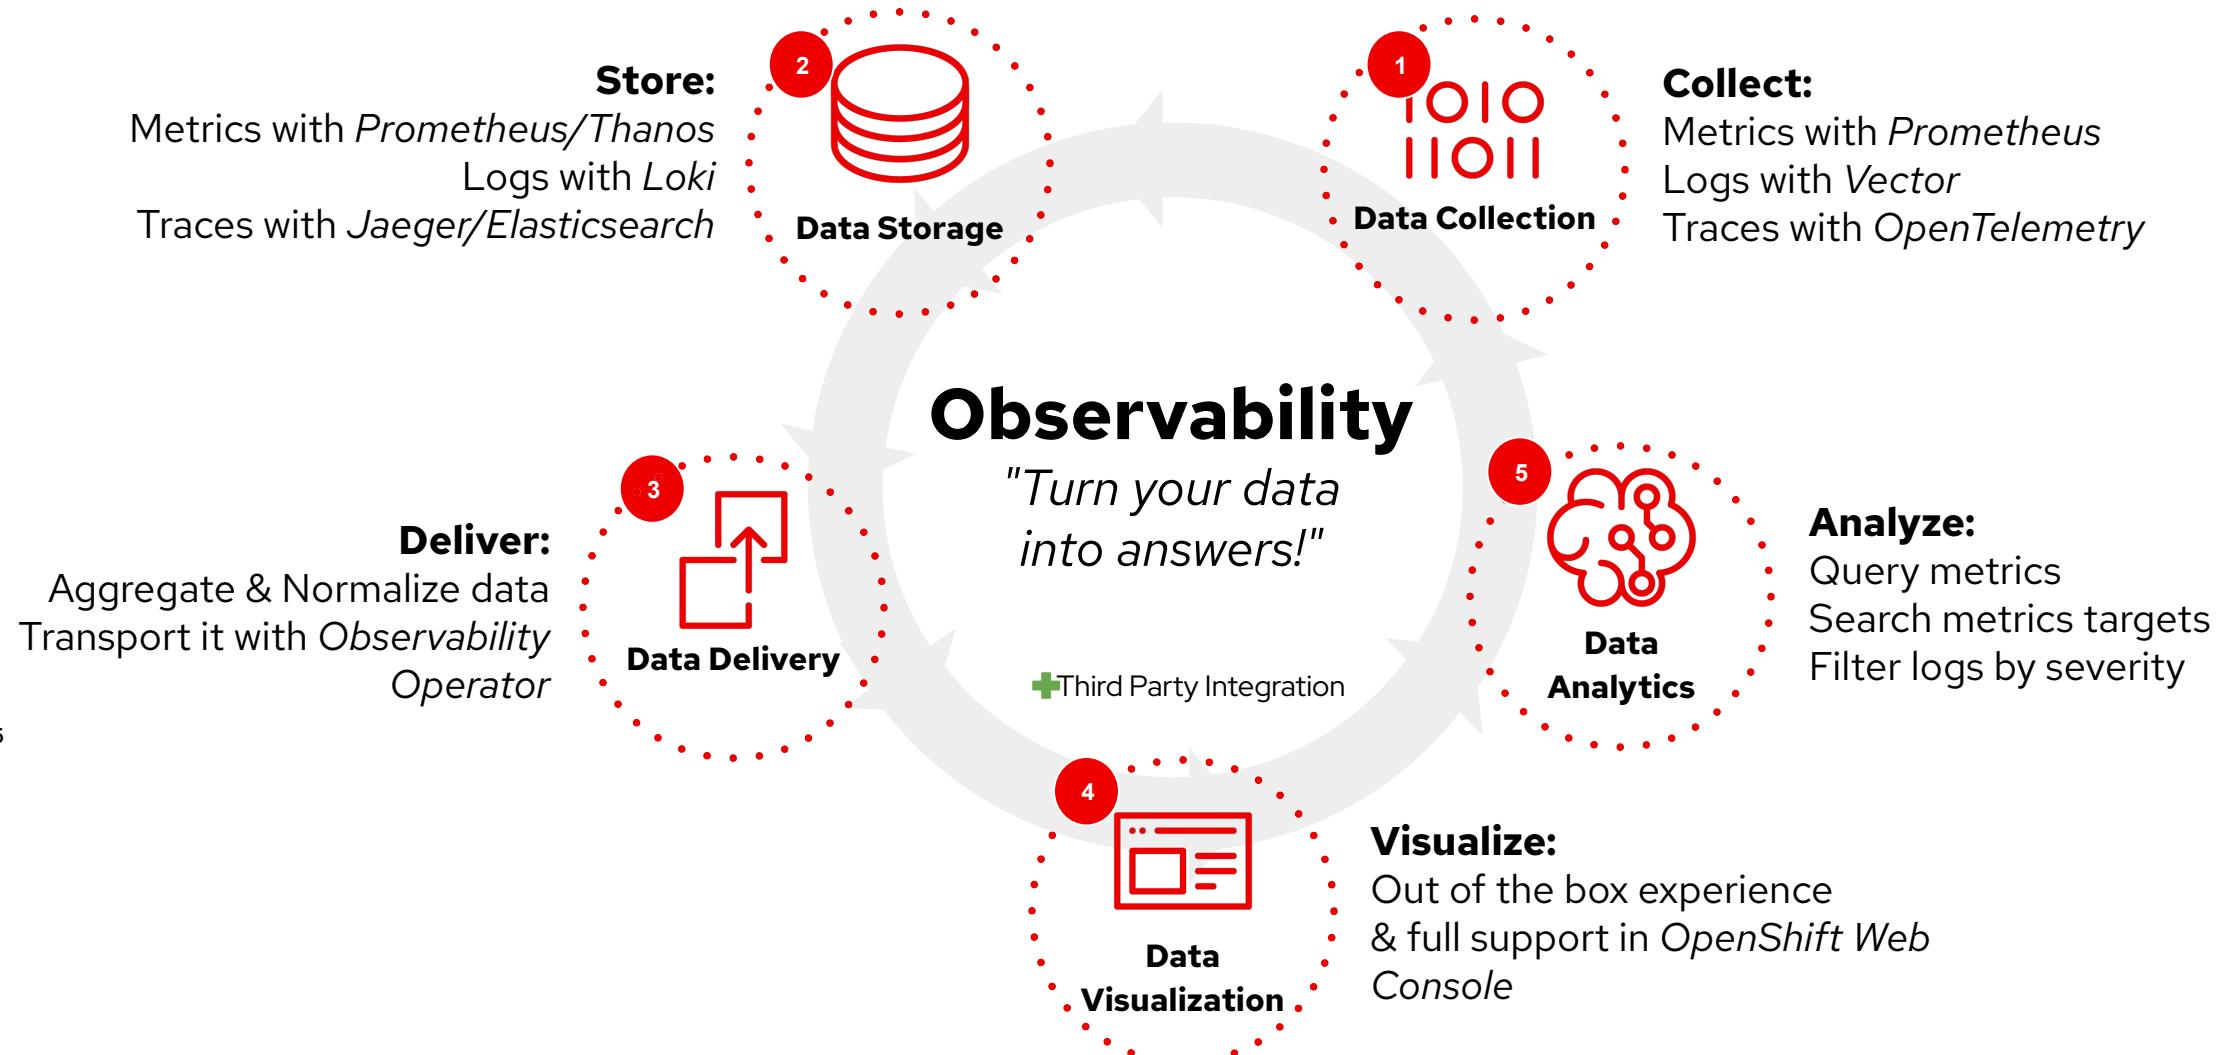


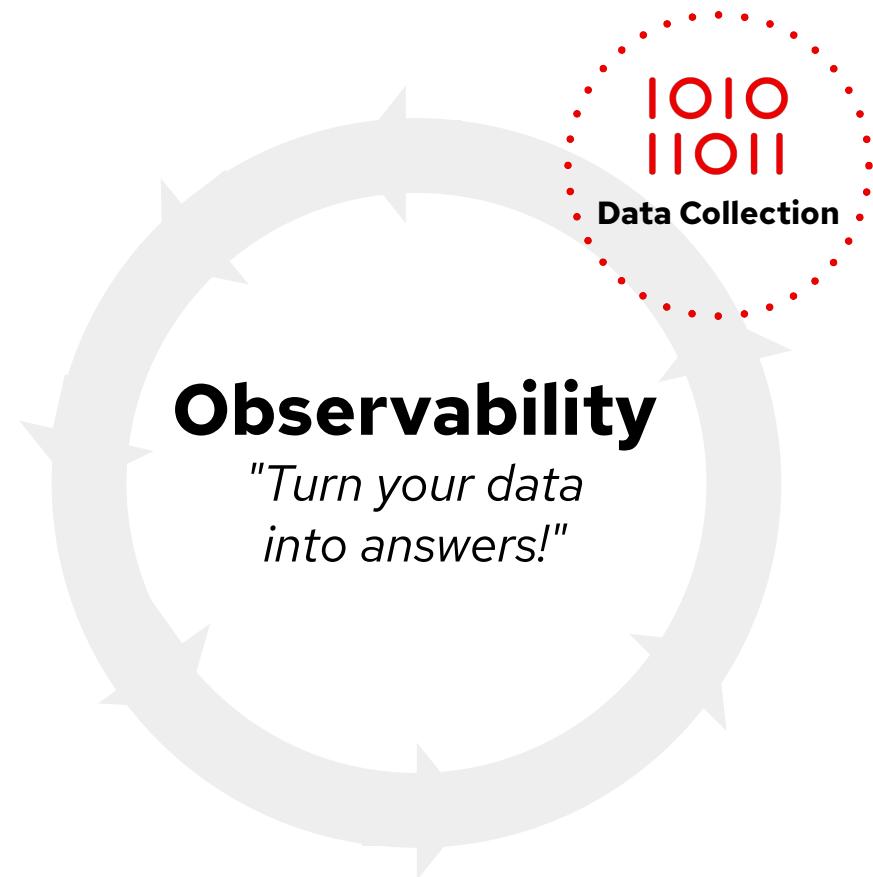
Dynamic Storage Provisioning





Observability



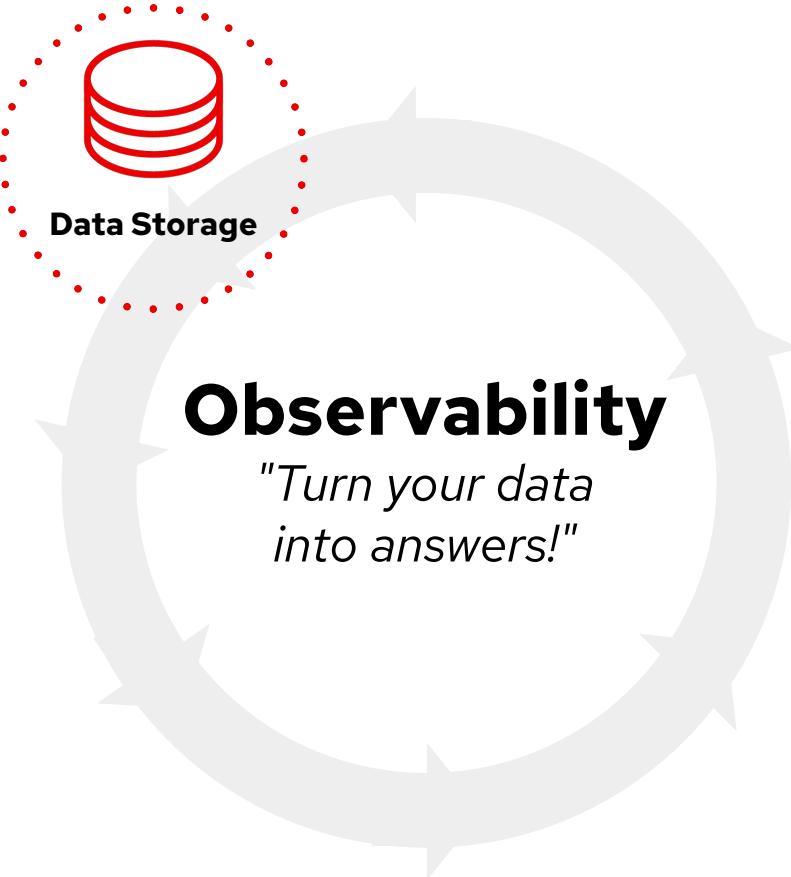


OpenShift 4.12 Monitoring

- ▶ Option to specify Topology Spread Constraints for Prometheus, Alertmanager & Thanos Ruler
- ▶ Option to improve consistency of prometheus-adapter CPU and RAM time series

Logging 5.6

- ▶ GA release of Vector as an alternate collector to Fluentd



OpenShift 4.12 Monitoring

- ▶ Version updates to Monitoring stack components & dependencies

Logging 5.6

- ▶ Exposed stream-based retention capabilities in the Loki Stack custom resource for OpenShift Application owners and OpenShift Administrators



OpenShift 4.12 Monitoring

- ▶ Tech Preview: Allow admin users to create new alerting rules based on platform metrics

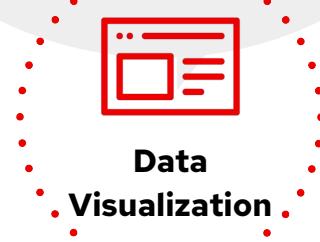
Logging 5.6

- ▶ Support for forwarding logs to Splunk



Observability

*"Turn your data
into answers!"*



69

OpenShift 4.12 Monitoring

- ▶ Improved UX experience in OpenShift Web Console:
Easier selection of records in Metrics UI
- ▶ Support for Alertmanager's negative matchers

Logging 5.6

- ▶ Log Exploration UI available in OpenShift Web Console
 - Developer Perspective: Observe > Aggregated Logs
- ▶ Improved UX experience in OpenShift Web Console:
Custom time range & Predefined filters to easily search logs (namespace, pod, container)

OpenShift Observability

CONFIDENTIAL designator

The screenshot shows the Red Hat OpenShift Developer Console with the 'Aggregated Logs' view selected. The interface includes a sidebar with developer tools like Topology, Observe, Search, Builds, Helm, Project, ConfigMaps, and Secrets. The main area shows a histogram of log events over time (12:09 to 13:09) and a detailed log table with columns for Date, Message, and Severity. A red box highlights the 'Aggregated Logs' tab in the navigation bar, and another red box highlights the search and filter controls at the top of the log table.

New entry:
Aggregated Logs
view in Developer
Console

Improved UX:
Filter by content
(namespace, pod,
container) AND
Search by content
AND Filter by
severity



Observability

*"Turn your data
into answers!"*



Data
Analytics

OpenShift 4.12 Monitoring

- ▶ Runbooks URLs enabled in the Alerting UI of OpenShift Console

Logging 5.6

- ▶ Add the OpenShift cluster ID to log records so that clusters can be uniquely identified in aggregated logs

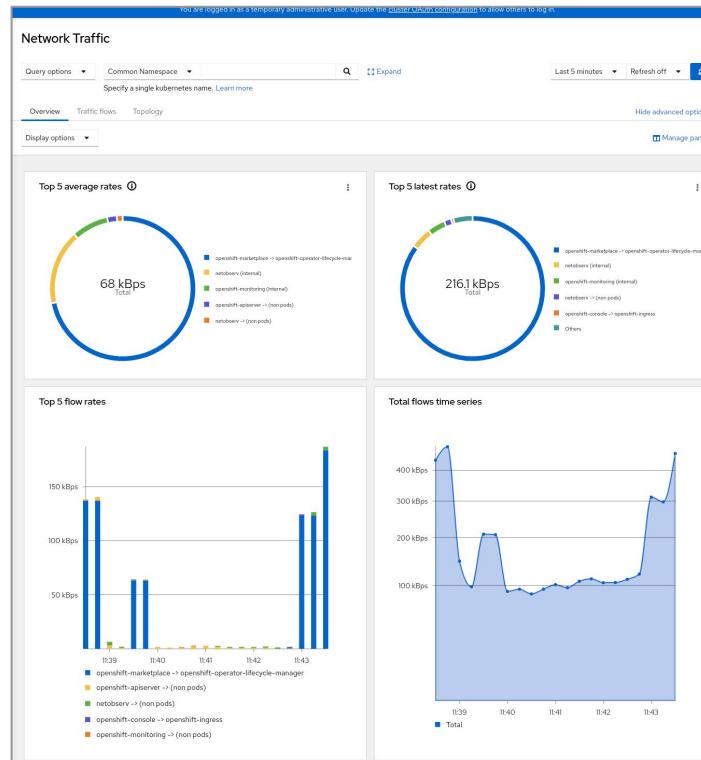


Network Observability

CONFIDENTIAL designator

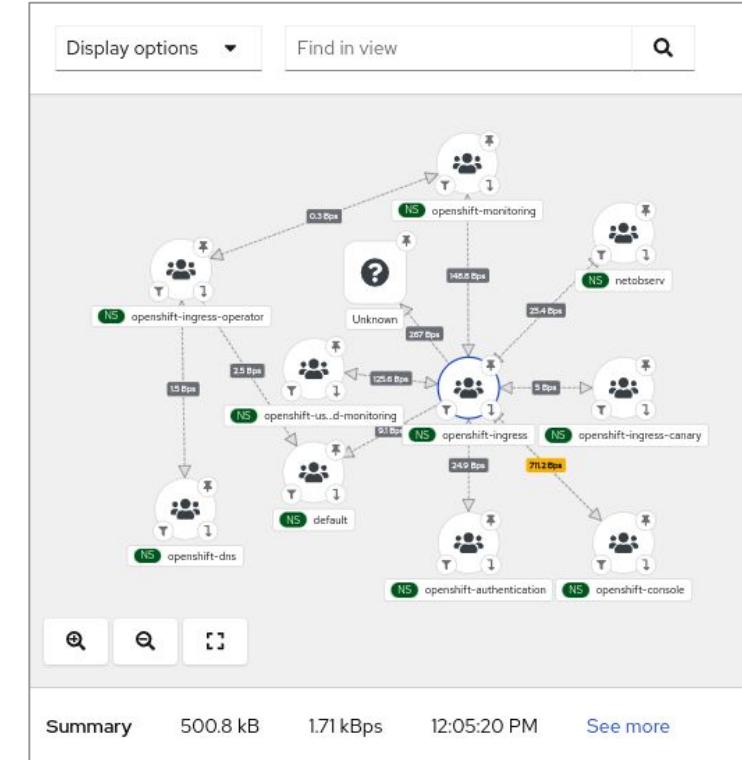
Network Observability GAs at 4.12 for all supported versions of OpenShift at 4.10 or newer

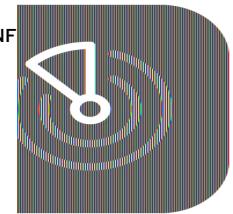
- Integrated with the larger Observability ecosystem, this optional Operator focuses on networking information for a single cluster
- Uses an **eBPF-based** agent on cluster nodes to collect metrics
- Provides observable network traffic metrics, flows, topology and tracing



A table listing network flows from Nov 22, 2022, to Nov 22, 2022. It includes columns for End Time, Source, Namespace, and Port. The table shows various flows between pods and services like installer, ip-10-0-131-74.us-east..., and apiserver.

End Time	Source	Namespace	Port
Nov 22, 2022, 11:51:25.209 AM	P installer-5-ip-10-0-131-74.us-east...	NS openshift-kube-scheduler	8443
Nov 22, 2022, 11:51:25.206 AM	N ip-10-0-131-74.us-east...	n/a	39616
Nov 22, 2022, 11:51:25.206 AM	N ip-10-0-203-192.us-east...	n/a	30377
Nov 22, 2022, 11:51:25.205 AM	P apiserver-7764cf65b-kkb4d	NS openshift-apiserver	8443
Nov 22, 2022, 11:51:25.205 AM	P apiserver-7764cf65b-r5lqg	NS openshift-apiserver	8443
Nov 22, 2022, 11:51:25.204 AM	N ip-10-0-203-192.us-east...	n/a	53633
Nov 22, 2022, 11:51:25.203 AM	P apiserver-7764cf65b-r5lqg	NS openshift-apiserver	8443
Nov 22, 2022, 11:51:25.203 AM	N ip-10-0-131-74.us-east...	n/a	22033
Nov 22, 2022, 11:51:25.202 AM	N ip-10-0-203-192.us-east...	n/a	2380





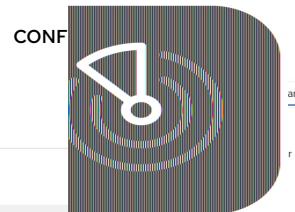
Insights Advisor for OpenShift

- ▶ **Free service leveraging Red Hat experience with supporting and operating OpenShift**
- ▶ New recommendations based on analysis of Kubernetes YAML files (available for managed OpenShift only ATM)
- ▶ Alerts in OpenShift WebConsole for most critical recommendations
- ▶ **New recommendations focused on storage performance, etcd issues etc.**
- ▶ **Improved internal integrations for more stable upgrades**

<https://console.redhat.com/openshift/advisor>

<https://console.redhat.com/settings/notifications/openshift>

Name	Added	Category	Total risk	Clusters
Cluster upgrade will fail when default SCC gets changed	2 years ago	Service Availability	Important	1
Workloads are still using the deprecated APIs which will be removed in the next release	5 months ago	Service Availability	Important	6
SystemMemoryExceedsReservation alerts when the system daemons memory usage on nodes exceeds 90% of the reservation for them	4 months ago	Service Availability	Important	1
Workloads are using the deprecated PodSecurityPolicy API	5 months ago	Performance	Moderate	3
CVE-2021-30465: runc vulnerable to privilege escalation	9 months ago	Security	Moderate	1
Nodes will become Not Ready due to a CRI-O PID leak in the running OpenShift Container Platform version	5 months ago	Service Availability	Moderate	13
The running OpenShift version has reached its End of Life	2 years ago	Service Availability	Moderate	1
Pods could fail to start if openshift-samples is degraded due to FailedImageImport which is caused by a hiccup while talking to the Red Hat registry	2 years ago	Service Availability	Moderate	1
Prometheus metrics data will be lost when the Prometheus pod is restarted or recreated	1 year ago	Service Availability	Moderate	50
The authentication operator is degraded when cluster is configured to use a cluster-wide proxy	2 years ago	Security	Moderate	1
An OCP node behaves unexpectedly when it doesn't meet the minimum resource requirements	2 years ago	Performance	Moderate	2



Insights Cost Management

- ▶ **Free service to monitor per-project and per-cluster spending**
- ▶ Currency support
- ▶ Marketplace services reported including ROSA, ARO, RHEL, ODF, 3rd parties, etc
- ▶ ROSA and ARO costs distributed to projects
- ▶ Costs now distributed according to the same resource consumption criteria in every view
- ▶ Cost of unallocated capacity accounted (both workers and platform)
- ▶ Filtering gained exclude capabilities ("negative filtering")
- ▶ AWS costs default to amortized when Savings Plans are involved
- ▶ Previous month report and custom date picker in Cost Explorer
- ▶ Performance improvements for OCP clusters running on GCP
- ▶ Integration with console.redhat.com notifications

Cost Management Overview

OpenShift Infrastructure

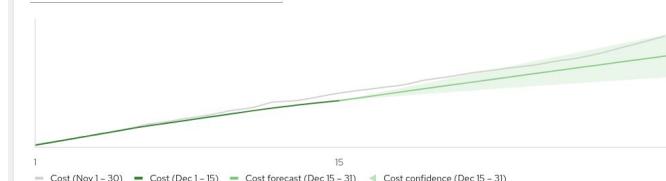
Perspective All OpenShift

All OpenShift cost

\$15,685.32

Cost

All OpenShift cumulative cost comparison (\$)



Top projects	DKK (DKK) - Danish Krone
analytics	EUR (€) - Euro
catalog	GBP (£) - British Pound
install-test	HKD (HK\$) - Hong Kong Dollar
78 Others	JPY (¥) - Japanese Yen
All projects	NOK (NOK) - Norwegian Krone
	NZD (NZ\$) - New Zealand Dollar
	SEK (SEK) - Swedish Krona
	SGD (SGD) - Singapore Dollar
	USD (\$) - United States Dollar
	ZAR (ZAR) - South African Rand

CPU usage and requests

7,159 core-hours

Usage

Memory usage and requests

37,149 GB-hours

Usage

Volume usage and requests

195 GB-month

Usage

Cost Explorer

Perspective All OpenShift Group by Project

Project includes

Custom 2022-11-16 to 2022-12-15

All OpenShift - Top 5 Costliest



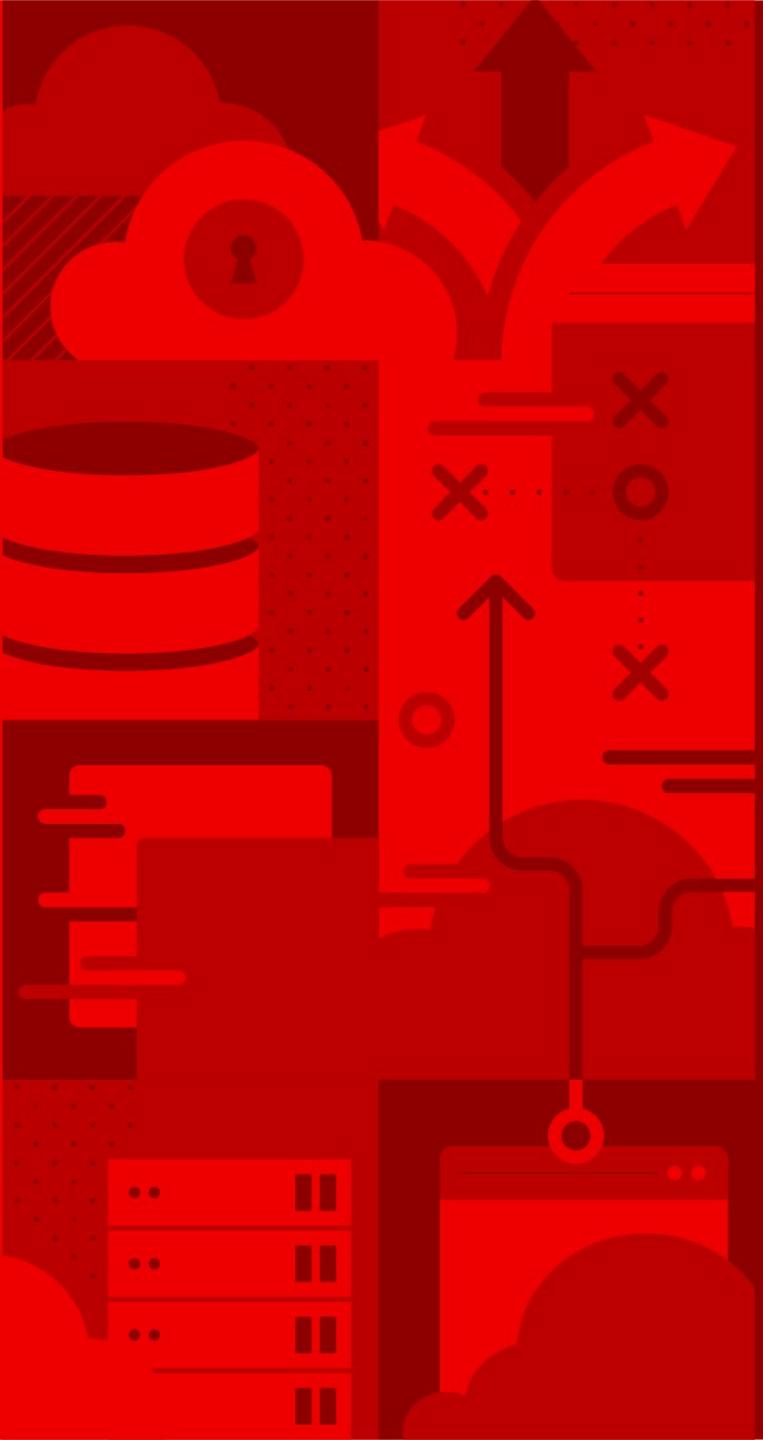
Project names	Nov 16	Nov 17	Nov 18	Nov 19	Nov 20	Nov 21	Nov 22	Nov 23	Nov 24	Nov 25	Nov 26
catalog	\$188.61	\$148.95	\$195.26	\$459.56	\$186.60	\$314.69	\$176.43	\$154.41	\$121.43	\$338.65	\$211.43

<https://console.redhat.com/openshift/cost-management>

<https://console.redhat.com/settings/applications/cost-management>

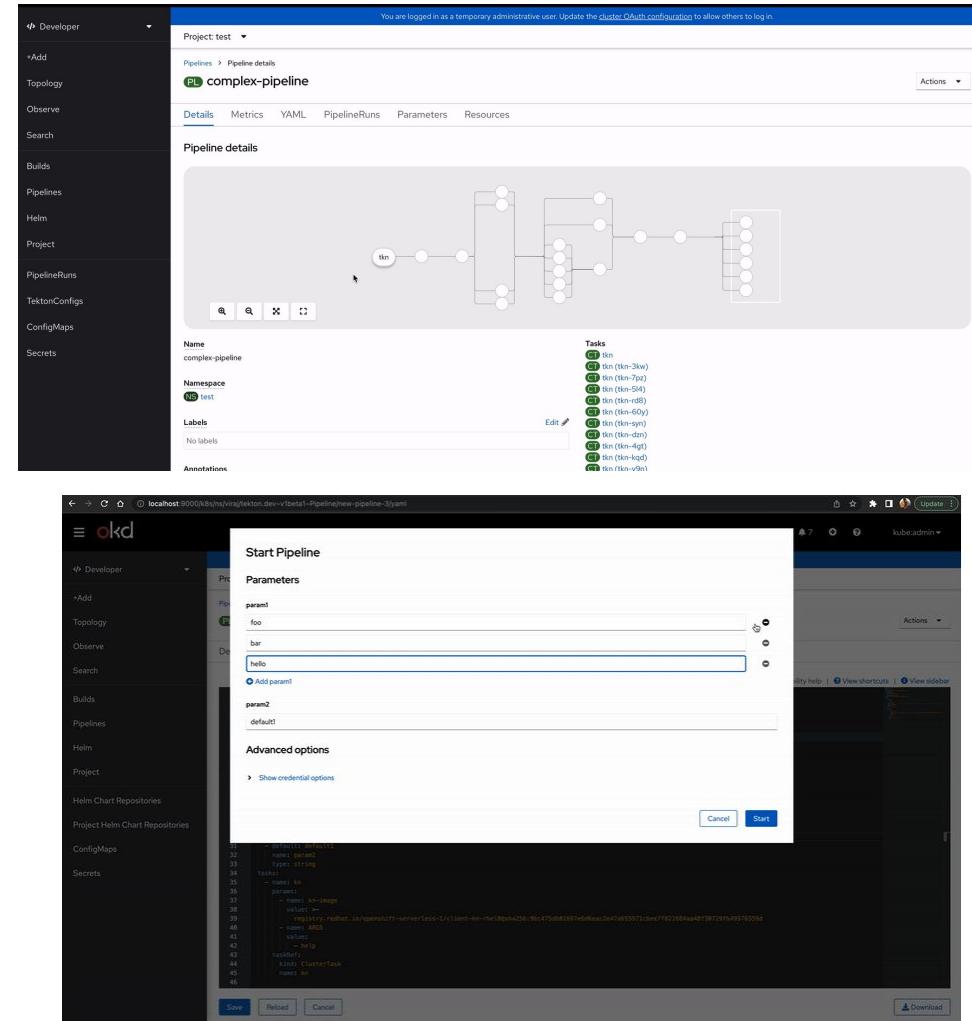
<https://console.redhat.com/settings/notifications/openshift>

Platform Services



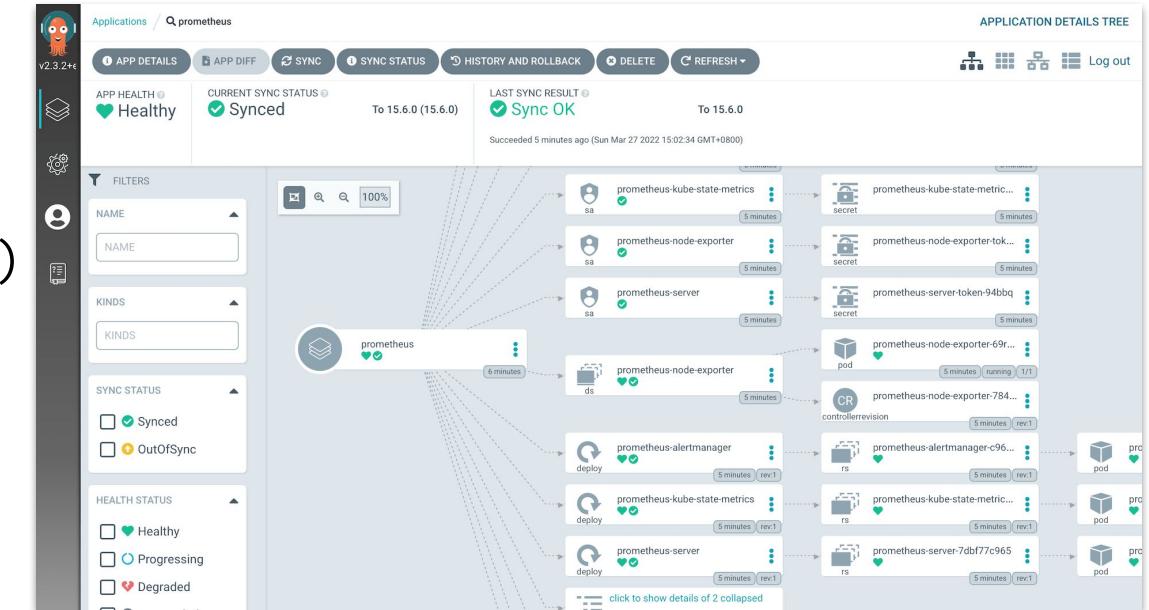
OpenShift Pipelines

- ▶ OpenShift Pipelines 1.9
- ▶ Reference pipelines/tasks in Git, TektonHub, ArtifactHub, etc (Tech Preview)
- ▶ Pipelines as code GA
 - ▶ PAC concurrency control
 - ▶ Support for advanced event matching on filepath/PR title
 - ▶ Ability to enable pac for all [new] repos in a GitHub org
 - ▶ Better errors tooling in Pipelines as Code CLI
 - ▶ Rich PipelineRun details in GitHub Checks UI
- ▶ Support for CSI and projected volume for workspace
- ▶ New CLI: Openshift Pipelines CLI (opc) - Tech Preview
- ▶ Pipelines on Dev Sandbox
- ▶ Dev Console UX improvements : Pipeline topology view, Support of array in Param



OpenShift GitOps

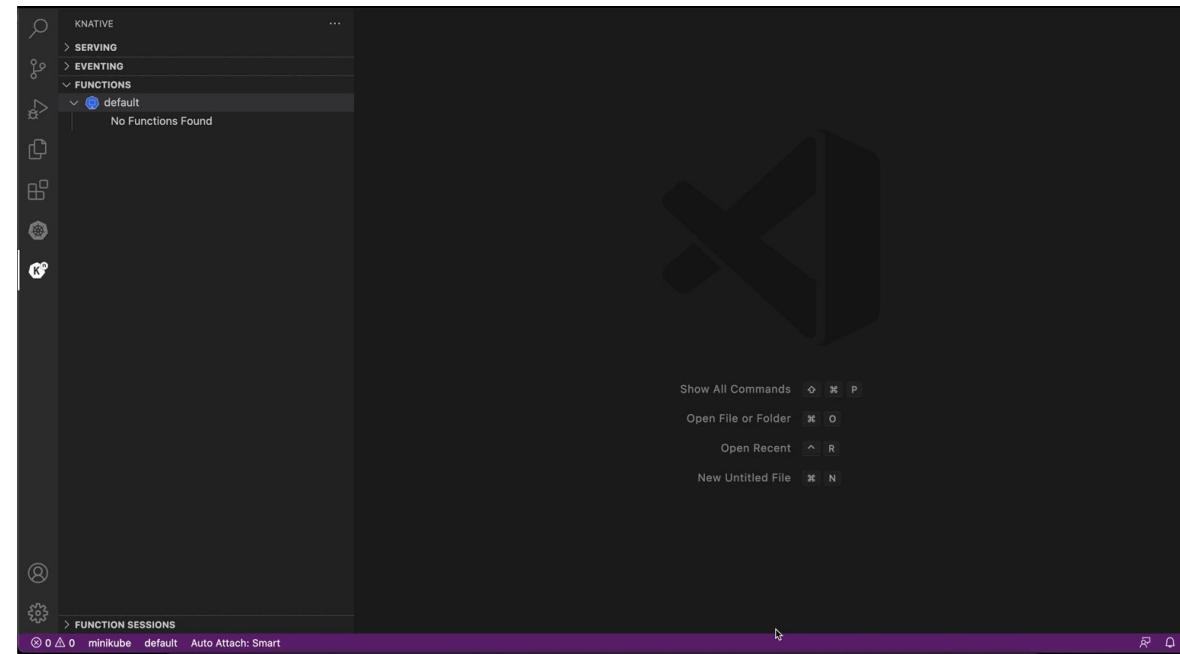
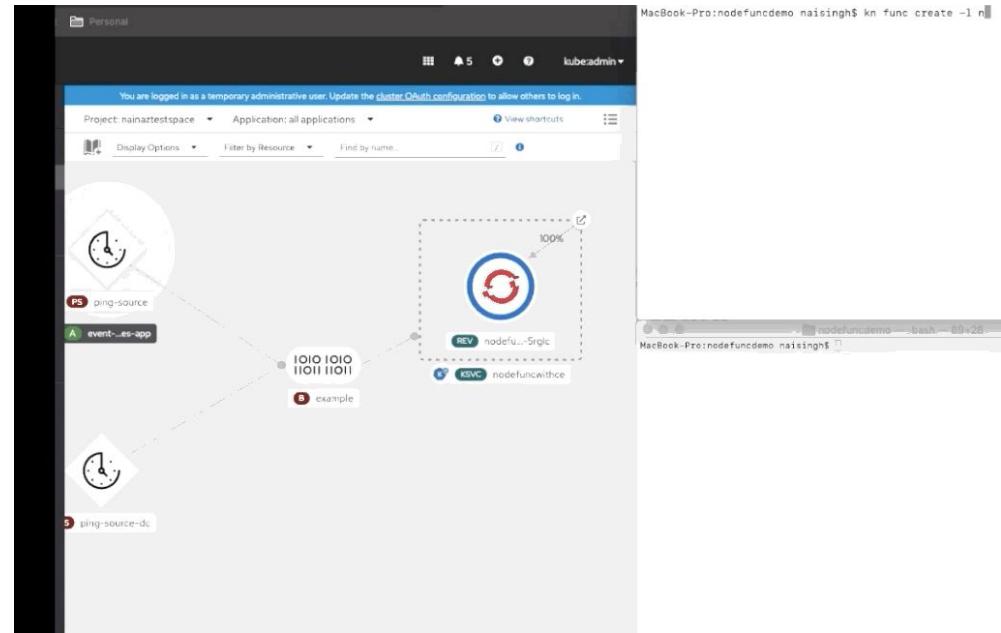
- ▶ OpenShift GitOps 1.7
- ▶ Includes Argo CD 2.6
- ▶ Patching existing resources with Server Side Apply
- ▶ Applications in non-control plane namespaces (TP)
- ▶ Operator improvements:
 - ▶ Custom node selectors
 - ▶ RBAC match mode 'regex'
 - ▶ Sub-keys for resource customizations
 - ▶ Enable/Disable cluster Argo CD console link



OpenShift Serverless

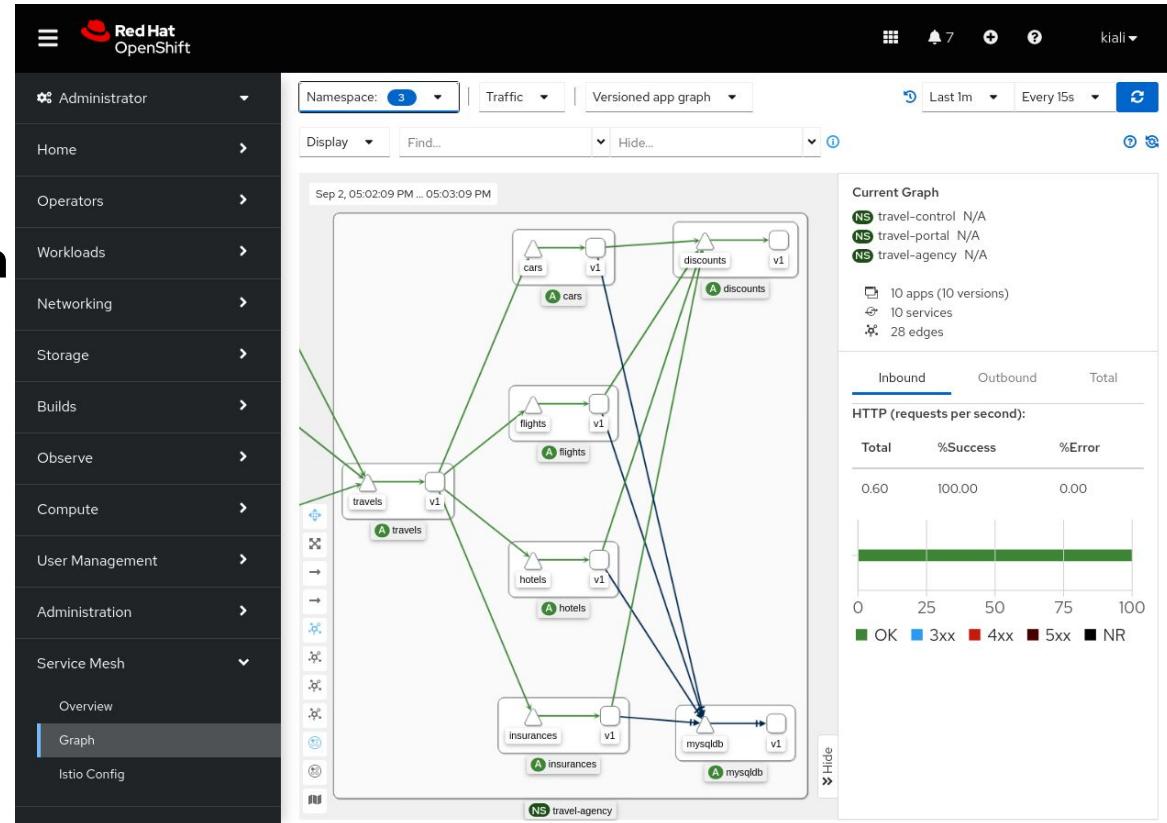
Key Features & Updates

- ▶ Update to Knative 1.6
- ▶ Serverless functions with Quarkus (GA)
 - ▶ In Cluster build using OpenShift Pipelines
 - ▶ Local experience with CLI and IDE (VScode and IntelliJ)
- ▶ Knative Kafka Broker & Knative Kafka Sink (GA)
- ▶ Support for Init Containers and PVC (GA)
- ▶ mTLS natively in Serverless (Tech Preview)
- ▶ Serverless Logic (Dev Preview)
 - ▶ Orchestration for Functions and Services
 - ▶ CLI and Workflow Editor(UX)



OpenShift Service Mesh

- ▶ OpenShift Service Mesh 2.3 is now available
- ▶ Based on **Istio 1.14** and **Kiali 1.57**
- ▶ Introduces GA support for **Gateway Injection**
- ▶ New Technology Preview features:
 - ▶ OpenShift **Console Service Mesh Plugin**
 - ▶ **Cluster-wide mesh** installation option
 - ▶ Kubernetes **Gateway API** (Kiali support added)
- ▶ Service Mesh federation is now supported on Azure Red Hat OpenShift (ARO)



Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.

 [linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)

 [youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)

 [facebook.com/redhatinc](https://www.facebook.com/redhatinc)

 twitter.com/RedHat